A yellow dashed arc composed of four segments, curving from the top-left towards the bottom-right, positioned to the right of the title.


Spatial Transformer Networks

Jaderberg, Max, Karen Simonyan, and Andrew Zisserman.
Advances in neural information processing systems 28 (2015).



2022.11.11
Advanced Computer Vision
KENTECH Energy AI
VIEW Lab
Sohee Kim


◆ Introduction



- The Convolutional Neural Network (CNN) define an exceptionally powerful class of models, but has **Limitation** : **lack of ability to be Spatially invariant to the input data**
== invariance of the model towards spatial transformations of images (rotation, translation, scaling)
- Max-pooling layer – satisfy this, but 2x2 pixel-wise operation is difficult to cope with various spatial variability
⇒ **Spatial transformer module** – include in to standard neural network

Spatial transformer module

- A learnable module that can be placed in a CNN, to increase the spatial invariance in an efficient manner
- **A dynamic mechanism** – unlike pooling layer (receptive fields are fixed and local)
- Input image → transformation is performed on the entire feature map (can include **scaling, cropping, rotations, non-rigid deformations**)
- Networks select regions that are most **relevant**(attention), transform those regions to a **canonical**(일반적인) pose

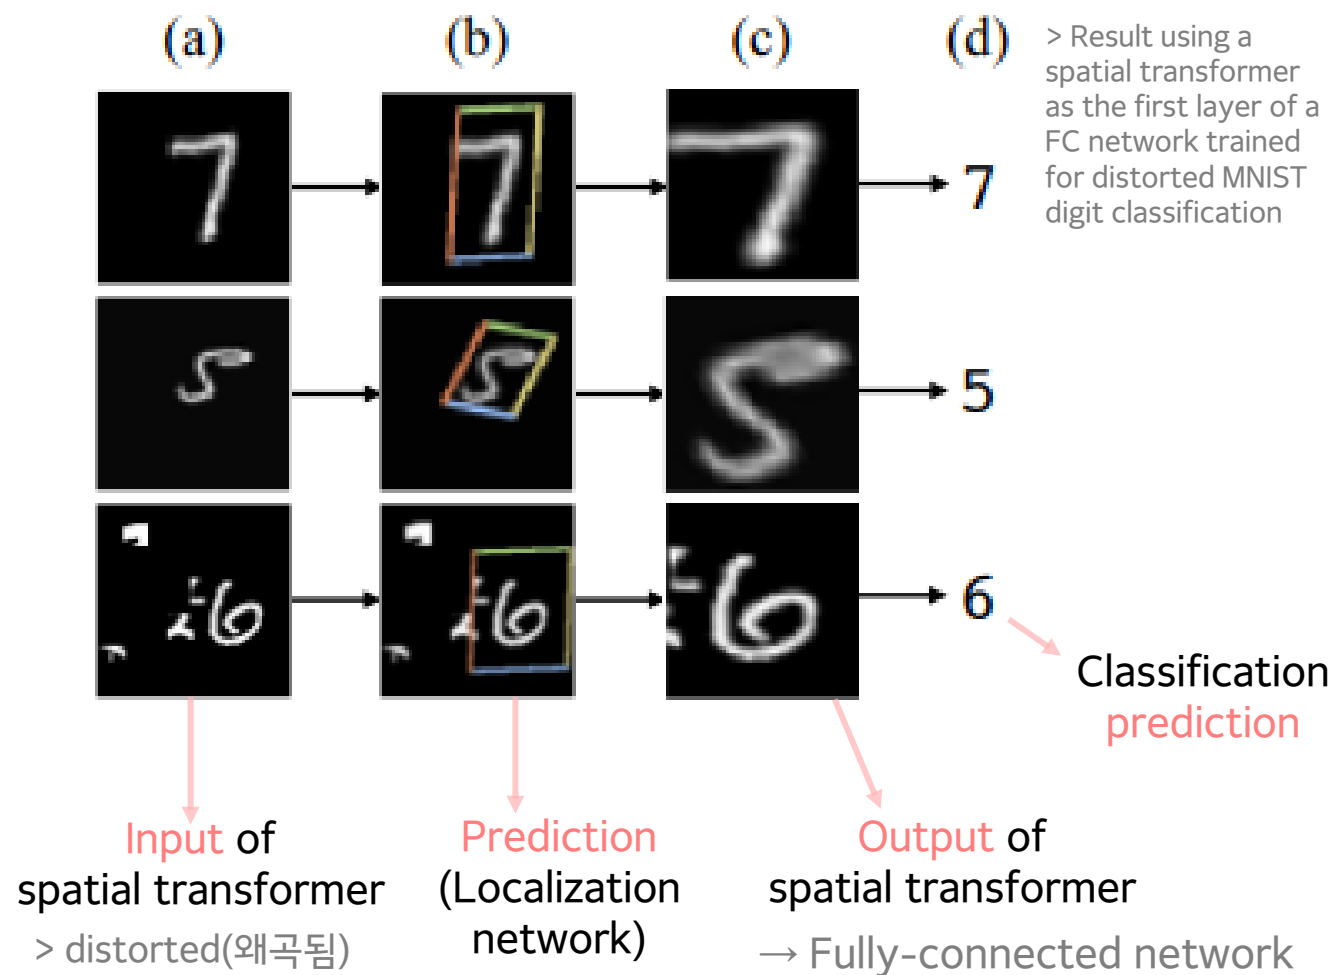


◆ Introduction

❖ Image classification

Using Spatial Transformer

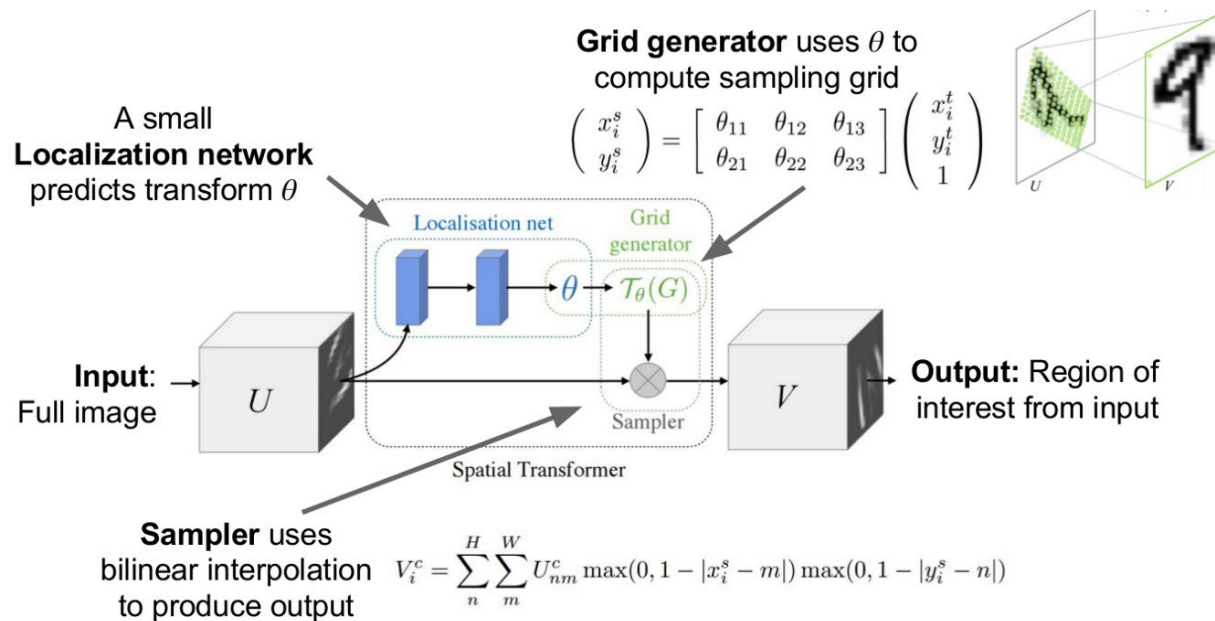
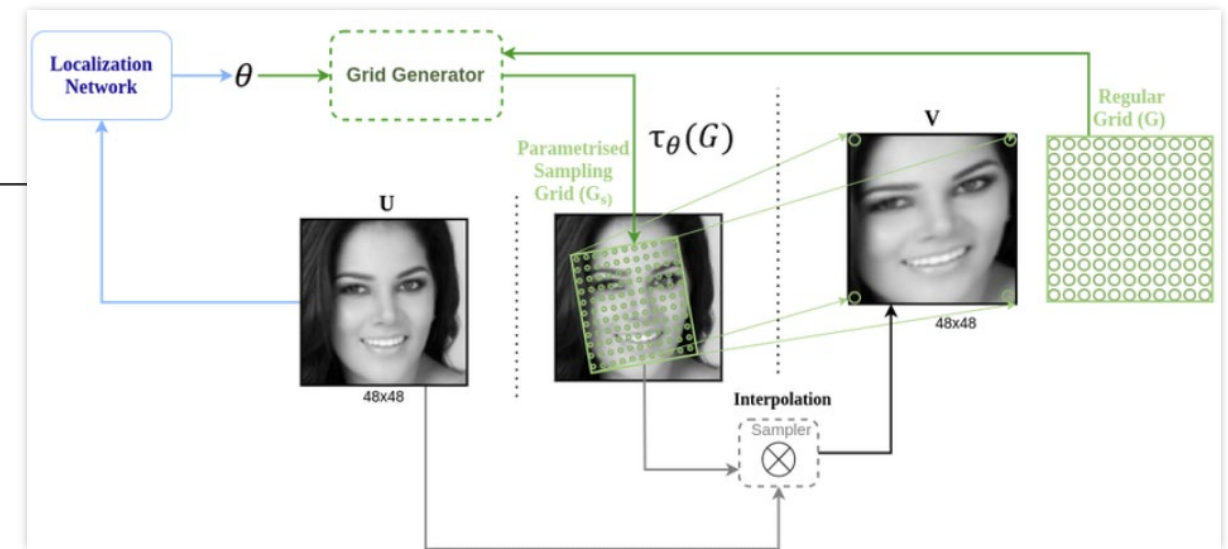
- The action of the spatial transformer is conditioned on **individual data samples**, trained well without extra supervision
- Can be trained with standard **back-propagation** (allowing **end-to-end** training of models)
- Can be added into CNNs to help with a variety of tasks (image classification, co-localization, spatial attention)
- The framework they present in this paper can be seen as a **generalisation of differentiable attention to any spatial transformation**



◆ Spatial Transformers

❖ Spatial Transformer

1. Localisation network
2. Grid generator
3. Sampler



1. Input feature map $U \Rightarrow$ localisation net
 \Rightarrow Transformation parameter θ
2. $\theta \Rightarrow$ Grid generator \Rightarrow Sampling grid $T_\theta(G)$
 (include sampling point location)
3. $U + T_\theta(G) \Rightarrow$ Sampler
 $\rightarrow T_\theta(G)$ has sampling point – apply to U
 \Rightarrow Output feature map V

◆ Spatial Transformers

1. Localisation network

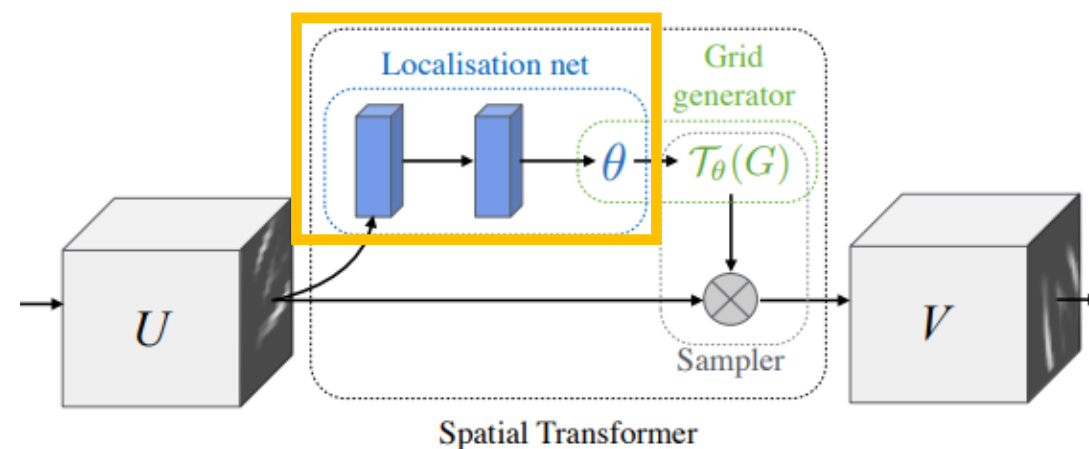
- The localisation network is composed of fully-connected layers or convolution layers
- Predict **Transformation parameter matrix** θ as output
- Input : $U \in R^{H \times W \times C}$ (width W, height H and C channels) $\Rightarrow \theta = f_{loc}(U)$
- The size of θ vary depending on transformation type (ex. affine=6-dimensional)
- Localisation network function $f_{loc}(U)$ can take any form (Convolution layer, FC layer) but **Must include a final regression layer** to produce the transform parameters.


$$A_{\theta} = \begin{bmatrix} s & 0 & t_x \\ 0 & s & t_y \end{bmatrix}$$

> Only scaling + translation

$$\begin{pmatrix} x_i^s \\ y_i^s \end{pmatrix} = \mathcal{T}_{\theta}(G_i) = \underline{A_{\theta}} \begin{pmatrix} x_i^t \\ y_i^t \\ 1 \end{pmatrix} = \begin{bmatrix} \theta_{11} & \theta_{12} & \theta_{13} \\ \theta_{21} & \theta_{22} & \theta_{23} \end{bmatrix} \begin{pmatrix} x_i^t \\ y_i^t \\ 1 \end{pmatrix}$$

Transformation
parameter matrix





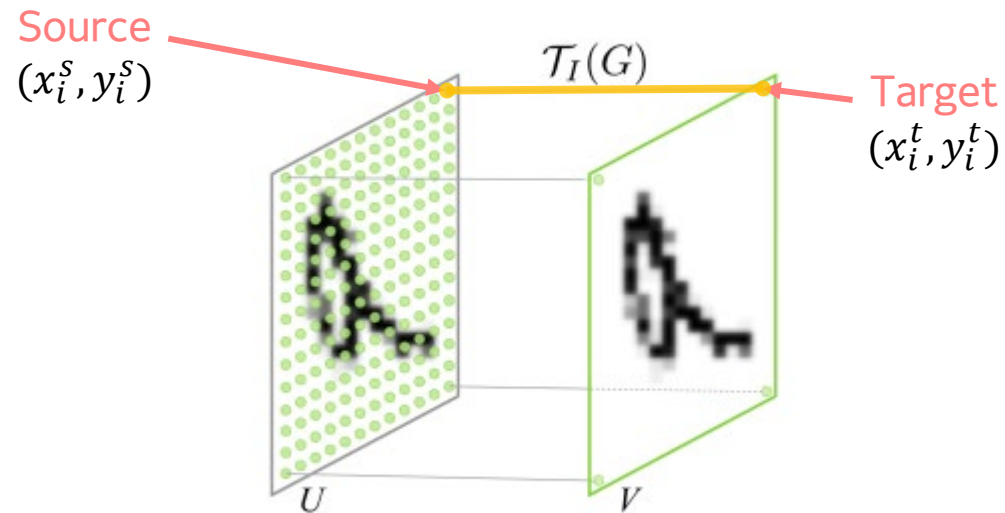
◆ Spatial Transformers

2. Parameterised Sampling Grid

- Regular grid $G = (x_i^t, y_i^t) \Rightarrow$ Apply transformation T_θ on $G \Rightarrow T_\theta(G)$
- Sampling Grid $T_\theta(G)$** = result of warping the regular grid G with transformation T_θ

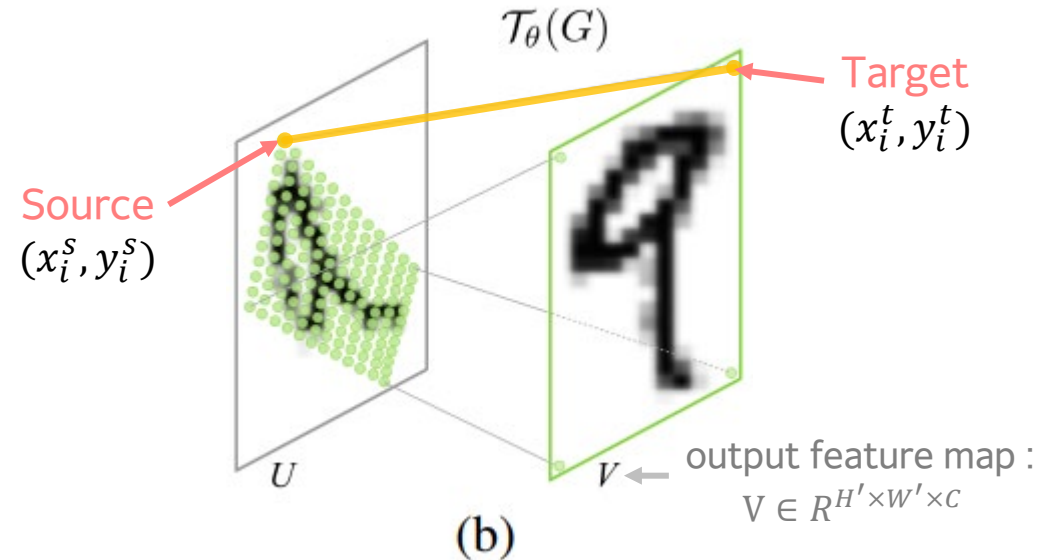
$$\begin{pmatrix} x_i^s \\ y_i^s \end{pmatrix} = \mathcal{T}_\theta(G_i) = A_\theta \begin{pmatrix} x_i^t \\ y_i^t \\ 1 \end{pmatrix} = \begin{bmatrix} \theta_{11} & \theta_{12} & \theta_{13} \\ \theta_{21} & \theta_{22} & \theta_{23} \end{bmatrix} \begin{pmatrix} x_i^t \\ y_i^t \\ 1 \end{pmatrix}$$

✓ Grid generator – **Identity** transformation




(a) Sampling Grid :
Regular grid $G = T_I(G)$

✓ Grid generator – **Affine** transformation



(b) Sampling Grid : $G = T_\theta(G)$



◆ Spatial Transformers

2. Parameterised Sampling Grid

✓ Grid generator Examples

- The transformation can have **any parameterised form**, provided that it is **differentiable** with respect to the parameters

Affine transform

scale, rotation, translation, skew, cropping

$$\begin{pmatrix} x_i^s \\ y_i^s \end{pmatrix} = \mathcal{T}_\theta(G_i) = \mathbf{A}_\theta \begin{pmatrix} x_i^t \\ y_i^t \\ 1 \end{pmatrix} = \begin{bmatrix} \theta_{11} & \theta_{12} & \theta_{13} \\ \theta_{21} & \theta_{22} & \theta_{23} \end{bmatrix} \begin{pmatrix} x_i^t \\ y_i^t \\ 1 \end{pmatrix}$$


Source value Target value

Attention model

isotropic scale, translation, cropping


$$\begin{pmatrix} x_i^s \\ y_i^s \end{pmatrix} = \mathcal{T}_\theta(G_i) = \mathbf{A}_\theta \begin{pmatrix} x_i^t \\ y_i^t \\ 1 \end{pmatrix} = \begin{bmatrix} s & 0 & t_x \\ 0 & s & t_y \end{bmatrix} \begin{pmatrix} x_i^t \\ y_i^t \\ 1 \end{pmatrix}$$

◆ Spatial Transformers



3. Differentiable Image Sampling

- $Sample(T_{\theta}(G), U) \rightarrow V$
- Sampler : the set of sampling points $T_{\theta}(G)$ + input feature map $U \Rightarrow$ produce sampled **output feature map V**
- Each (x_i^s, y_i^s) coordinate in $T_{\theta}(G)$ defines the spatial location in the input U
where **sampling kernel** is applied to get the value at a particular pixel in the output V



- With high probability, location points are not an exact integer (like (1, 2))
→ need **interpolation** of surrounding values
- Any sampling kernel can be used.
- Interpolation function: **Sampling kernel $k()$**

$$V_i^c = \sum_n^H \sum_m^W U_{nm}^c k(x_i^s - m; \Phi_x) k(y_i^s - n; \Phi_y) \quad \forall i \in [1 \dots H'W'] \quad \forall c \in [1 \dots C]$$

Parameters of $k()$

Output value (x_i^t, y_i^t)
Output feature map channel c

Input value
(n, m) value

Sampling grid
coordinate

> General form

◆ Spatial Transformers

3. Differentiable Image Sampling

■ Nearest Integer Sampling


$$V_i^c = \sum_n^H \sum_m^W U_{nm}^c \delta(\lfloor x_i^s + 0.5 \rfloor - m) \delta(\lfloor y_i^s + 0.5 \rfloor - n)$$

Target value

Source value

$\lfloor \cdot \rfloor$: rounding to the nearest integer

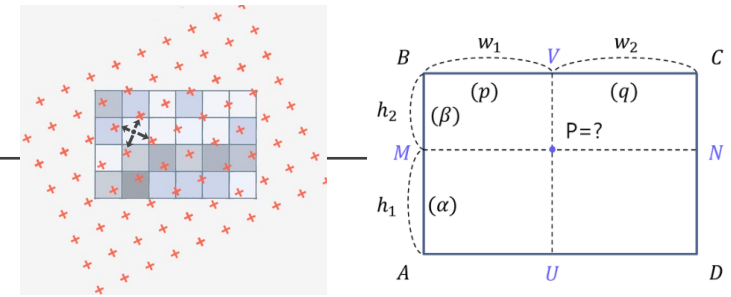
Original Image



Nearest Neighbor



Bilinear Interpolation



■ Bilinear Sampling → linear interpolation on the x, y axes respectively

$$V_i^c = \sum_n^H \sum_m^W U_{nm}^c \max(0, 1 - |x_i^s - m|) \max(0, 1 - |y_i^s - n|)$$

$$\frac{\partial V_i^c}{\partial U_{nm}^c} = \sum_n^H \sum_m^W \max(0, 1 - |x_i^s - m|) \max(0, 1 - |y_i^s - n|)$$

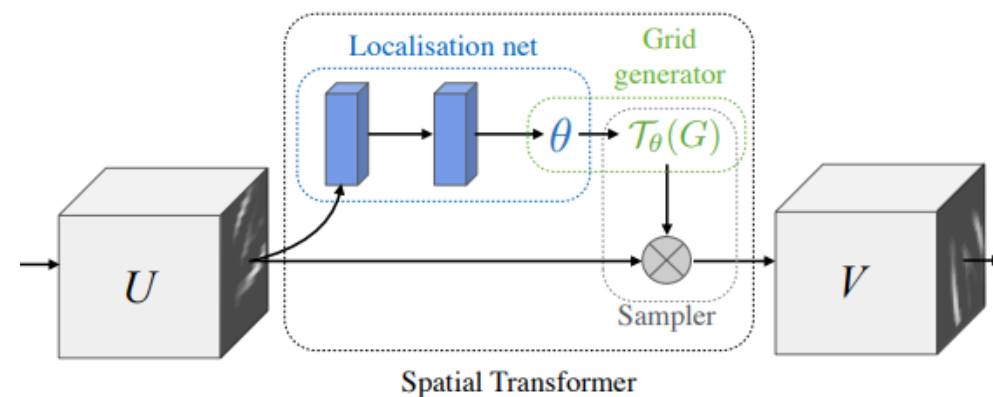
$$\frac{\partial V_i^c}{\partial x_i^s} = \sum_n^H \sum_m^W U_{nm}^c \max(0, 1 - |y_i^s - n|) \begin{cases} 0 & \text{if } |m - x_i^s| \geq 1 \\ 1 & \text{if } m \geq x_i^s \\ -1 & \text{if } m < x_i^s \end{cases}$$

- It is a (sub-)differentiable sampling mechanism so that it is convenient for backpropagation
- If non-differentiable, backpropagation can be calculated by dividing it by sections.

◆ Spatial Transformers

❖ Spatial Transformer Networks

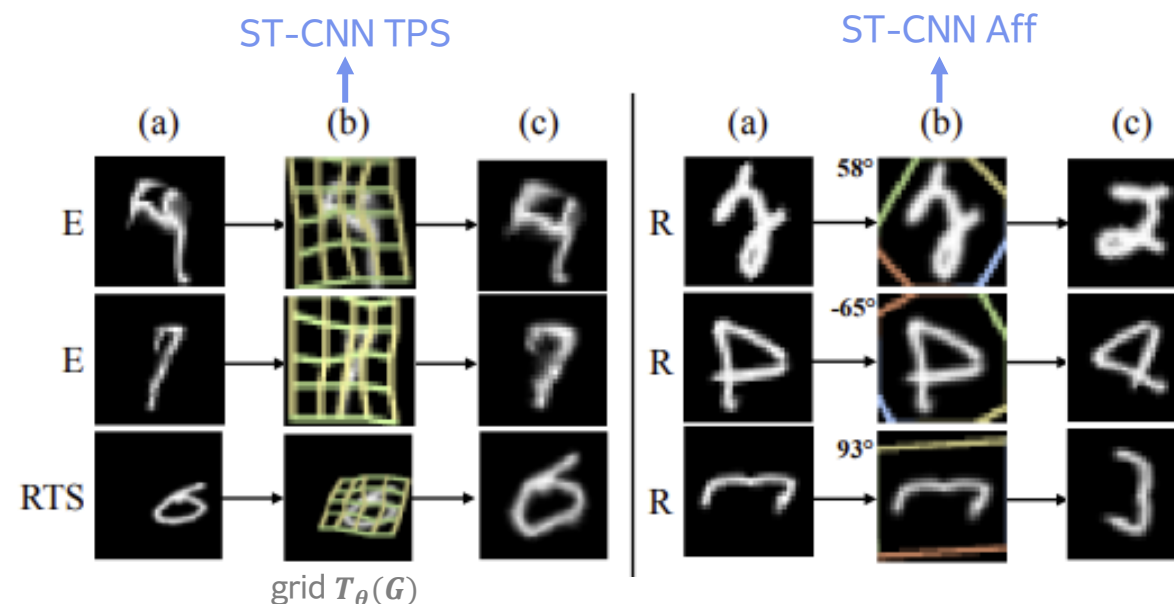
- The combination of Localisation network, grid generator and sampler form **spatial transformer module**
+ CNN architecture \Rightarrow **Spatial Transformer Network**
- Spatial transformer module can be dropped into a CNN architecture **at any point, and in any number**
- Learn **during the training process** to minimize the overall cost function of CNN \Rightarrow it has little effect on training speed
- The knowledge(how to transform) is compressed and cached in the weights of the localisation network during training
- It is possible to have **multiple** spatial transformers in a CNN.
 - To put the ST layer just before the CNN input : most efficient



◆ Experiments

1. Distorted MNIST

✓ Model	Model	MNIST Distortion			
		R	RTS	P	E
• FCN, CNN	FCN	2.1	5.2	3.1	3.2
	CNN	1.2	0.8	1.5	1.4
• ST + FCN	Aff	1.2	0.8	1.5	2.7
	Proj	1.3	0.9	1.4	2.6
	TPS	1.1	0.8	1.4	2.4
• ST + CNN	Aff	0.7	0.5	0.8	1.2
	Proj	0.8	0.6	0.8	1.3
	TPS	0.7	0.5	0.8	1.1

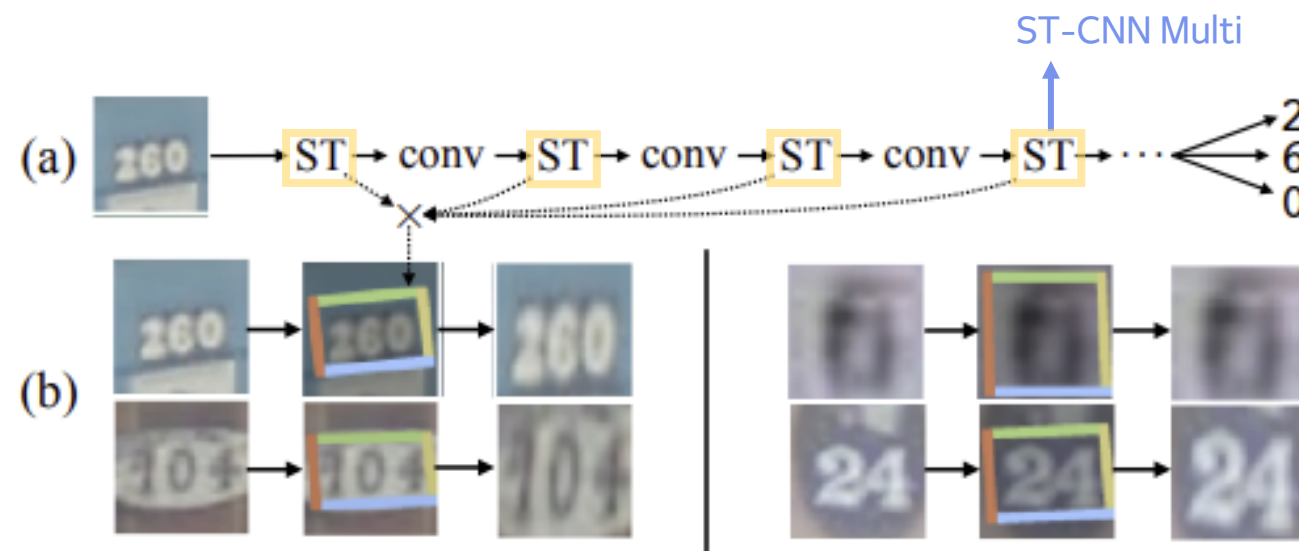


- MNIST dataset – distorted in various ways (4)
 - rotation (R), rotation-translation-scale (RTS), projective transformation (P), elastic warping (E)
- As we can see, ST-FCN outperforms FCN and ST-CNN outperforms CNN.
 - **ST-CNN models** consistently perform better than ST-FCN models
- **CNN** > FCN \Rightarrow The CNN models include two max-pooling layers \rightarrow Maxpooling layer makes spatial invariance high
- Spatial transformation : affine (AFF), projective (Proj), thin plate spline transformation (TPS)
 - Between different classes of transformation, the **TPS** is the most powerful

◆ Experiments

2. Street View House Numbers

Model	Size	
	64px	128px
Maxout CNN [10]	4.0	-
CNN (ours)	4.0	5.6
DRAM* [1]	3.9	4.5
ST-CNN	Single	3.7
	Multi	3.6

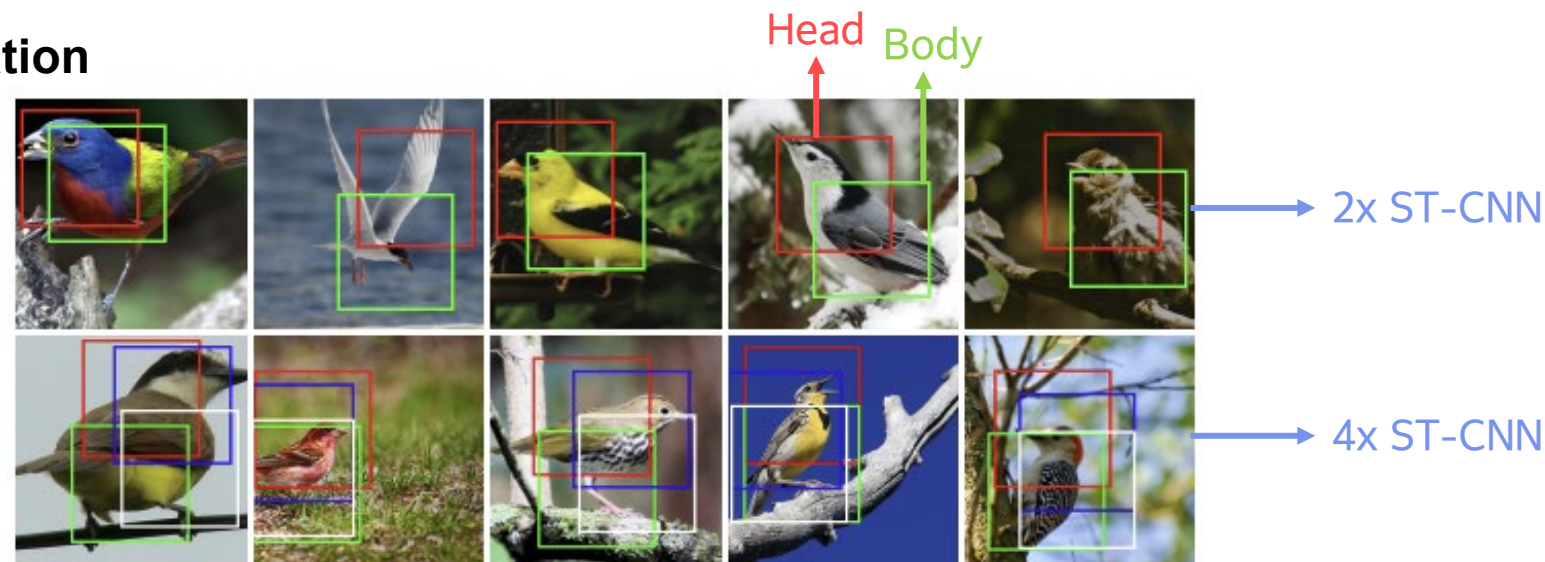


- This dataset contains around 200k real world images of house numbers
- Put ST layer → CNN convolutional stack
 - ST-CNN Single : Only one ST at the beginning of network
 - ST-CNN Multi : One ST before each convolutional layer
- Affine transformation, bilinear sampler is used
- **ST-CNN** outperforms Maxout and CNN
 - **ST-CNN Multi** outperforms ST-CNN Single a bit- 속도 6% 느려짐

◆ Experiments

3. Fine-Grained Classification

Model		
Cimpoi '15 [4]		66.7
Zhang '14 [30]		74.9
Branson '14 [2]		75.7
Lin '15 [20]		80.9
Simon '15 [24]		81.0
CNN (ours) 224px		82.3
2×ST-CNN 224px		83.1
2×ST-CNN 448px		83.9
4×ST-CNN 448px		84.1



- Use a spatial transformer network with multiple transformers in parallel to perform fine-grained bird classification
- CUB 200-2011 birds dataset - 200 species of birds, 11,788 images
- Strong baseline CNN model : Inception-v2 (ImageNet pre-trained) for classifying 200 species (82.3% accuracy)
 - 2x ST-CNN , 4x ST-CNN : 2 or 4 parallel STs
- 4x ST-CNN achieves an accuracy of 84.1%, outperforms the baseline by 1.8%
- Interesting behaviour : Each box = ST found (without supervision)
 - 2x ST-CNN : red ST – head detector, green – detect central part of the body

◆ Conclusion

- **A new self-contained module for neural networks – the spatial transformer**
- This module can be dropped into a network and perform explicit spatial transformations of features
- **Differentiable** and learnt in an **end-to-end** fashion – **without making any changes to the loss function**
- While CNNs provide an incredibly strong baseline, they see gains in accuracy using spatial transformers across multiple tasks, resulting in **state-of-the-art performance**.

