



Advanced Computer Vision

Week 03

Sep. 12, 2022

Seokju Lee

Image Debugging with Python

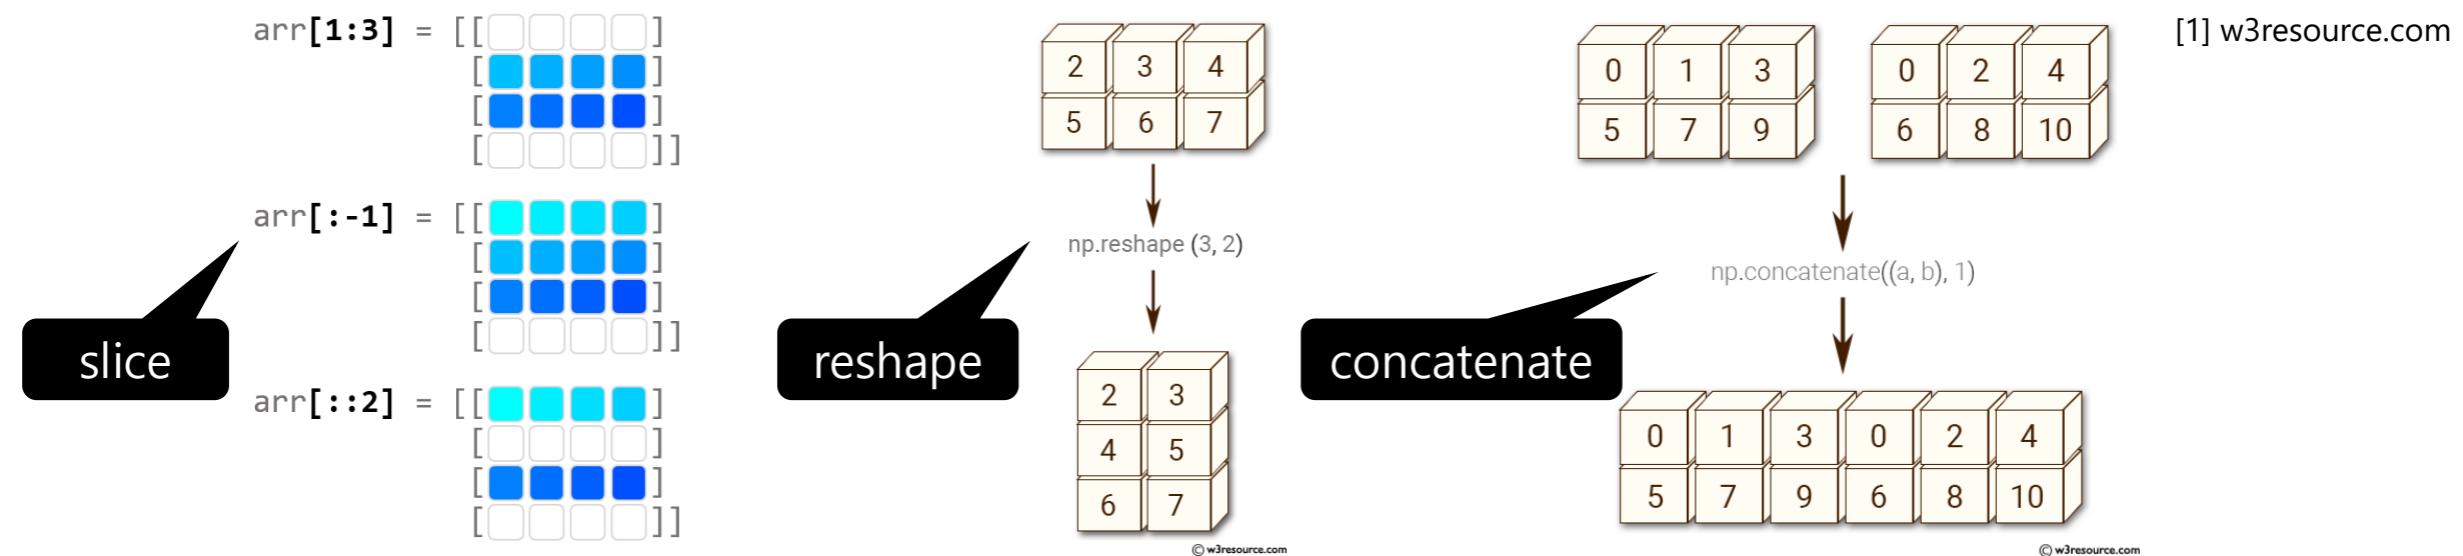
Summary of Previous Lesson

Basics of Python functions

- How to define and use

Basics of NumPy arrays

- Many useful operations for multi-dimensional arrays
- Please always check the current shape of arrays using the **shape** method!
- Basic math operations: add, subtract, min, max, mean, etc. along different axis
- Basic transformations: slicing, reshape, concatenate, stack, append, etc.



Now, Let's Play with Images!

Solve image processing puzzles.

- Multiple image processing missions to make specific images.
- If you need, please refer below pages.

Numpy: <https://numpy.org/doc/stable/reference/>

Scikit-learn image: <https://scikit-image.org/>

PIL image: <https://pillow.readthedocs.io/en/stable/reference/Image.html>

Codes are available at:

<https://view.kentech.ac.kr/a6fac1a2-7304-409e-85f0-78b1d527034f>

Installed packages
scikit-learn 1.0.2
4 dependencies
numpy
scipy
joblib
threadpoolctl
matplotlib 3.5.1
scikit-image 0.19.2
8 dependencies
numpy
scipy
networkx
pillow
imageio
tifffile
PyWavelets
packaging
numpy 1.22.3

Image Processing Puzzle

Now we will try visualization

→ Enables image debugging

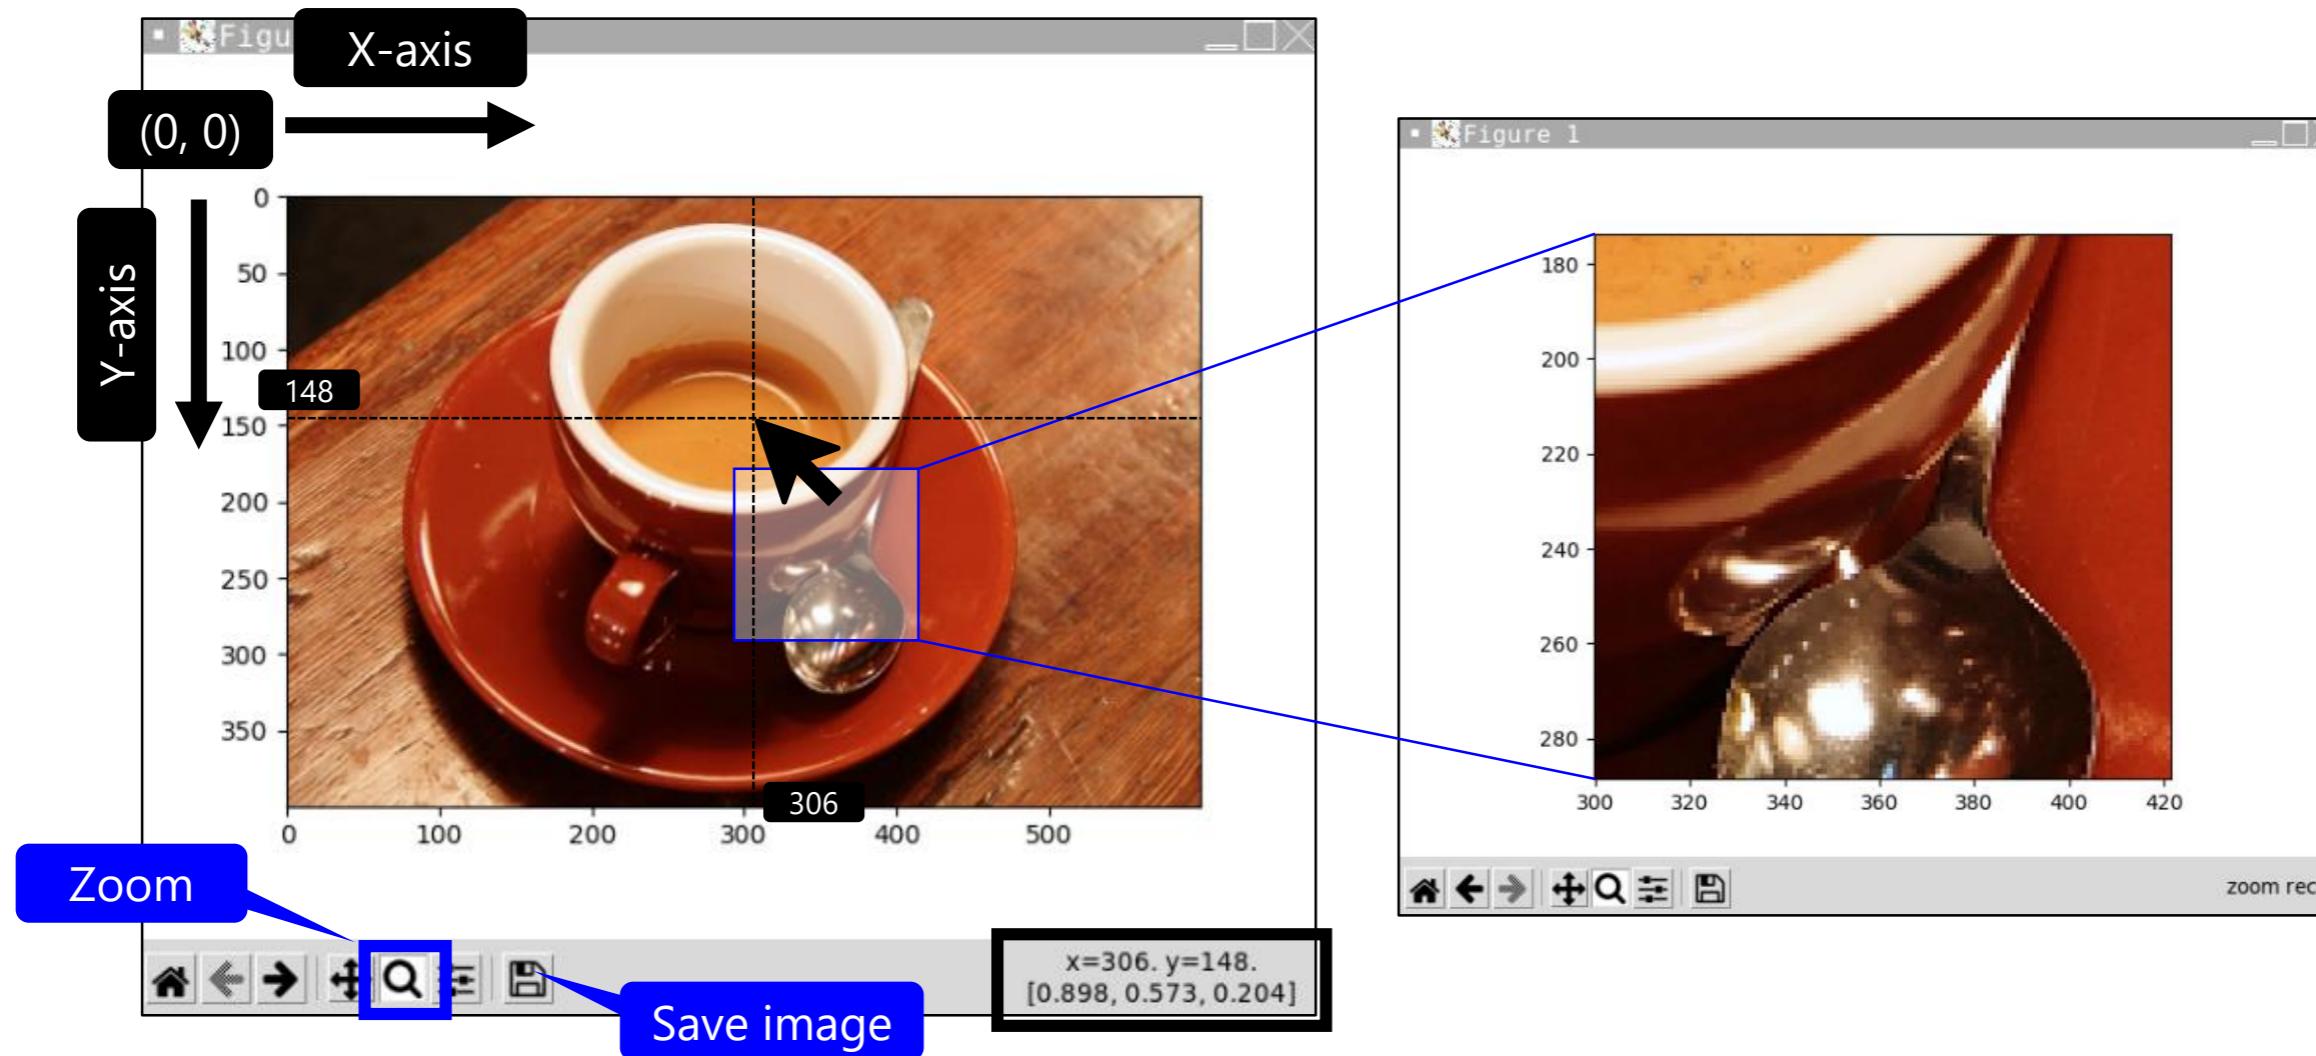


Image Processing Puzzle

Discussion 2: shifting image over x-axis

→ Use loop to generate the target image

```
hh, ww, ch = im.shape
```

Check the shape

```
for i in range(ww):
```

Search space: maximum width

```
    print('index:', i)
```

```
    im1 = np.zeros(im.shape)
```

Initialize

```
    im1[:, i:, :] = im[:, :ww-i, :]
```

Assign

```
    if (im1 == gt1).all():
```

Break the loop if the target is matched

```
        break;
```

```
check_same(gt1, im1)
```

Check if the output is correct

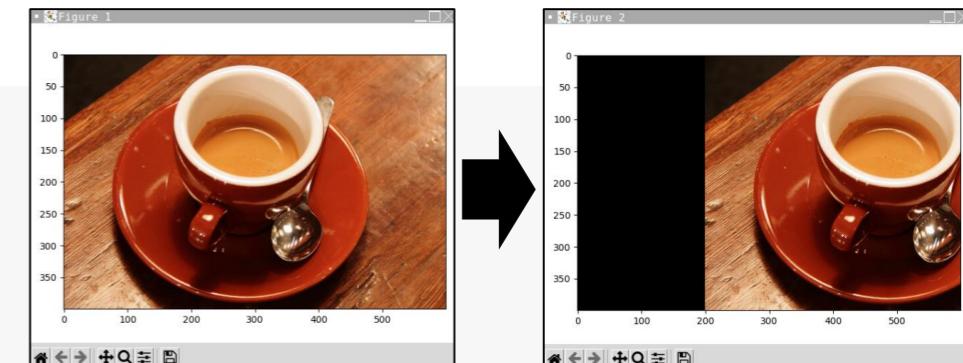


Image Processing Puzzle

Discussion 2-2: shifting over both x- and y-axis

→ Use loop to generate the target image → Too slow!

```
flag = False
hh, ww, ch = im.shape
for i in range(hh):
    for j in range(ww):
        print('index:', i, j)
        im1 = np.zeros(im.shape)
        im1[i:, j:, :] = im[:hh-i, :ww-j, :]
        error = np.abs(im1 - gt1).mean()
        if error == 0:
            flag = True
            break;
        if flag:
            break;
check_same(gt1, im1)
```

Search space: both maximum height + width
Use "**nested loop**"

You can check the processing time using the below code lines:

```
import time
start = time.time()
...
print(time.time() - start)
```

Image Processing Puzzle

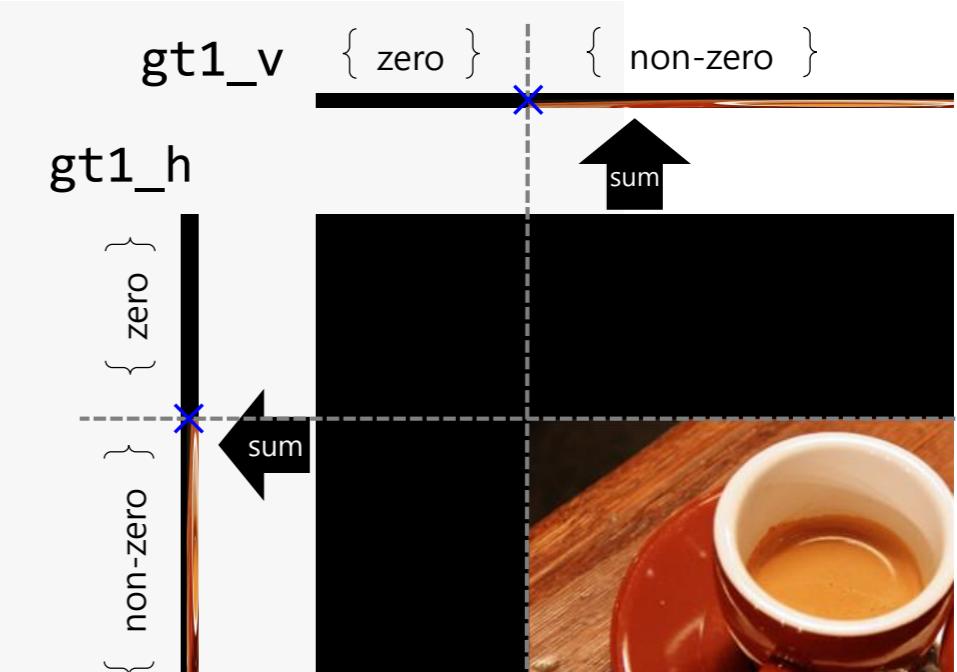
Discussion 2-2: shifting over both x- and y-axis

→ Dimension reduction. Please discuss this code in your report!

```
hh, ww, ch = im.shape  
  
gt1_h = gt1.sum(axis=(0,2))  
gt1_v = gt1.sum(axis=(1,2))  
  
jj = (gt1_h == 0).sum()  
ii = (gt1_v == 0).sum()  
  
im1[ii:, jj:, :] = im[:hh-ii, :ww-jj, :]  
  
check_same(gt1, im1)
```

Dimension reduction

Automatic binary sum
(= count the number of zeros)



→ Once you are familiar with the “for-loop” statement,
Try to implement your code **efficiently** through **dimension reduction!**

Image Processing Puzzle

Discussion 4: Image multiplication

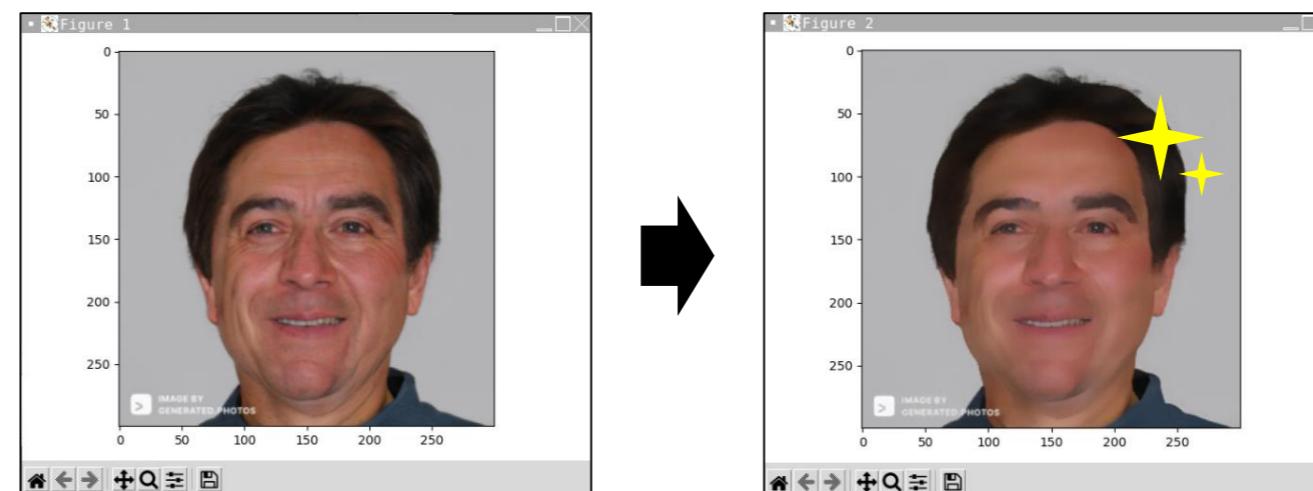
→ Multiply a binary mask to black out unnecessary parts.

Discussion 5: Image cropping and resizing

→ Crop a specific region with slicing, and resize the image.

Discussion 6: Photoshop - bilateral filter

→ Let's see how to do "photoshop" with a basic filtering algorithm.



Randomly generated face image (AI model) from
<https://generated.photos/face-generator/new>

Discussion

<Image Processing Puzzle 2>

Please paste your result images (screenshot or save) in this report if needed.

Discussion 2-3: Shifting image

2.7. Try yourself! Please try to use "dimension reduction".

2.8. Please visualize your output image.

Discussion 4 - Image multiplication

4.1. Please load a ground truth (GT) array from 'samples/puzzle2/gt1.npy'

4.2. Please convert 'im' to look like 'gt1' by multiplying the binary mask.

4.3. What is broadcasting?

4.4. What are the rules for allowing broadcasting?

Discussion 5: Image cropping and resizing

5.1. Crop the image (cat's eye) by array slicing.

5.2. Increase the size of the "cat's eye" image to 400 x 400. Please discuss the difference between "rescale" and "resize".

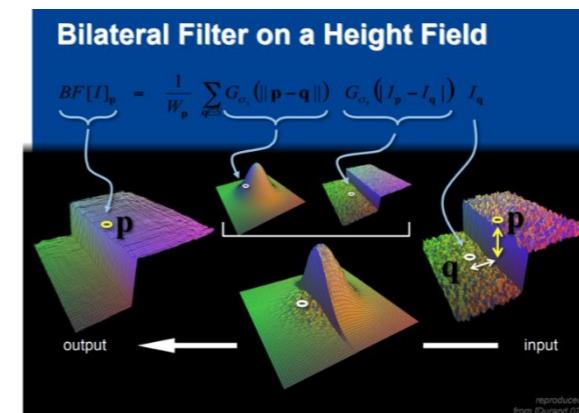
Discussion 6: Photoshop - bilateral filter

6.1. Please capture random faces from "<https://generated.photos/face-generator/new>"

6.2. Please apply a "bilateral filter" on the image.

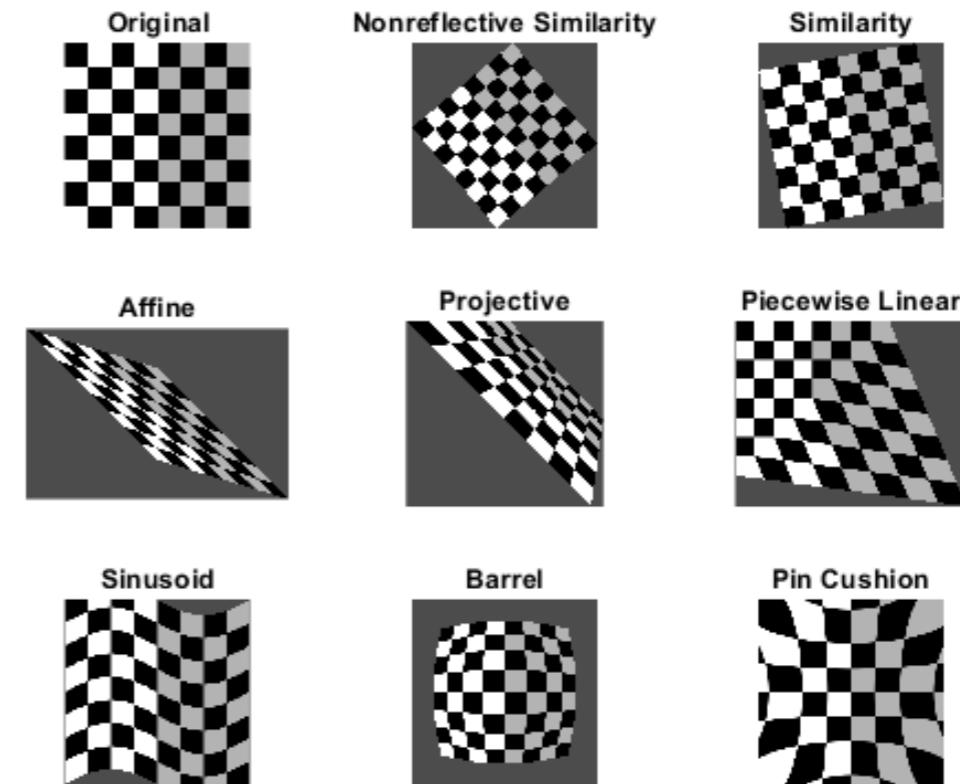
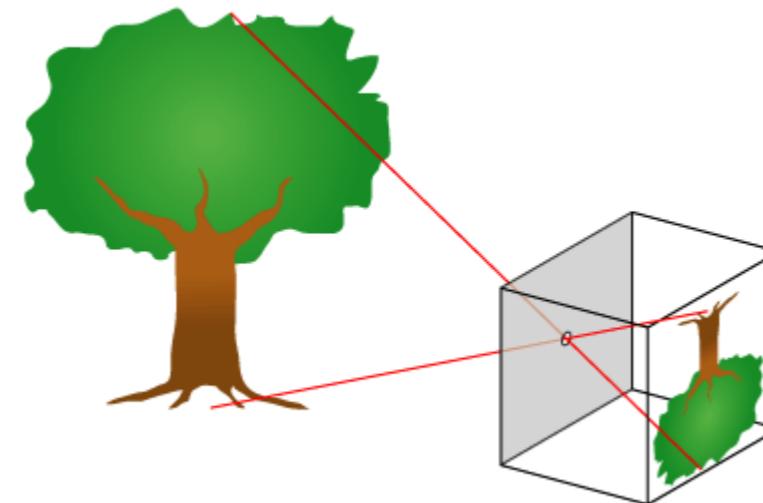
6.3. Please analyze the effect of the bilateral filter by changing the input parameters.

6.4. What are the pros and cons of filtering?



Next Contents

- Basic camera theory
 - Pinhole camera model
- Camera geometry

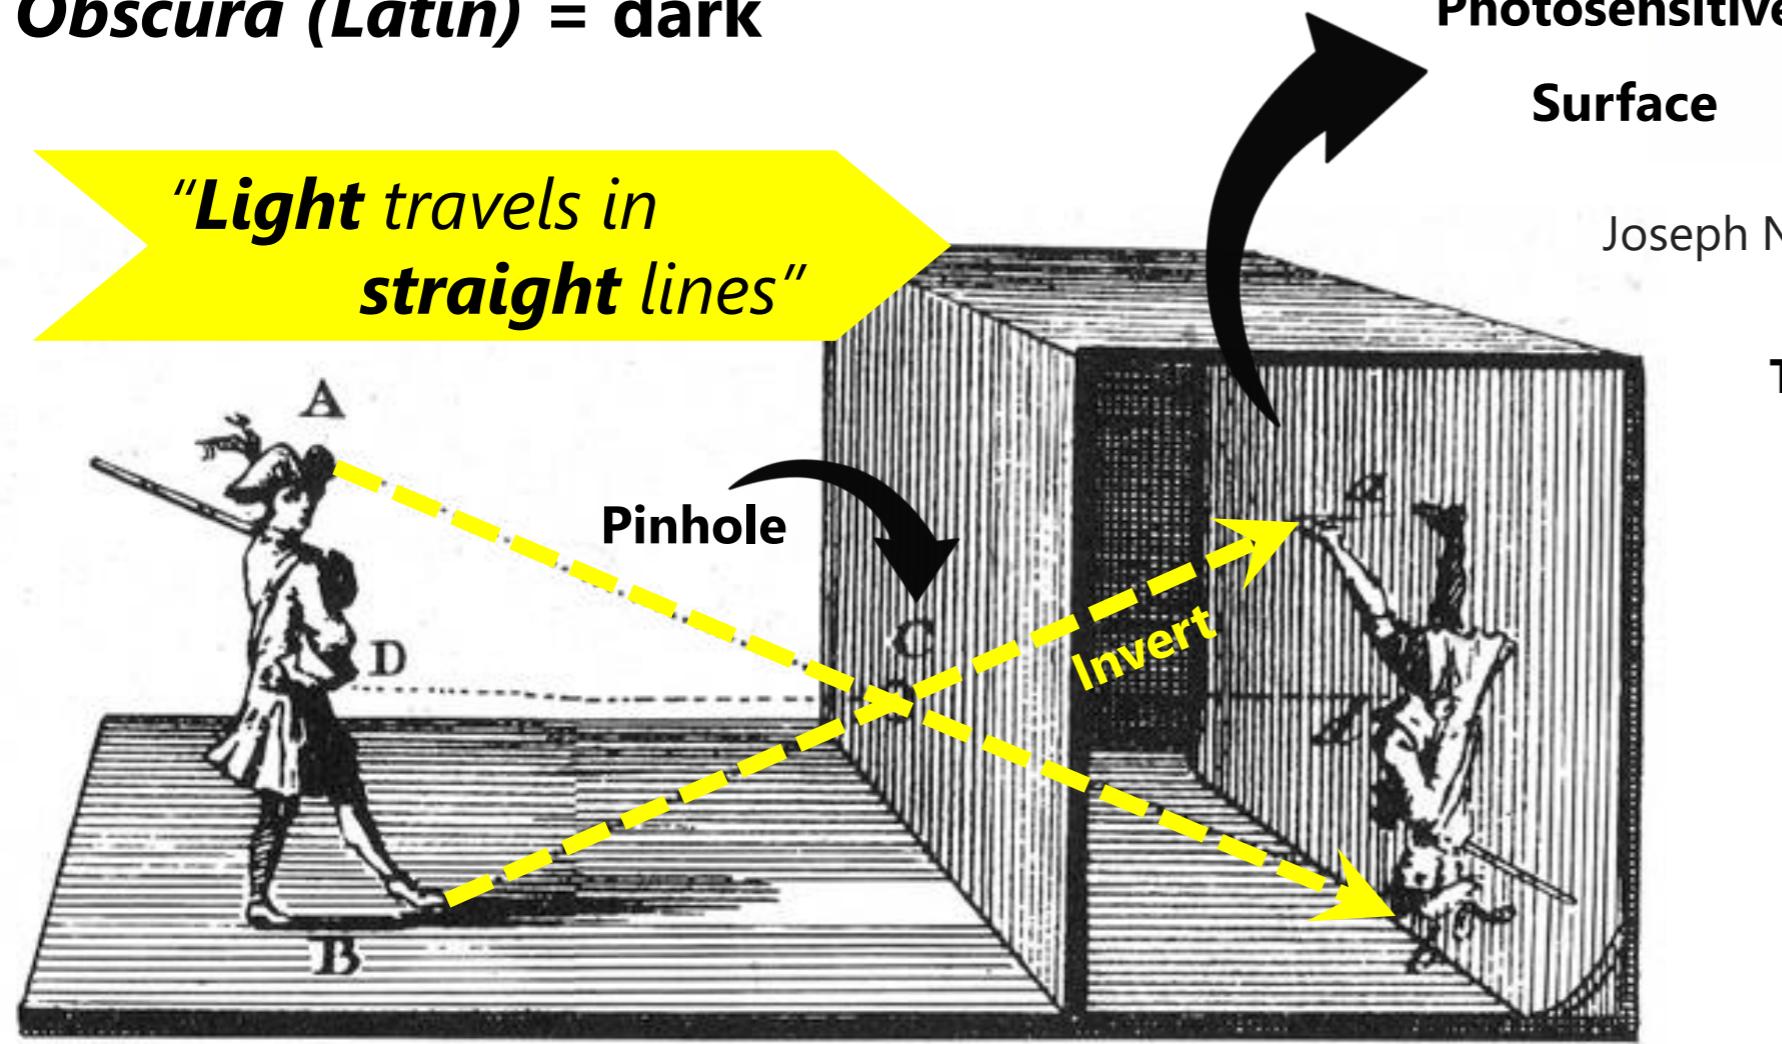


Brief Overview of Pinhole Camera Model

Camera Obscura – *Darken Room*

Camera (Latin) = room or chamber

Obscura (Latin) = dark



Joseph Nicéphore Niépce
(1765 – 1833)

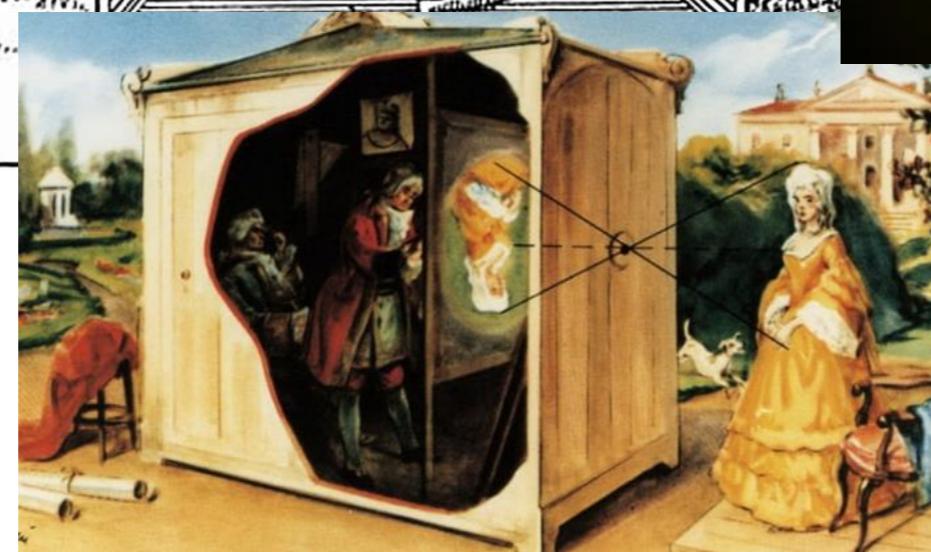
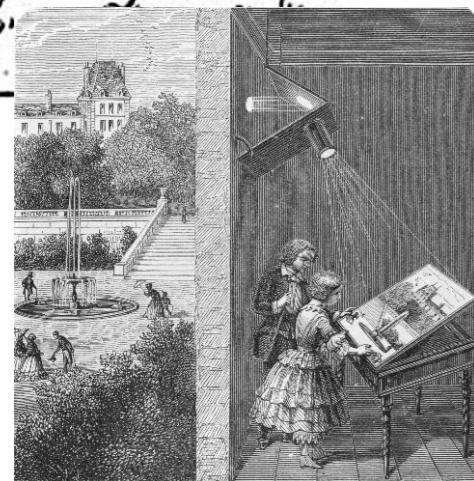
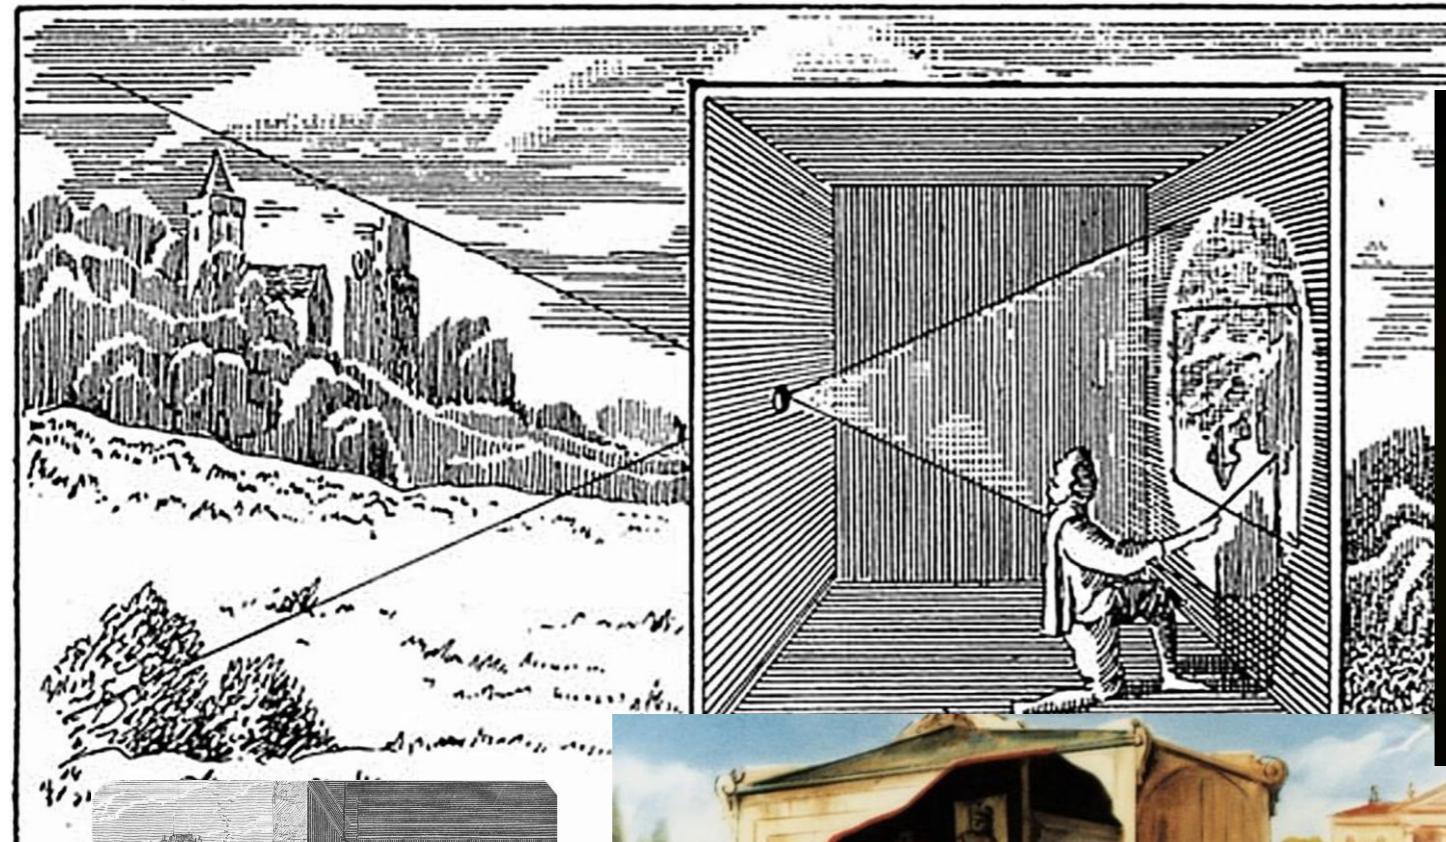


The inventor of photography



The first photograph
→ It took 8 hours

Camera Obscura – Based on Art



Amazing, cheap, fun way to experiment during COVID-19 quarantine



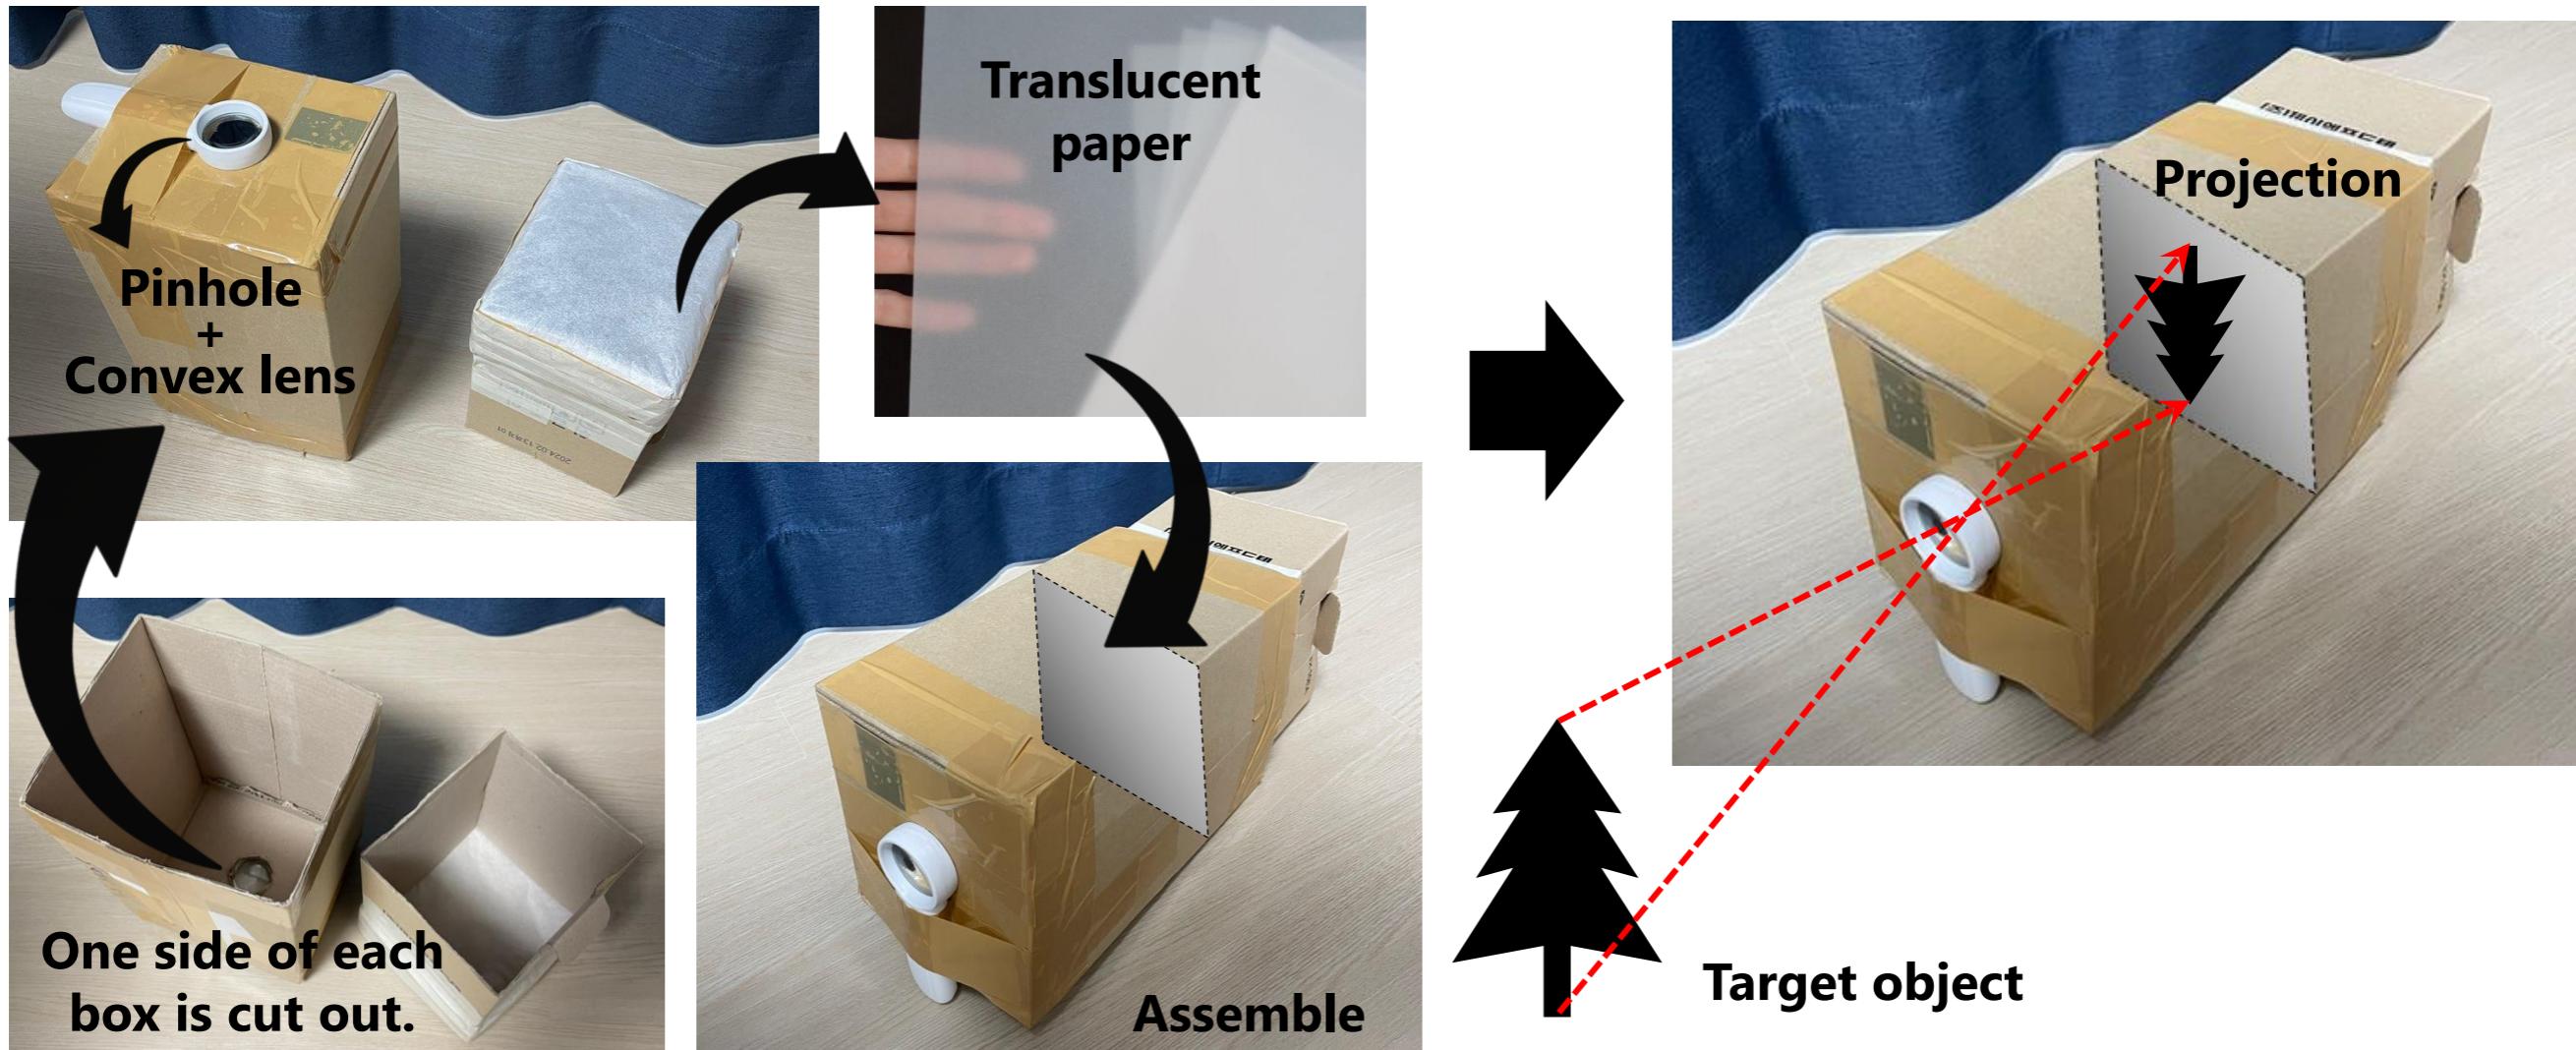
By Mathieu Stern

Inviting nature into your home!

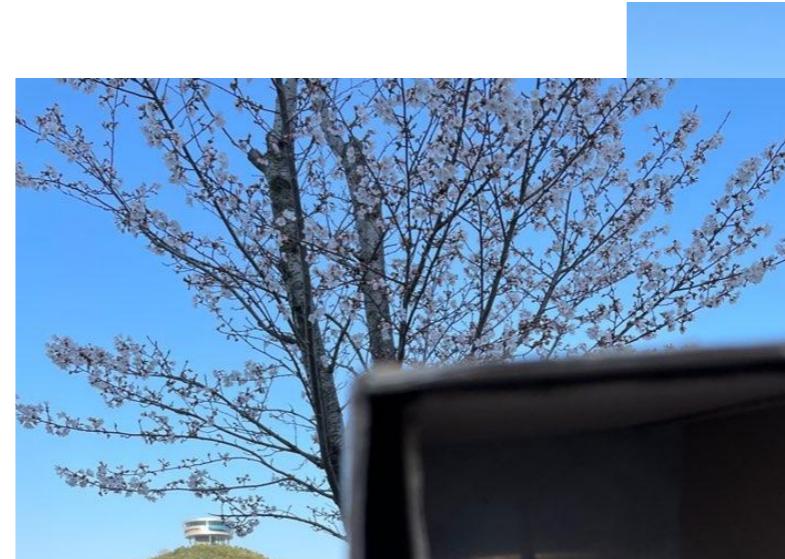
Agents of Change: Camera Obscura (Art Land Magazine)

Camera Obscura – Making

Preparation: two boxes, convex lens, translucent paper

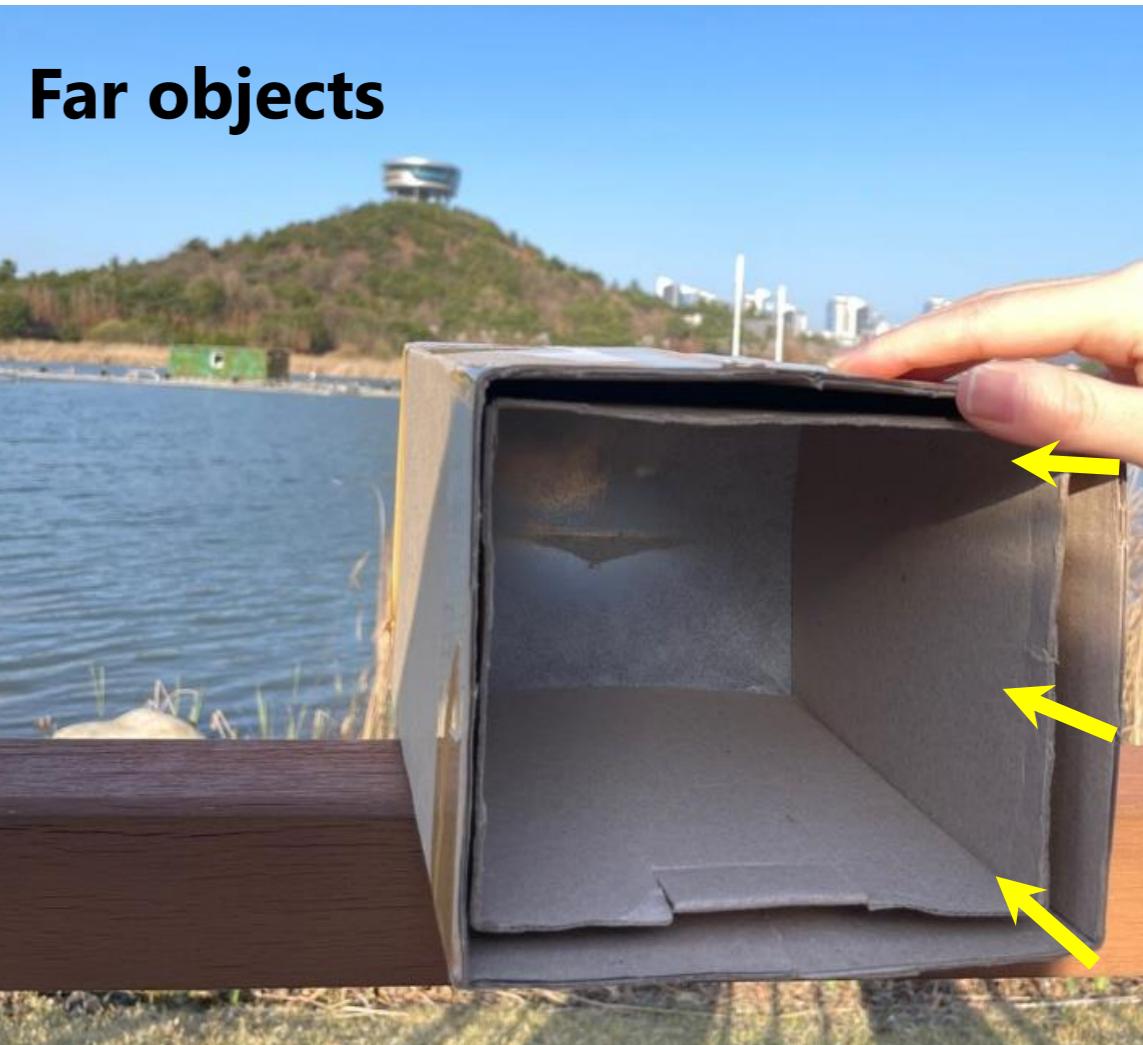


Camera Obscura – Experiments



Camera Obscura – About Focal Length

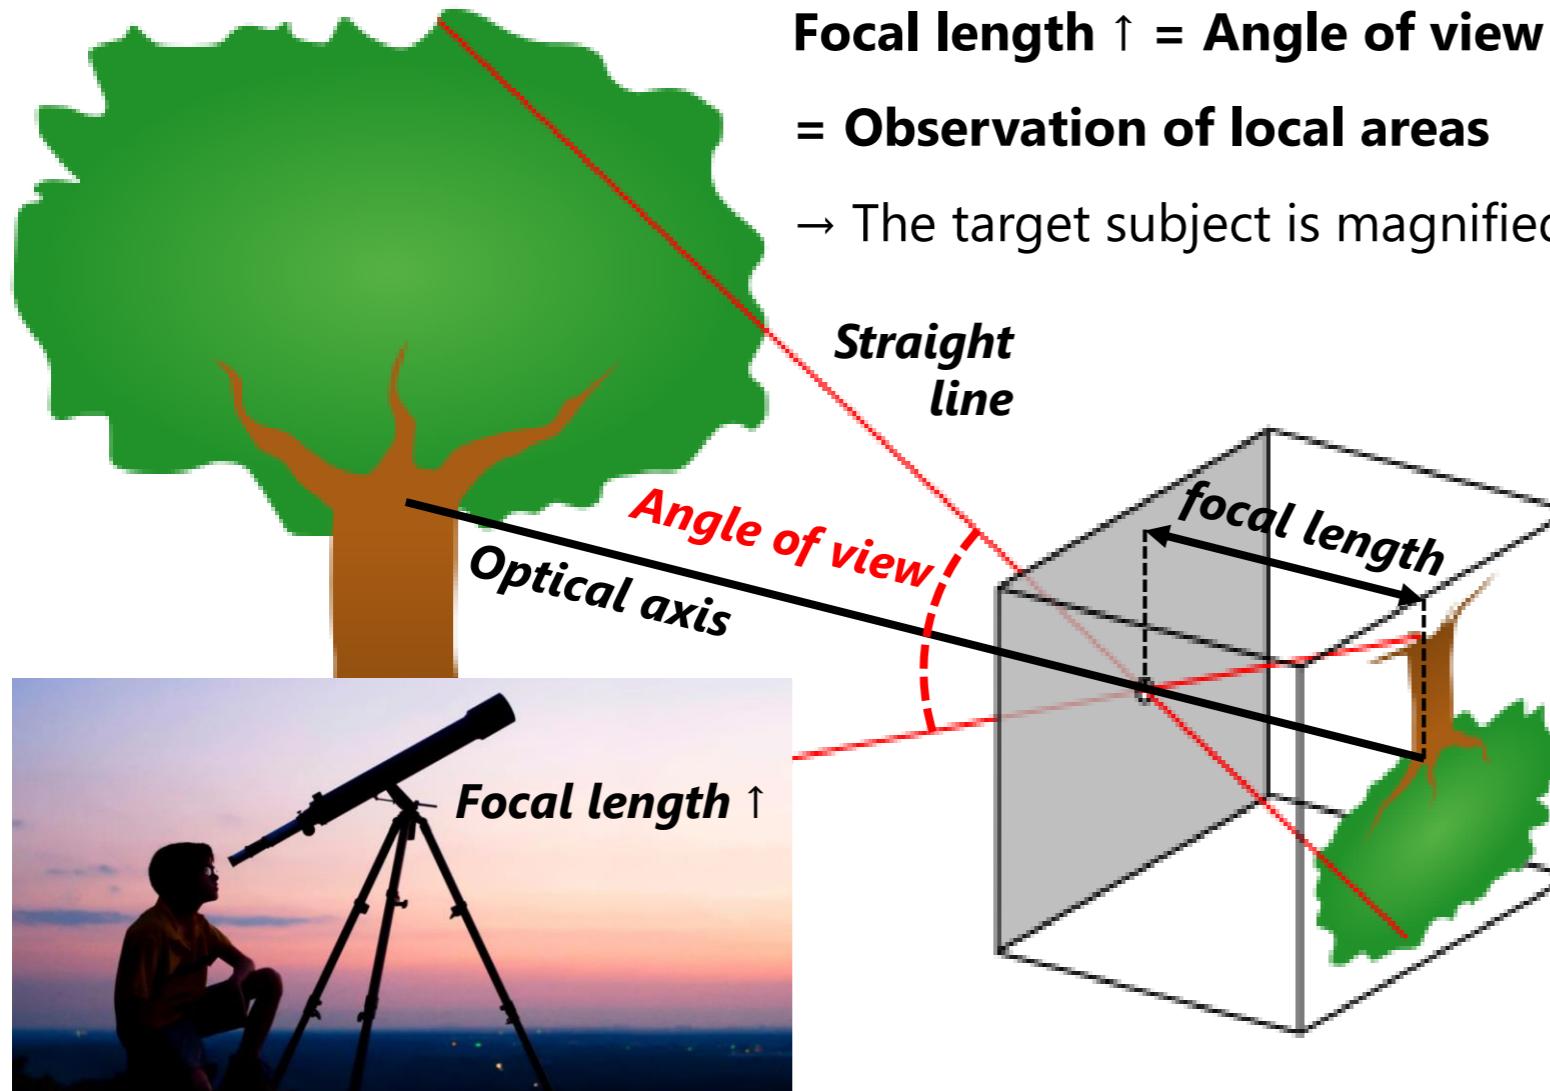
Near
objects



- **Near** objects: focal length ↑
- **Far** objects: focal length ↓

Pinhole Camera Model

Simplified example of ray optics



Focal length ↑ = Angle of view ↓
= Observation of local areas
→ The target subject is magnified!



We will make it!

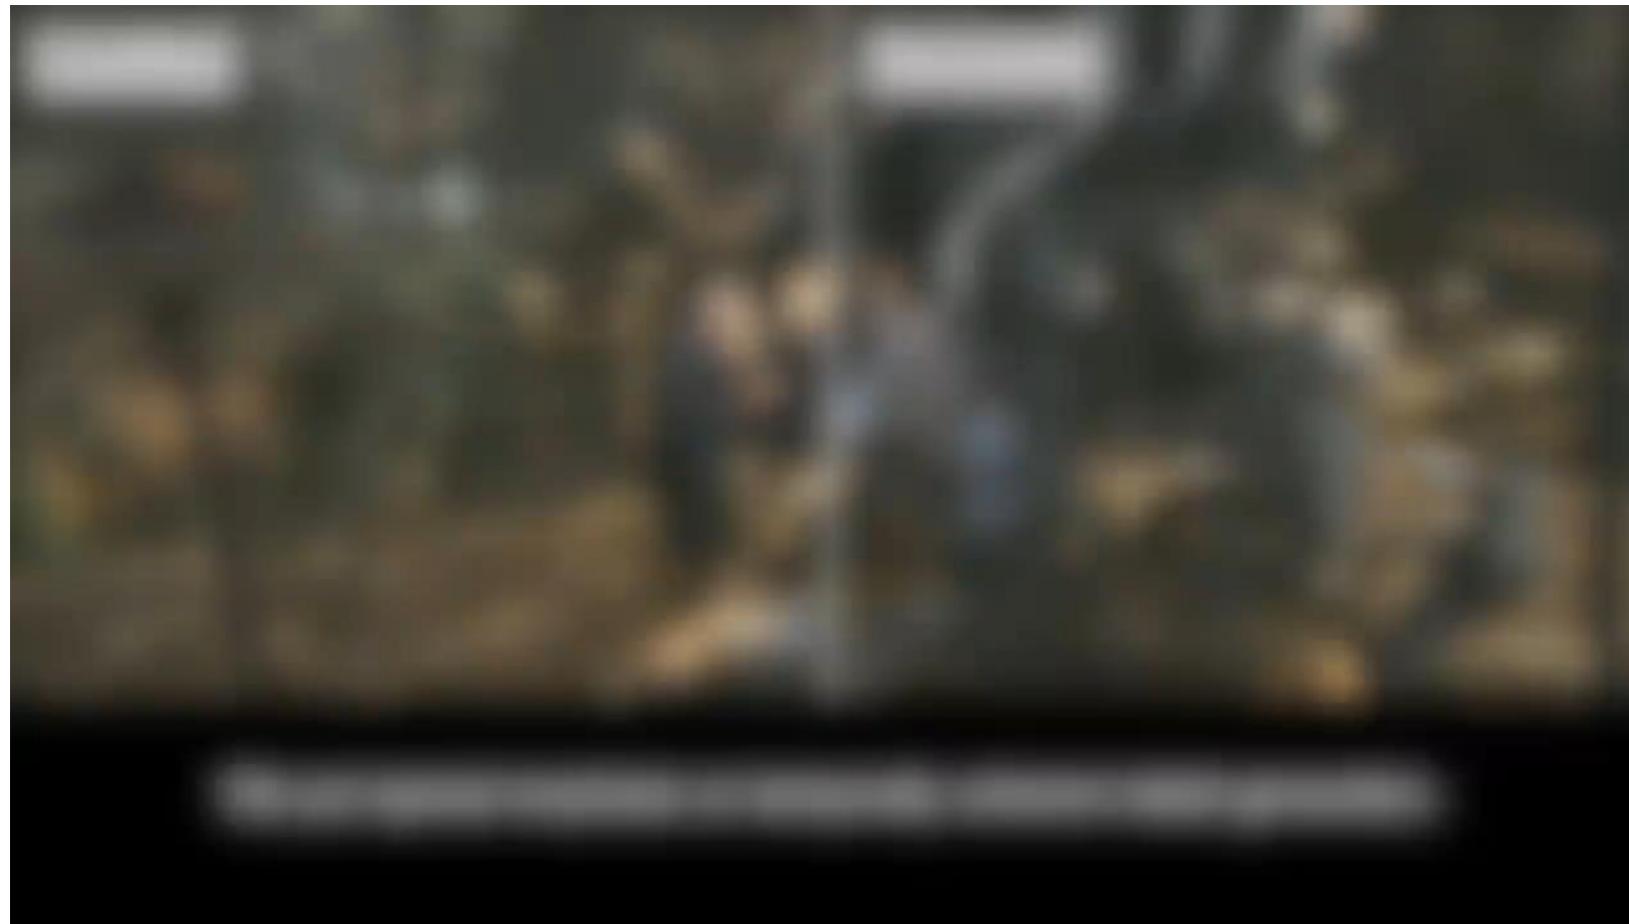
- ***Resolution** (for image, "spatial")
 - The number of pixels in each dimension
- ***Frame rate** (for video, "temporal")
 - The number of image frames per second

Screen = Photographic film = CCD sensor

About Resolution

Example of deep learning application:

→ Image super resolution, “**spatial** processing”



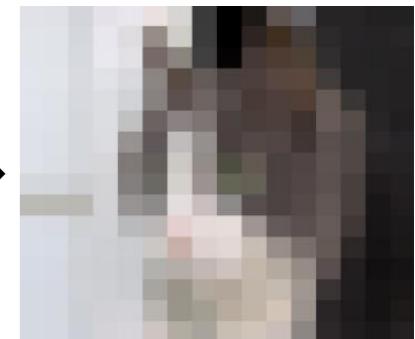
TecoGAN (SIGGRAPH'19)

How to train?

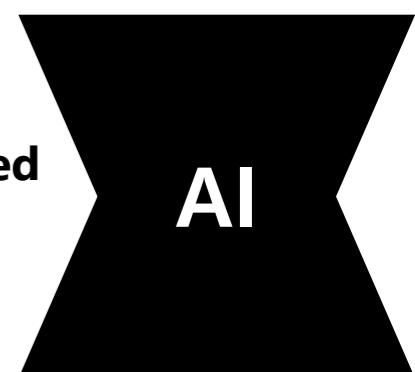
Original image



Resize!



Self-supervised
learning!



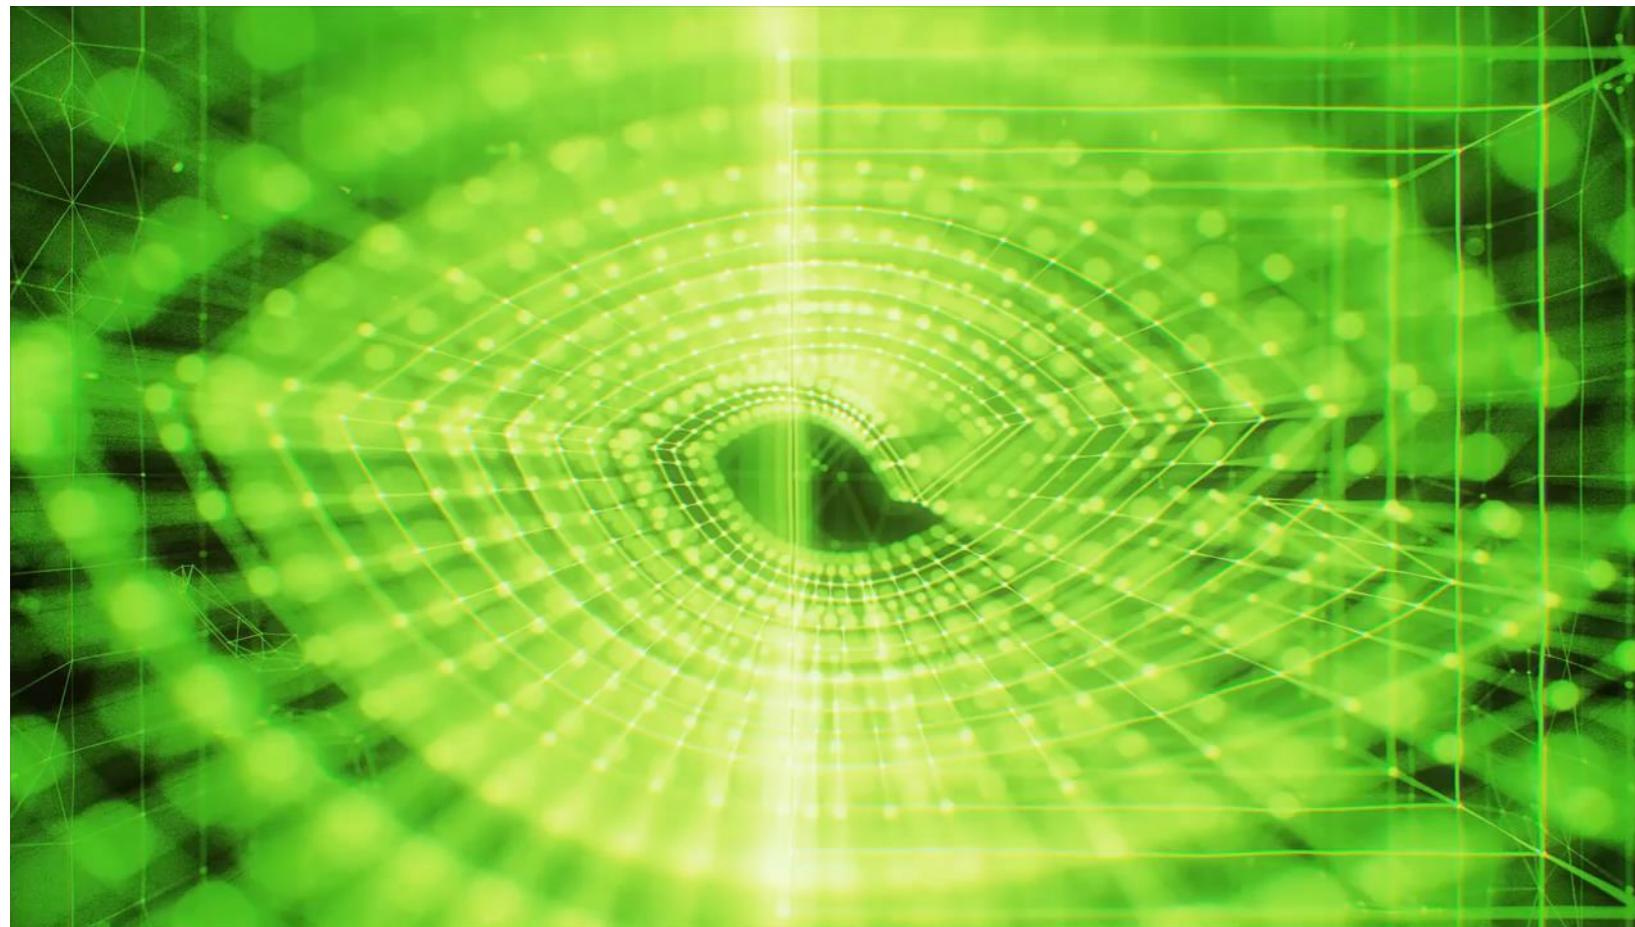
Compare!



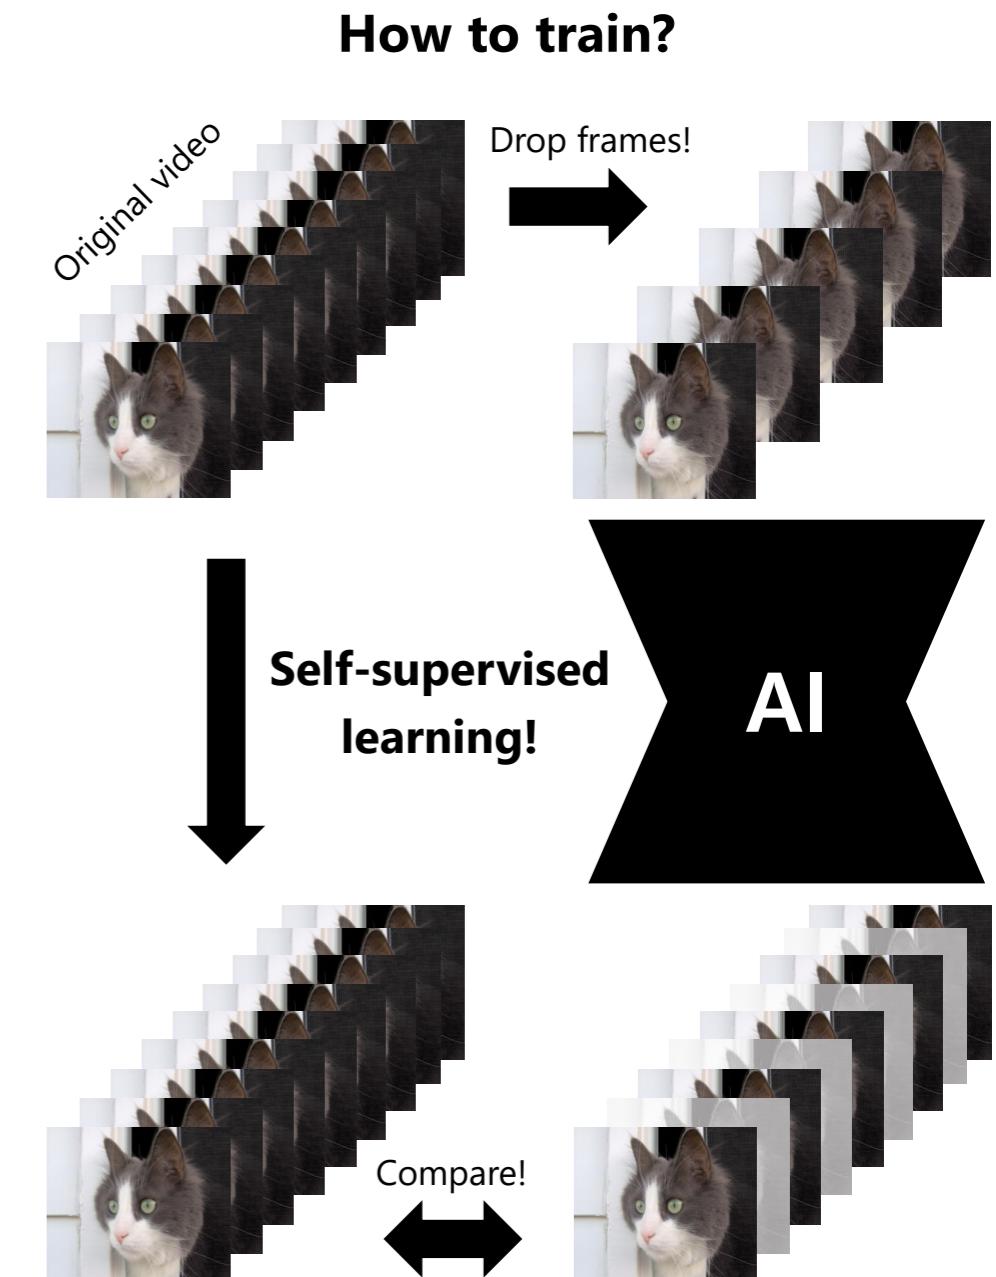
About Frame Rate

Example of deep learning application:

→ Video frame interpolation, "temporal processing"



Transforming Standard Video Into Slow Motion with AI
(Research at NVIDIA, 2018)



About Shutter Speed

A.k.a., exposure time:

The length of time that the digital sensor inside the camera is exposed to light



→ How shutter speed affects motion.

"The faster your shutter speed,
the sharper your image will be."

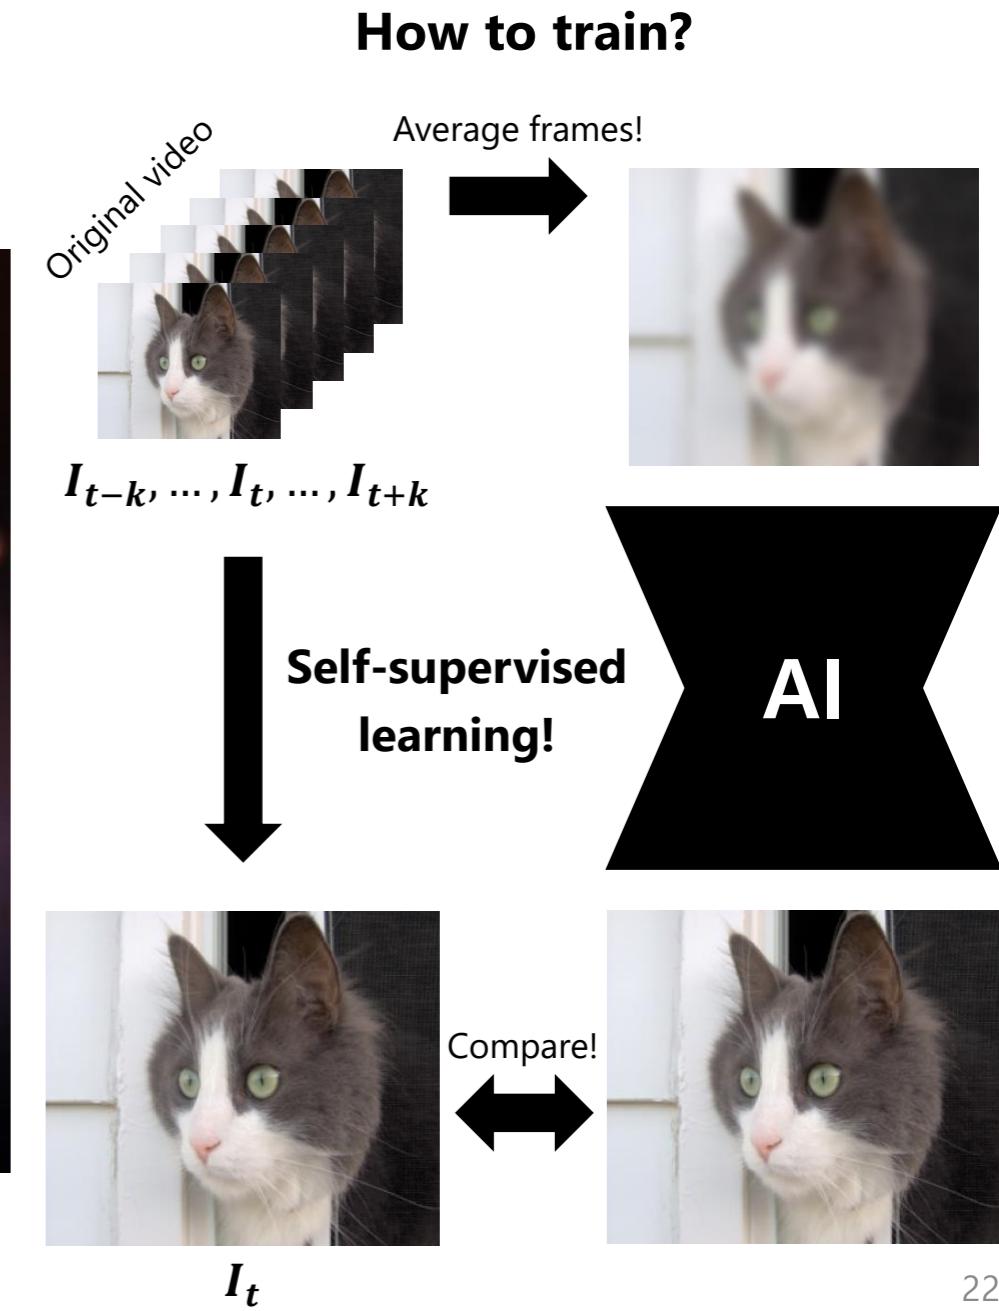
About Shutter Speed

Example of deep learning application:

→ Motion deblurring



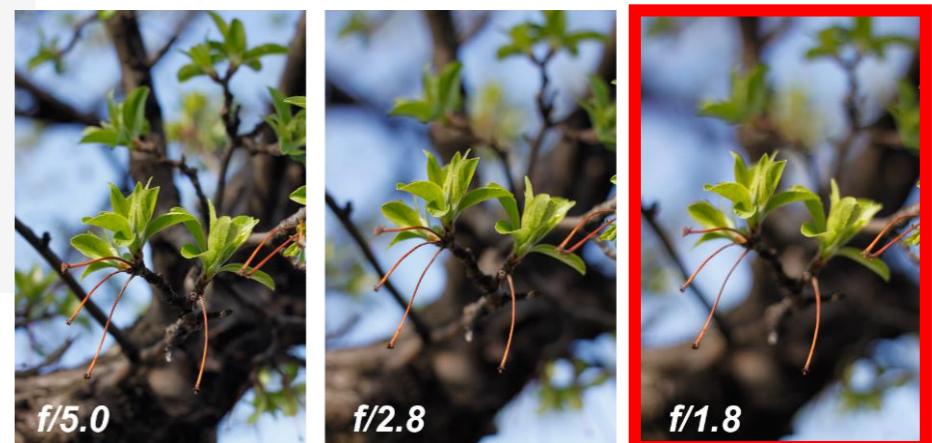
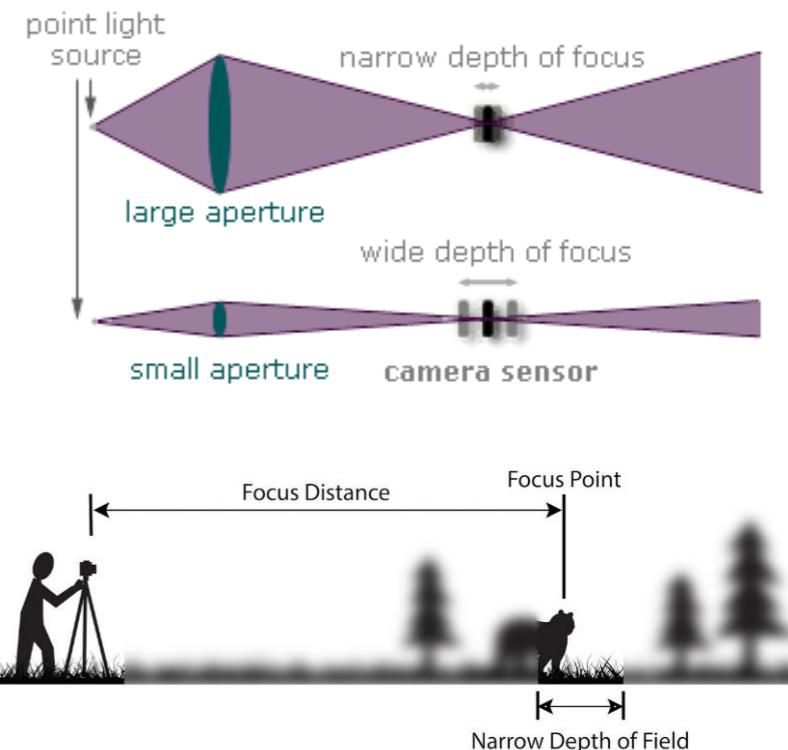
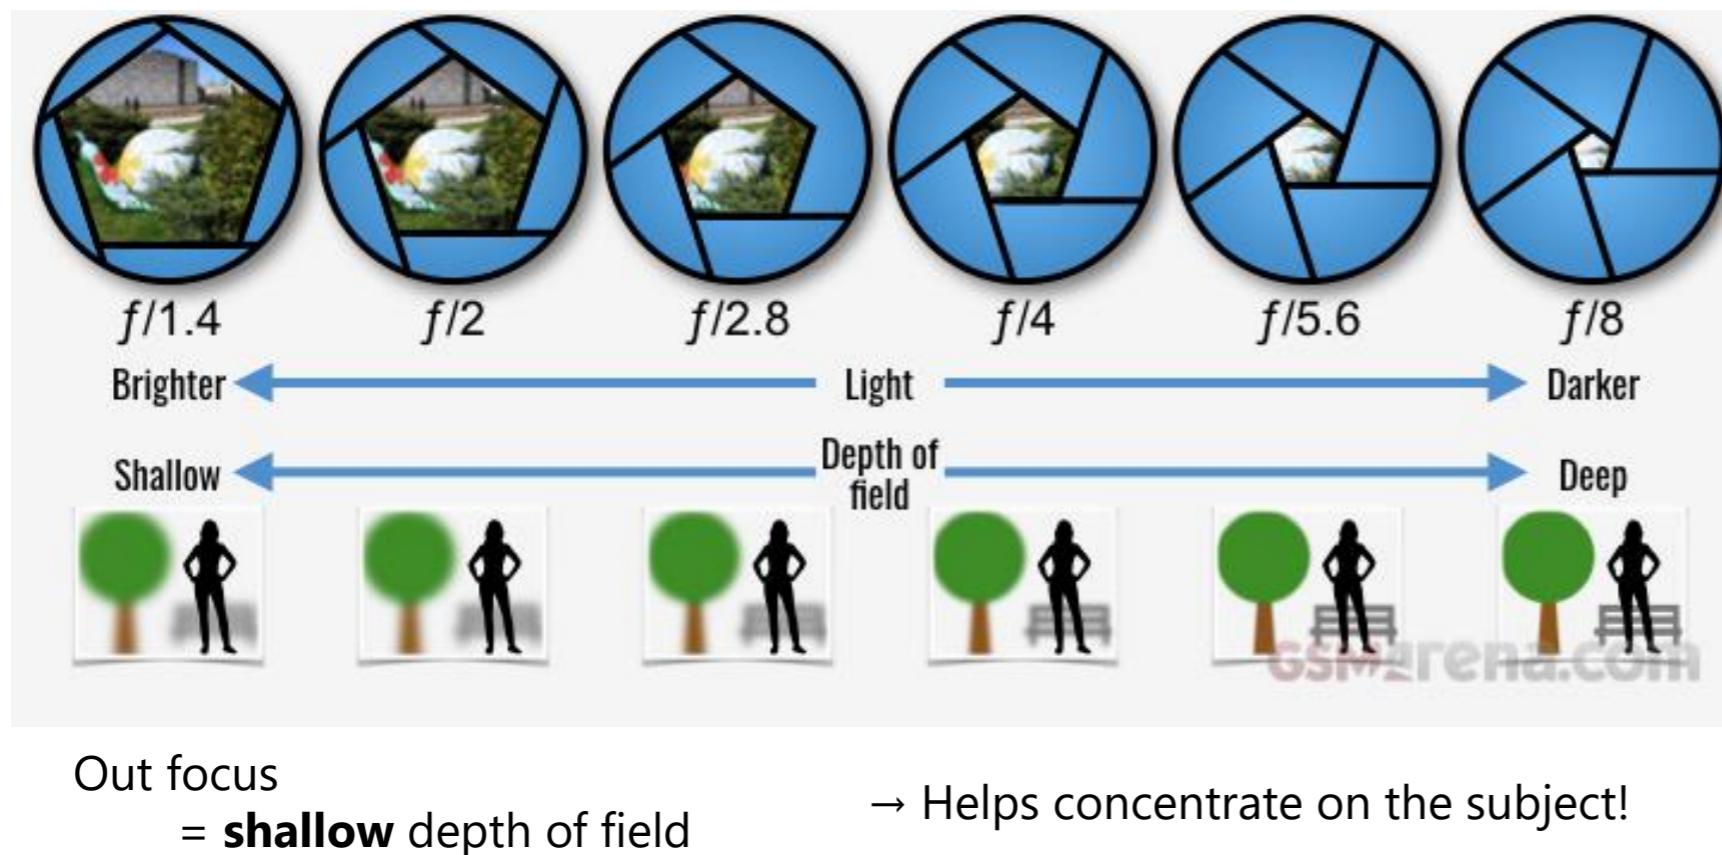
Fix Blurry Photos with Motion Deblur AI (by AKVIS, 2021)



Aperture

A hole through which light travels

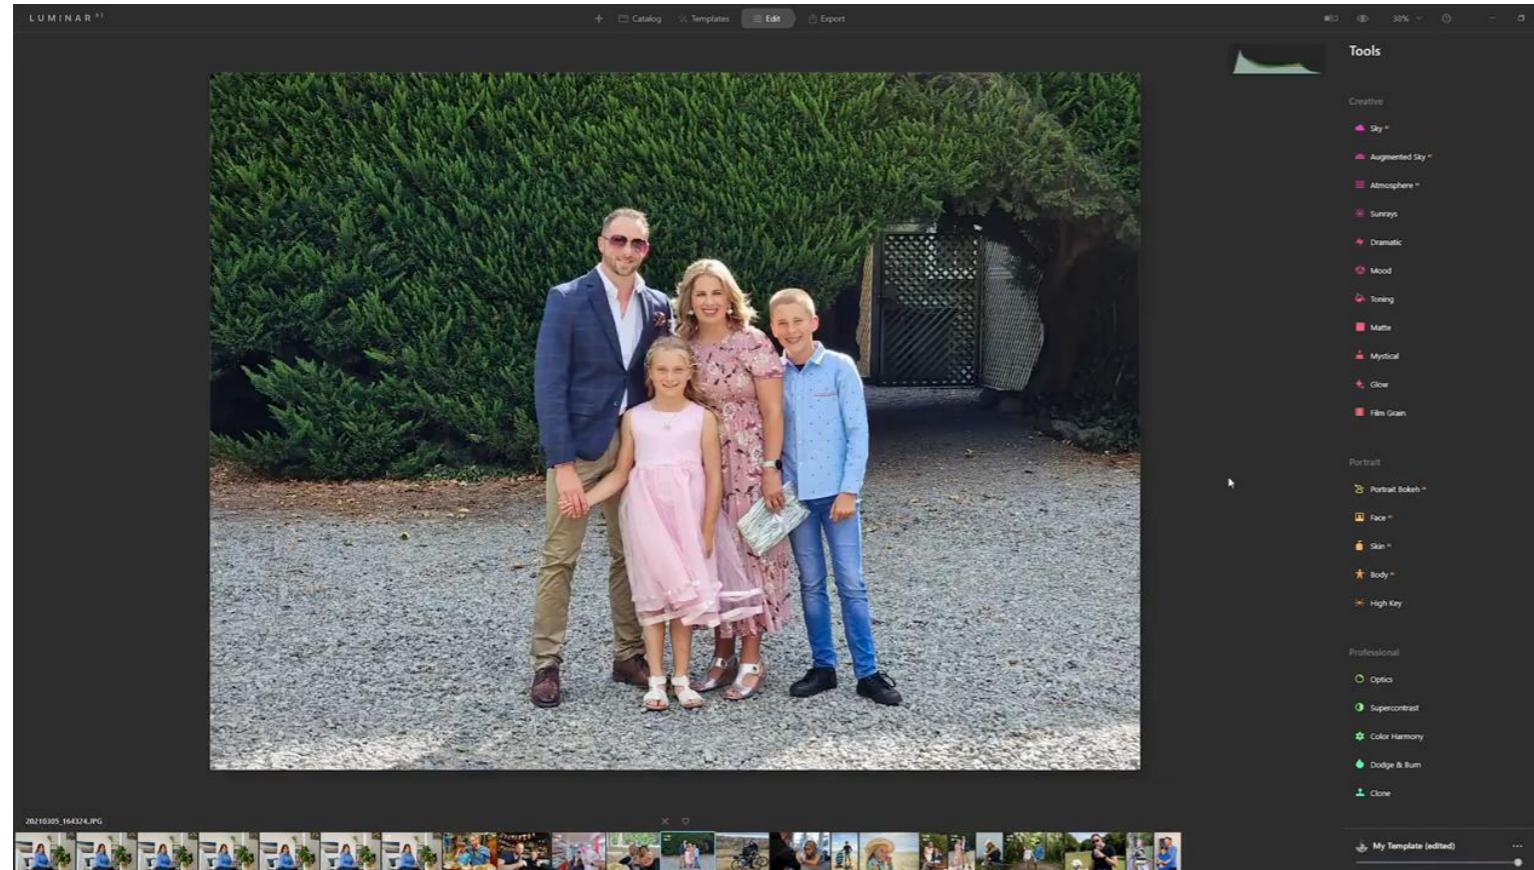
- Controls the amount of the light
- Controls depth-of-field (DoF)



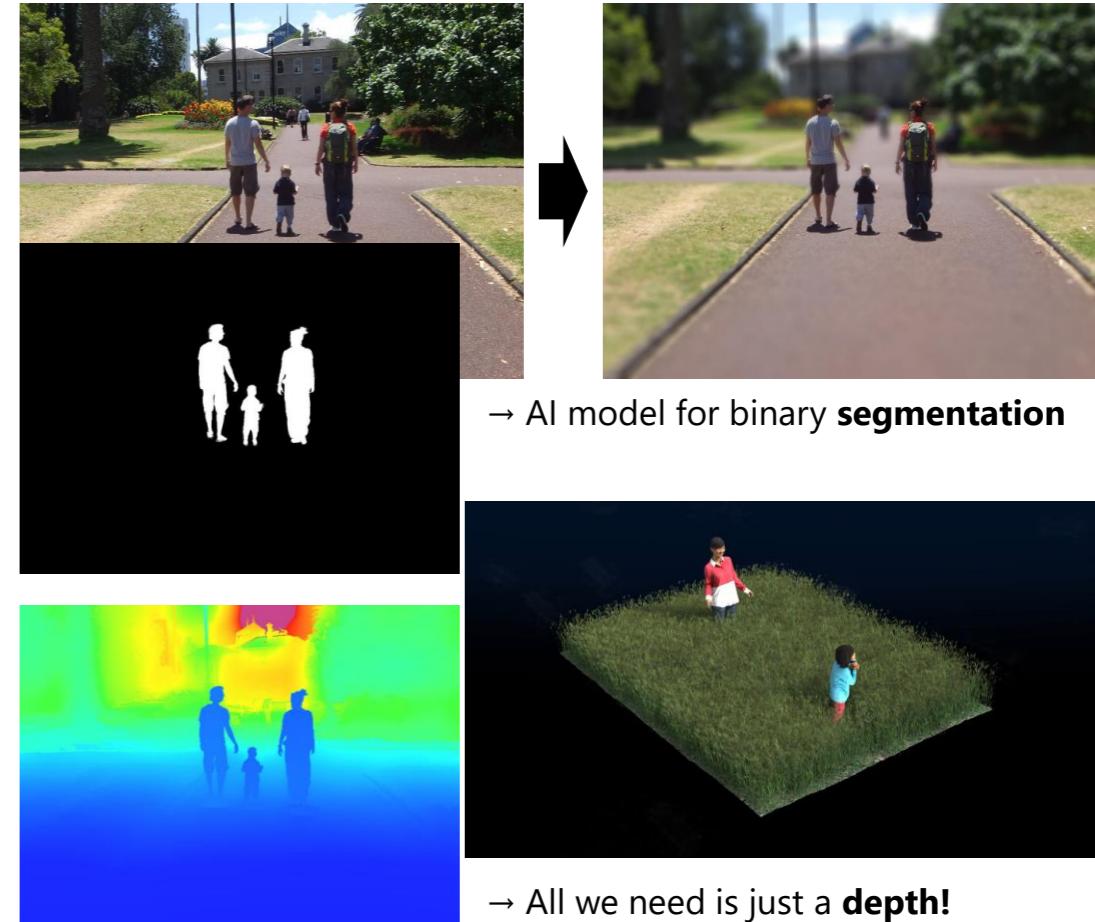
Out Focusing

Technically difficult for smartphone cameras 😱

→ Computer vision & AI make this possible!



Automatic bokeh by Luminar AI (2021)



[1] Liu, et al., “Bokeh Effects Based on Stereo Vision” (2015)

Summary of a Basic Camera Model

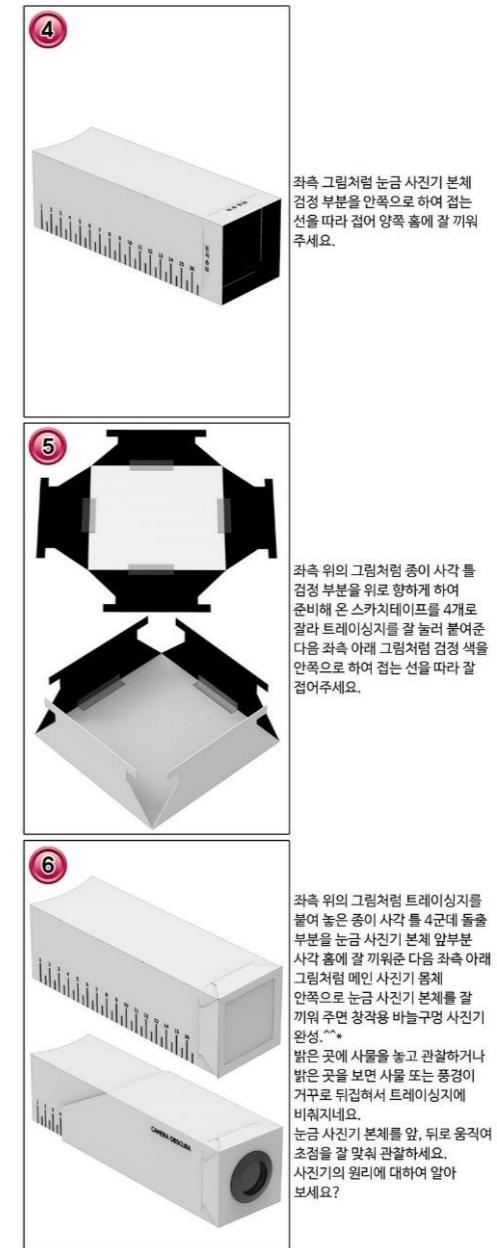
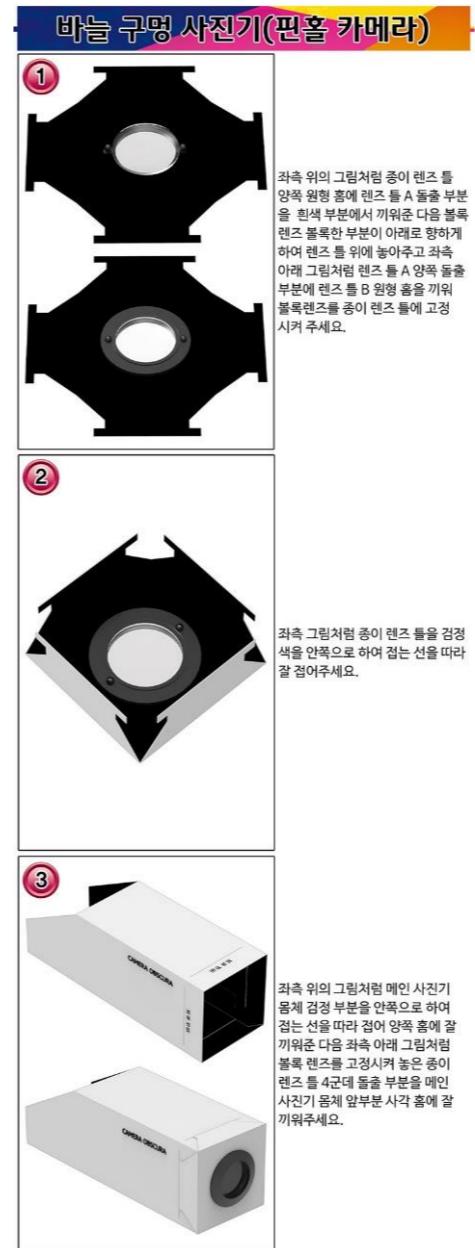
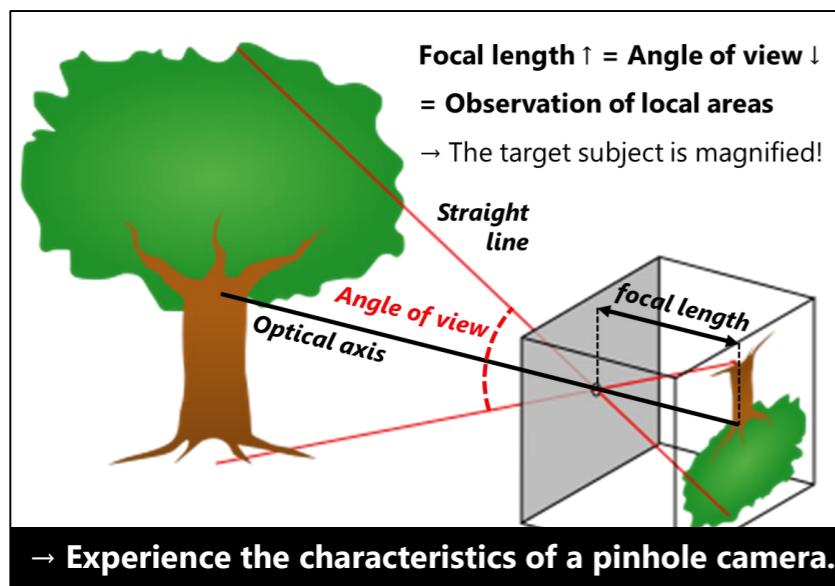
Characteristics of a pinhole camera

- Camera obscura
- Focal length
- Aperture & depth of field
- Shutter speed & exposure time

What AI can enhance images and videos

- Self-supervised learning
 - Image super resolution
 - Video frame interpolation
 - Motion deblurring
- Segmentation and depth
- 3D photography

Let's Make a Pinhole Camera



Discussion

<Pinhole Camera Quiz>

True or False?

- Q1. Decreasing the focal length makes the object smaller. [True / False]
- Q2. Increasing the focal length makes the field-of-view smaller. [True / False]
- Q3. Increasing the size of aperture makes the photo darker. [True / False]
- Q4. Increasing the size of aperture makes the depth-of-field shallower. [True / False]
- Q5. Decreasing the shutter speed makes the photo sharper. [True / False]
- Q6. Increasing the exposure time makes the photo brighter. [True / False]

Camera Calibration

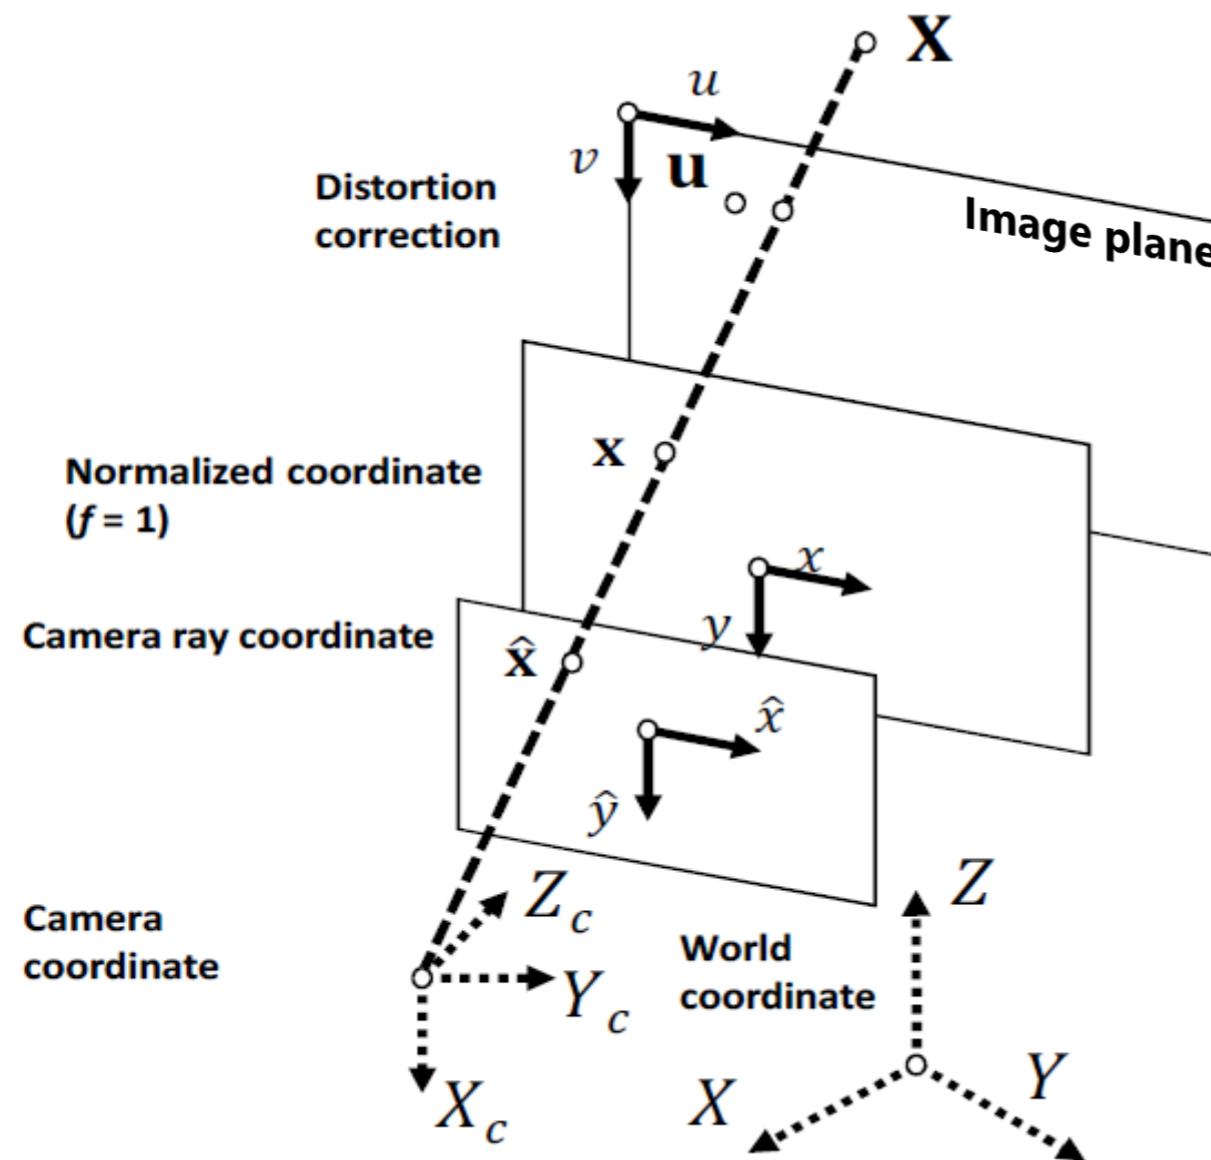
Camera Calibration?

- ✓ Process of estimating the **parameters** of a pinhole camera model.
- ✓ Determines which **incoming light (3D)** is **associated with** each pixel on the resulting **image (2D)**.
- ✓ In an ideal pinhole camera, **a simple projection matrix** is enough to do this.
- ✓ The camera projection matrix is derived from the **intrinsic** and **extrinsic** parameters of the camera.

$$\begin{aligned} \begin{bmatrix} \mathbf{x} \\ \mathbf{y} \\ 1 \end{bmatrix} &= \begin{bmatrix} f_x & \text{skew_cf}_x & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \mathbf{r}_{11} & \mathbf{r}_{12} & \mathbf{r}_{13} & \mathbf{t}_1 \\ \mathbf{r}_{21} & \mathbf{r}_{22} & \mathbf{r}_{23} & \mathbf{t}_2 \\ \mathbf{r}_{31} & \mathbf{r}_{32} & \mathbf{r}_{33} & \mathbf{t}_3 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \\ &= K[R|t] \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \end{aligned}$$

- (X, Y, Z) : 3D point in the world coordinate
- $[R|t]$: extrinsic parameters to convert the world coordinate into the camera coordinate
- K : intrinsic parameters to represent the camera characteristics
- $K[R|t]$: camera projection matrix
- (x, y) : 2D pixel location in the image plane
- s : scale factor

Camera Geometry: 3D → 2D



$$\mathbf{X}_c = \begin{bmatrix} R & T \\ 0^T & 1 \end{bmatrix} \mathbf{X}$$

$$(X_c, Y_c, Z_c)^T \mapsto (\hat{x}, \hat{y}) = \left(\frac{f X_c}{Z_c}, \frac{f Y_c}{Z_c} \right)^T$$

$$\hat{\mathbf{x}} = \begin{bmatrix} \hat{x} \\ \hat{y} \\ f \end{bmatrix} = \frac{1}{Z_c} \begin{bmatrix} f & 0 & X_c \\ 0 & f & Y_c \\ 0 & 0 & f \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

$$f = 1$$

$$\mathbf{x} = \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = [R \mid T] \mathbf{X}$$

$$\mathbf{u} = \mathbf{K} \mathbf{x}$$

$$\mathbf{u} = \mathbf{K}[\mathbf{R} \mid \mathbf{T}] \mathbf{X}$$

Intrinsic Parameters

- ✓ The intrinsic matrix K contains 5 intrinsic parameters of the specific camera model.
 - Focal length: f_x, f_y
 - Principal point: c_x, c_y
 - Skew coefficient (y-axis slope of the image sensor): $skew_c = \tan(\alpha)$

$$\begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & skew_cf_x & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \mathbf{r}_{11} & \mathbf{r}_{12} & \mathbf{r}_{13} & \mathbf{t}_1 \\ \mathbf{r}_{21} & \mathbf{r}_{22} & \mathbf{r}_{23} & \mathbf{t}_2 \\ \mathbf{r}_{31} & \mathbf{r}_{32} & \mathbf{r}_{33} & \mathbf{t}_3 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

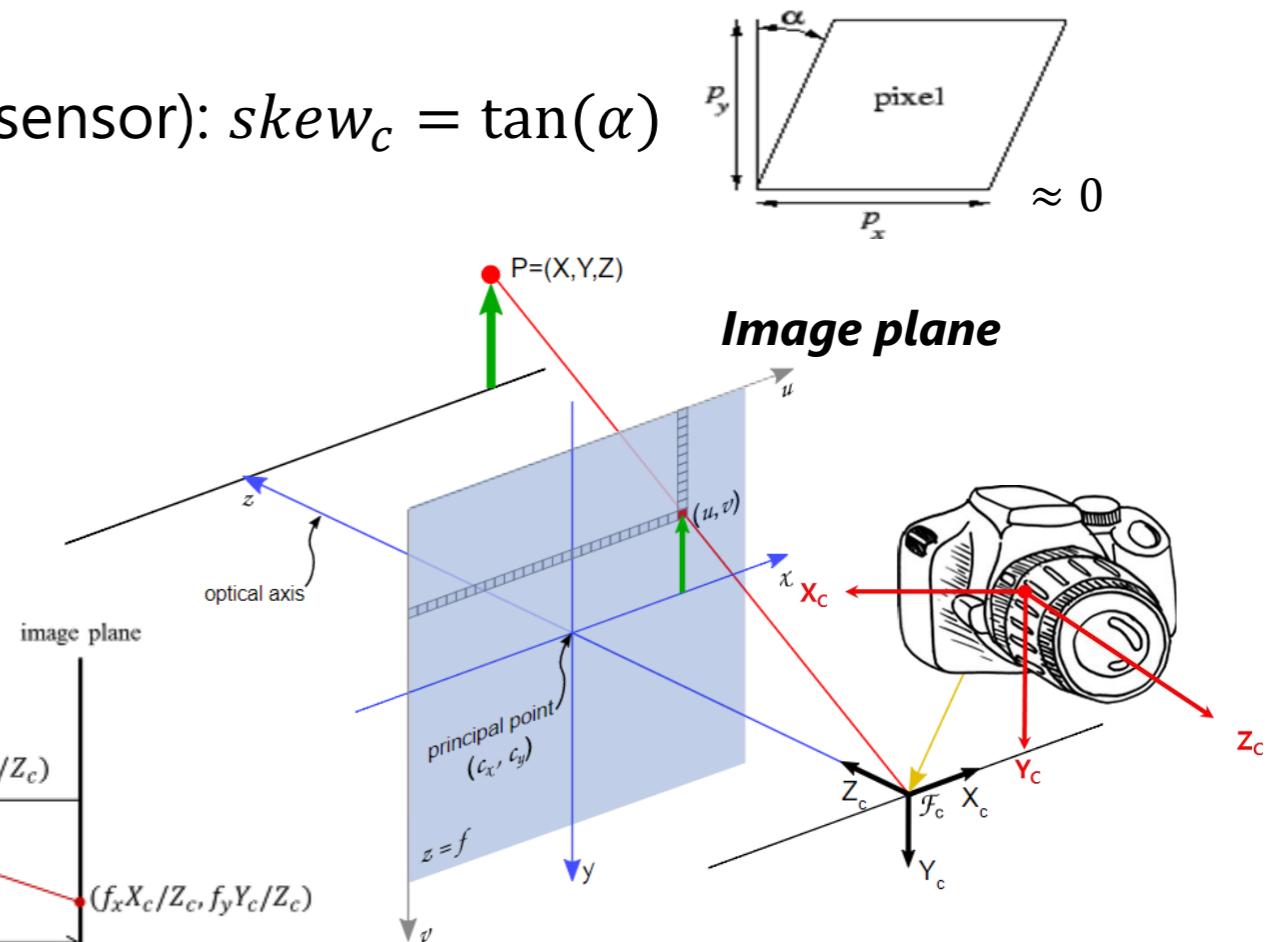
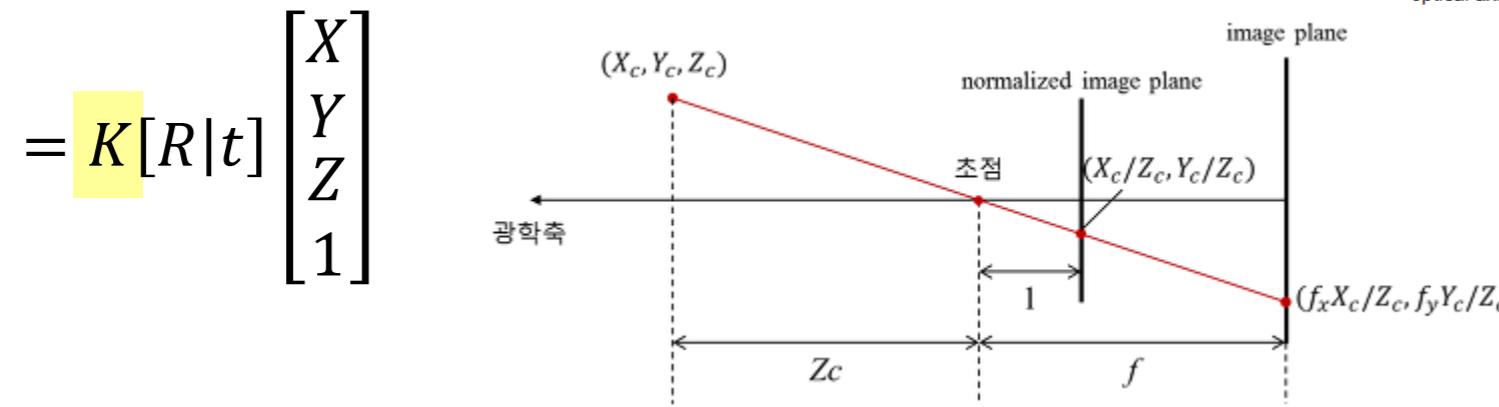


Figure by NVIDIA

Extrinsic Parameters

- ✓ Transformation matrix from 3D world coordinates to 3D camera coordinates
→ Rotation (R , 3×3 matrix) + Translation (t , 3×1 matrix)

$$\begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & \text{skew_cf}_x & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \mathbf{r}_{11} & \mathbf{r}_{12} & \mathbf{r}_{13} & \mathbf{t}_1 \\ \mathbf{r}_{21} & \mathbf{r}_{22} & \mathbf{r}_{23} & \mathbf{t}_2 \\ \mathbf{r}_{31} & \mathbf{r}_{32} & \mathbf{r}_{33} & \mathbf{t}_3 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

$$= K[R|t] \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

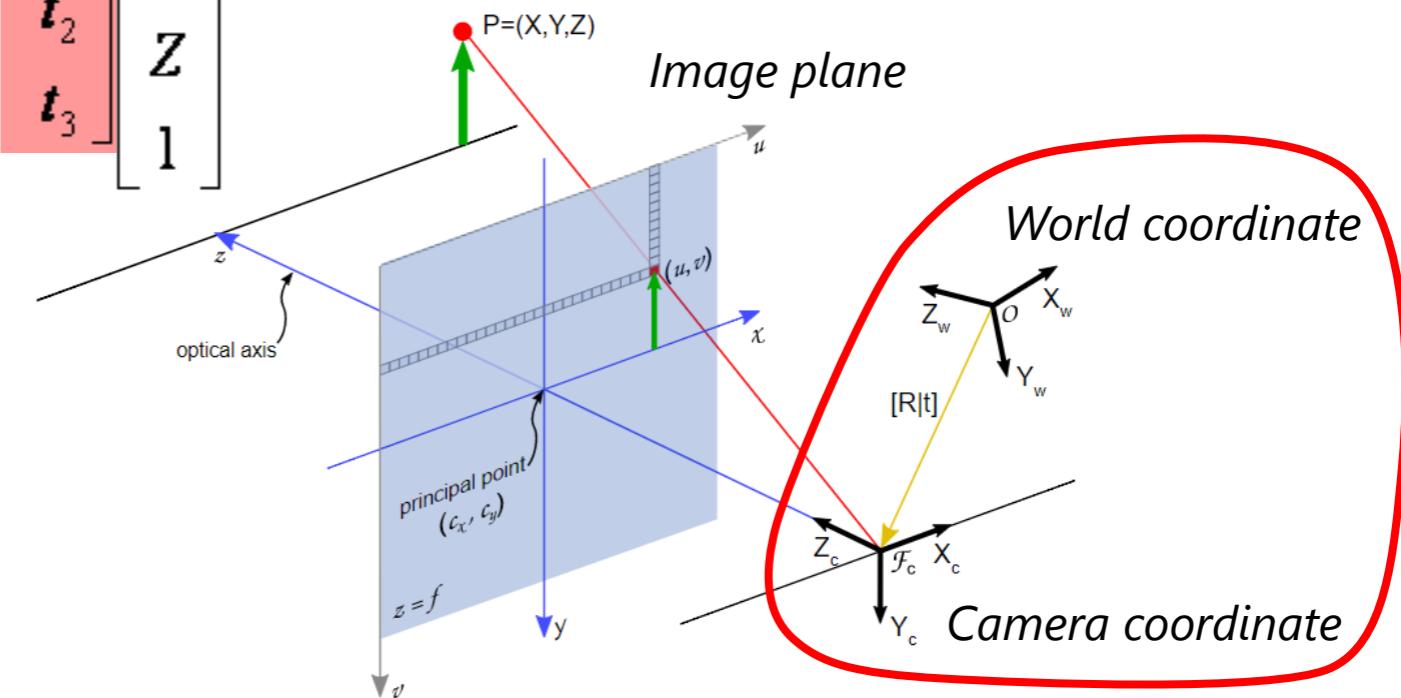
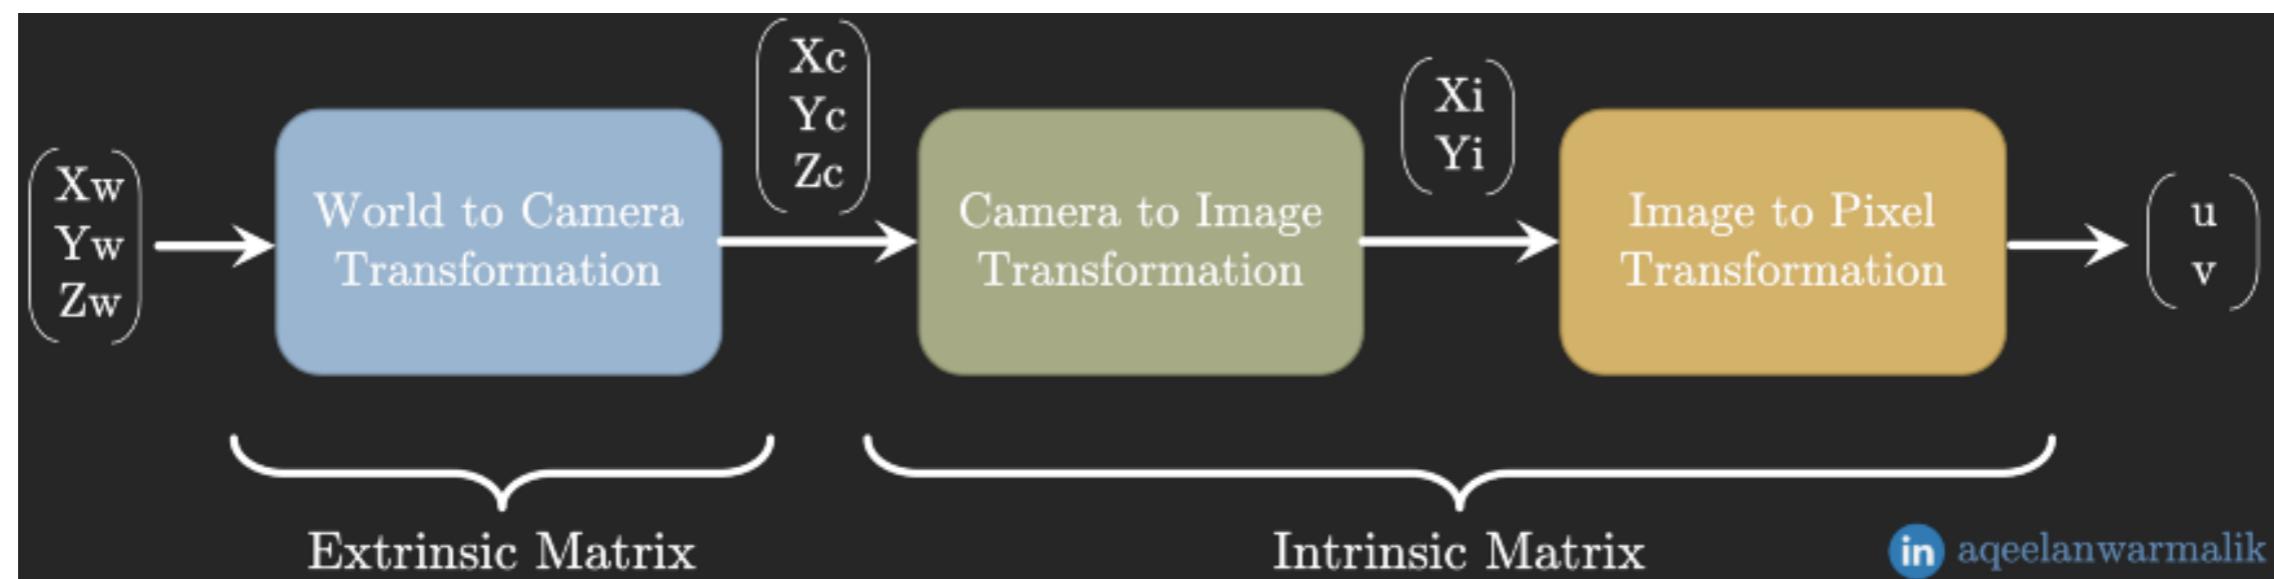


Figure by NVIDIA

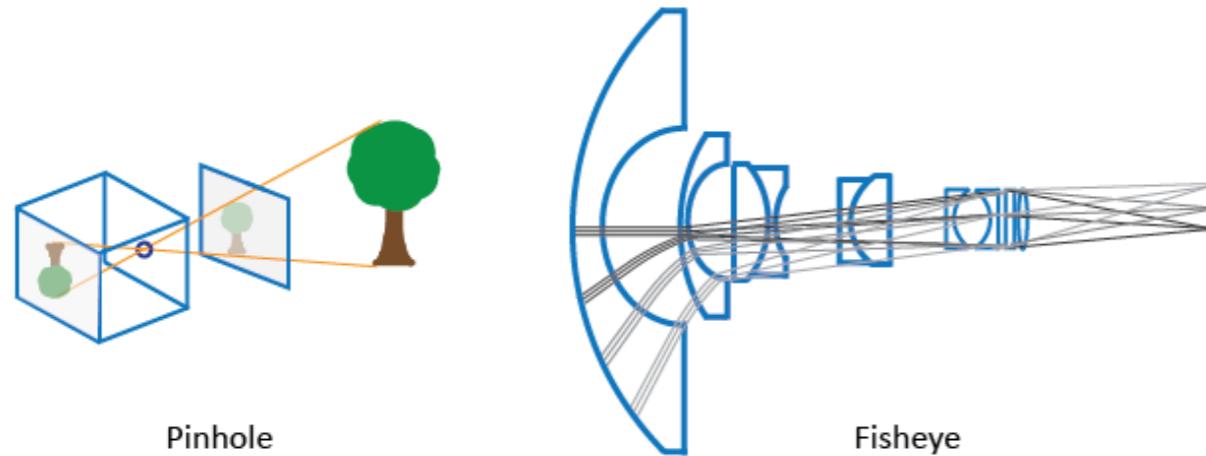
Summary: 3D to 2D Projection

1. **World-to-Camera**: 3D-3D projection. Representation of Rotation + Scaling + Translation.
→ **Camera extrinsic matrix**: depends on the position and orientation of the camera.
2. **Camera-to-Image**: 3D-2D projection. Loss of information. Requires pinhole camera model and its parameters.
→ **Camera intrinsic matrix**: depends on camera properties (e.g., focal length, pixel dimensions, resolution, etc.)
3. **Image-to-Pixel**: 2D-2D projection. Continuous to discrete. Quantization + Origin shift.



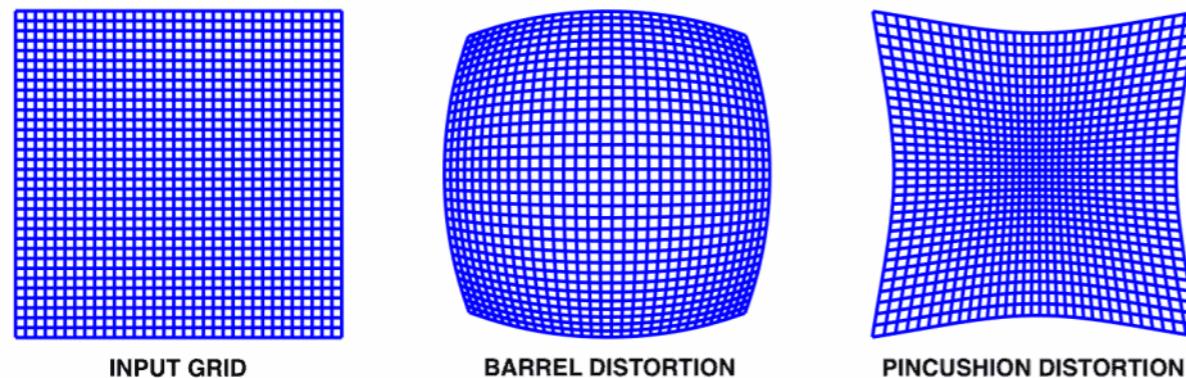
+ Camera Lens Distortion

- ✓ For the wider Field-of-View (FoV), various types of lens are used.



Pinhole

Fisheye



INPUT GRID

BARREL DISTORTION

PINCUSHION DISTORTION

Different types of the distortion



Distorted image with wide FoV

+ Camera Lens Distortion → Undistortion

- ✓ Before camera calibration, we need to correct the image distortion.



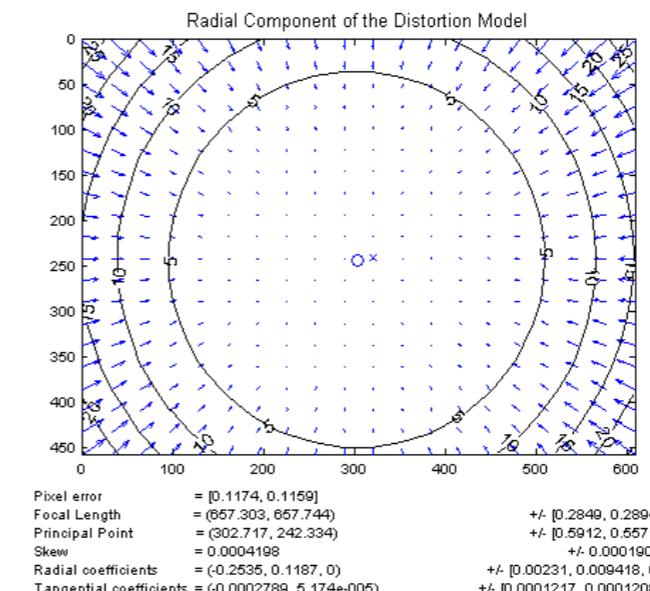
Undistortion



Radial distortion model

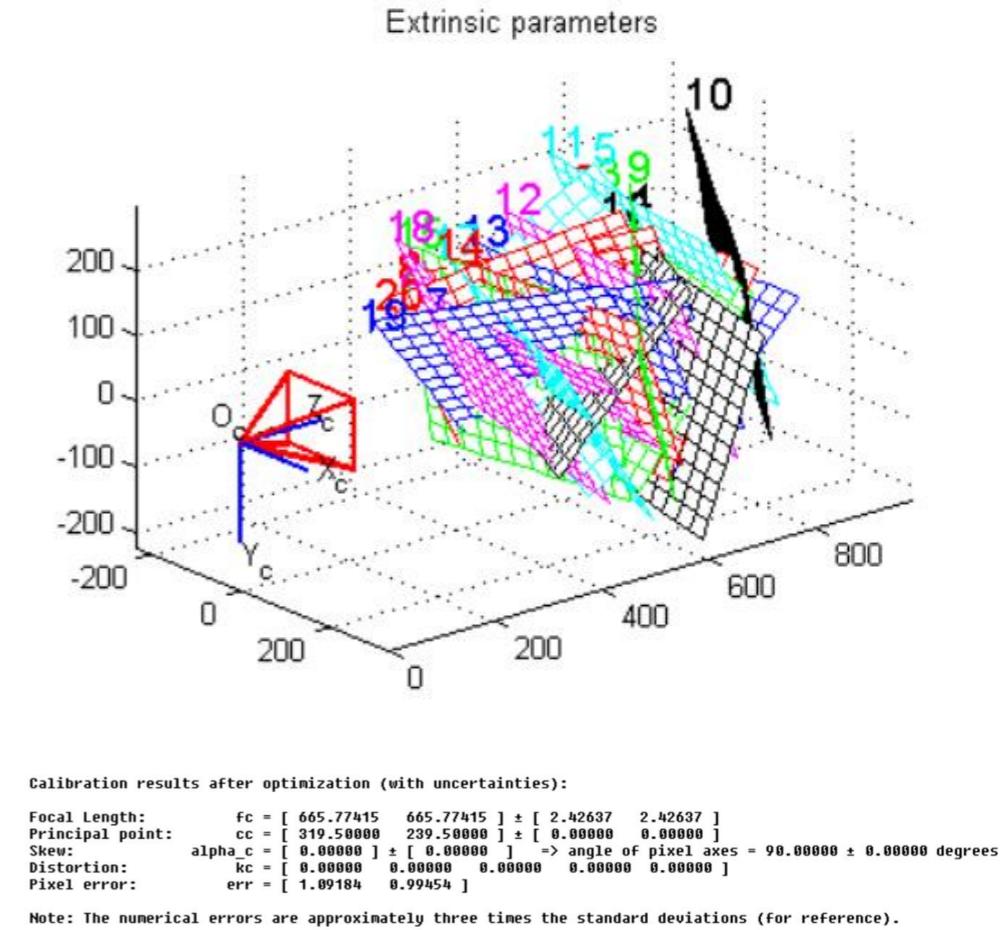
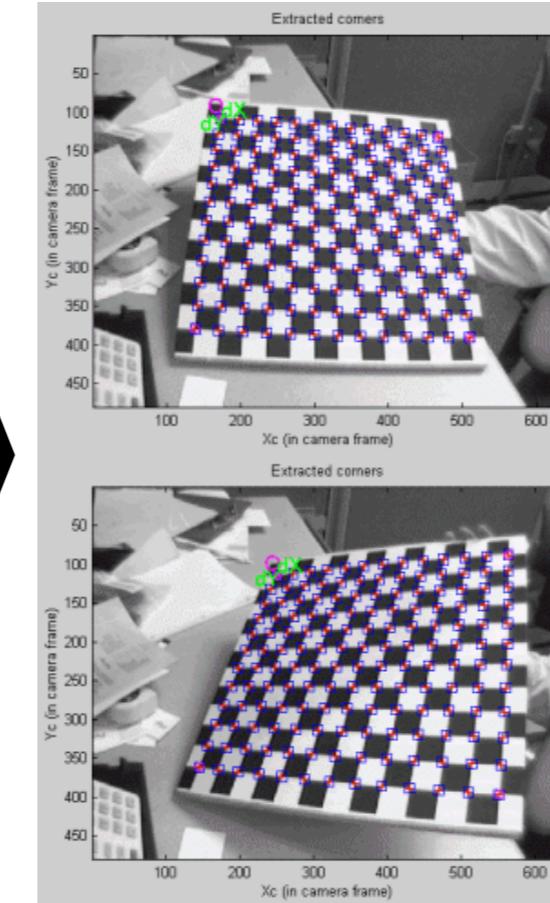
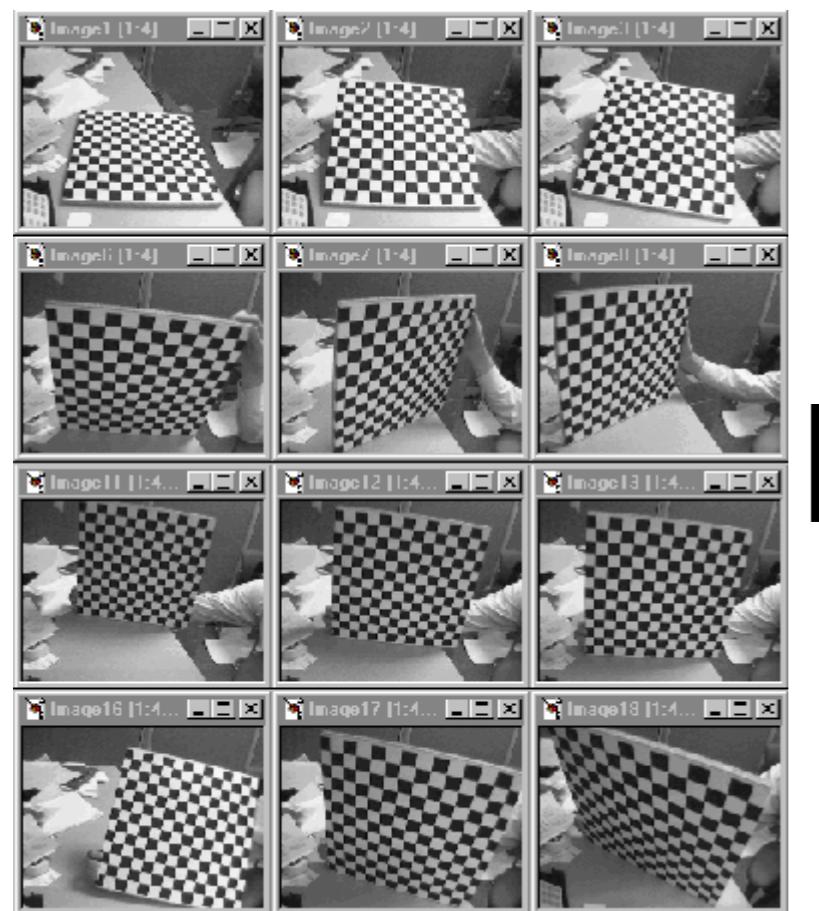
$$x_{\text{distorted}} = x \left(1 + k_1^* r^2 + k_2^* r^4 + k_3^* r^6 \right)$$

$$y_{\text{distorted}} = y \left(1 + k_1^* r^2 + k_2^* r^4 + k_3^* r^6 \right)$$



How to Calibrate Pinhole Camera?

- ✓ Camera Calibration Toolbox [1]



[1] <http://robots.stanford.edu/cs223b04/JeanYvesCalib/htmls/example.html>

Next Contents

- ✓ Python practice for basic camera geometry & image transformation
- ✓ Basic feature extraction & matching

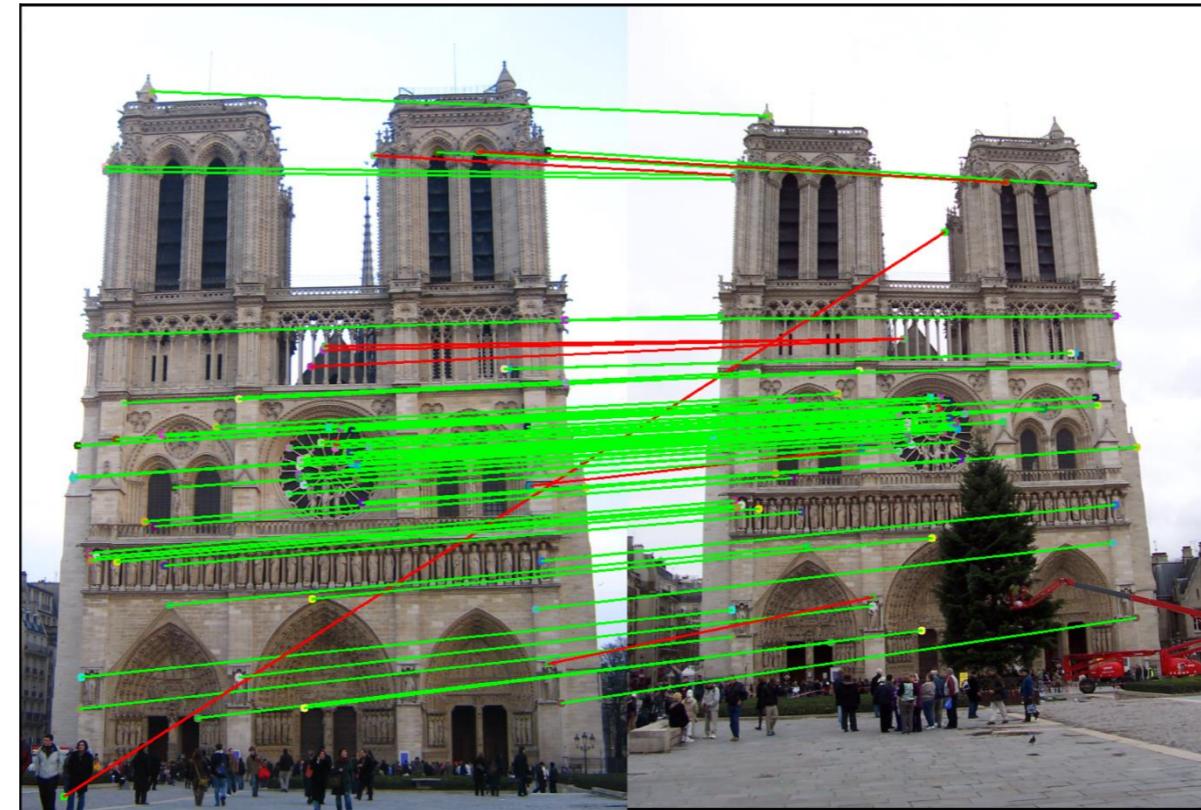


Figure by James Hays