KENTECH
Korea Institute of Energy Technology

Visionary Course – Energy AI
Week 09

Seokju Lee

KENTECH
Korea Institute of Energy Technology

Image Classification with Jetson

Configurations: Hello-AI-World by NVIDIA

- Follow **Quick Reference** in https://github.com/dusty-nv/jetson-inference/blob/master/docs/building-repo-2.md

```
$ sudo apt-get update
$ sudo apt-get install git cmake libpython3-dev python3-numpy
$ git clone --recursive https://github.com/dusty-nv/jetson-inference
$ cd jetson-inference
$ mkdir build
$ cd build
$ cmake -DENABLE_NVMM=OFF ../
$ make -j4
$ sudo make install
$ sudo ldconfig
```

./download_models.sh is available at **~/jetson_inference/tools**

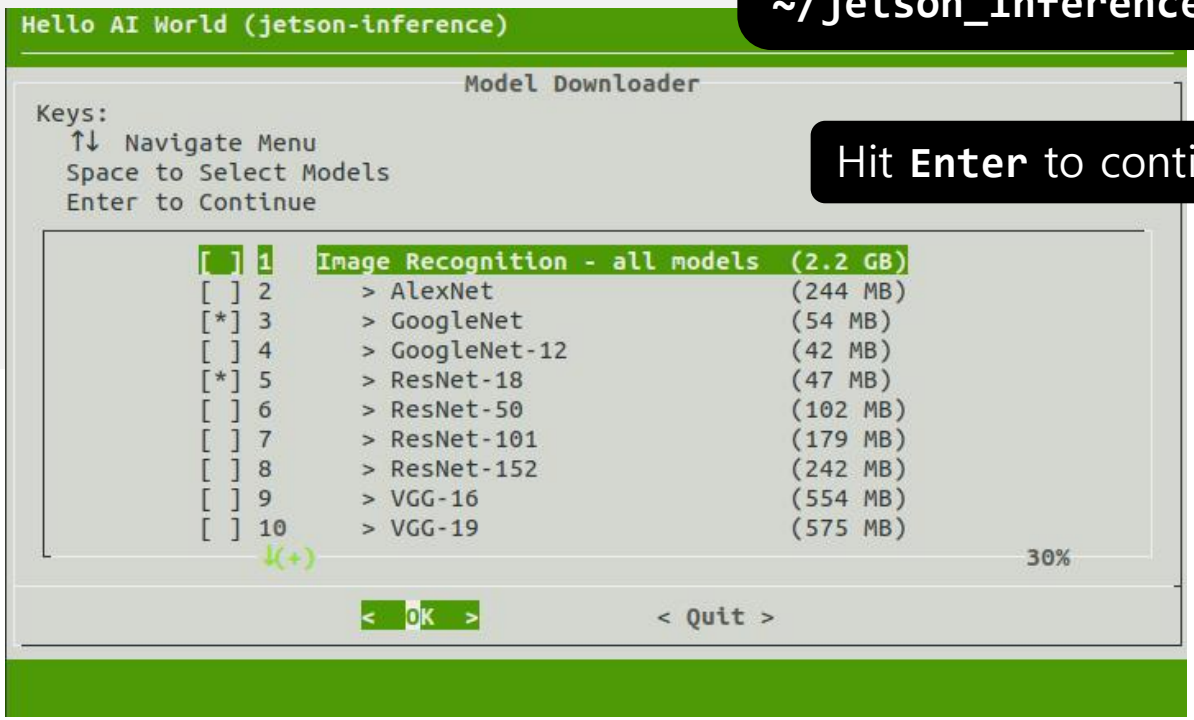

Hit **Enter** to continue!

You don't need to install PyTorch!

Experiments – Classification (Report Due ~11/3)

Before Starting
*Your basic workspace is here: "cd ~/jetson-inference/build/aarch64/bin" Every code is pre-built in this path.

Video Streaming

Q1.1. Run "python video-viewer.py csi://0" What is the output?

Q1.2. Run "python video-viewer.py --flip-method=rotate-180 csi://0" Discuss the differences.

Q1.3. Run "python video-viewer.py --flip-method=rotate-180 --input-width=640 --input-height=480 csi://0" Discuss the differences.

Q1.4. Run "python video-viewer.py --flip-method=rotate-180 --input-width=640 --input-height=480 --framerate=10 csi://0" Discuss the differences.

Live Demo for Image Classification

Q2.1. Run "python imagenet.py --flip-method=rotate-180" What is the output of the pop-up display? Let's check the terminal output. Please take a screenshot and paste it here. You can see some output values. What does each output (network name, class ID, floating-point number next to it, class name, each processing time, etc.) mean?

Experiments – Classification (Report Due ~11/3)

Q2.2. Go to the linked page (https://deeplearning.cms.waikato.ac.nz/user-guide/class-maps/IMAGENET/) and check that the class ID is matched to the class name. How many classes can the model distinguish in total? Please prepare your **own object (🕶️,🔒,✏️,👢,⚽)** corresponding to one of the above classes for further experiments (classification, detection, etc.).

Q2.3. Run "cd ~/jetson-inference/build; ./download-models.sh;" to download different CNN models (e.g., AlexNet, ResNet-50, etc.). Run "python imagenet.py --flip-method=rotate-180 --network=resnet-50". Please check the qualitative performance of each model.

Classify Your Own Objects or Images

Q3.1. Place the object in front of the camera and run the code (imagenet.py). Please take a screenshot of the result.

Q3.2. Position the object closer or further away from the camera. Please Analyze how confidence changes.

Q3.3. Download random images from Google and classify them. Please refer "python my-recognition.py images/banana_0.jpg --network=googlenet" and the below code:

```
import PIL
img = PIL.Image.open('jellyfish.jpg').resize((224,224))
img = np.array(img)
img_cuda = jetson.utils.cudaFromNumpy(img)      # CUDA image
class_id, confidence = net.Classify(img_cuda)   # Inference
class_desc = net.GetClassDesc(class_id)         # Predicted class
print(class_desc, confidence)
```

Q3.4. Please try other CNN models and repeat Q3.

Live Demo for Image Classification

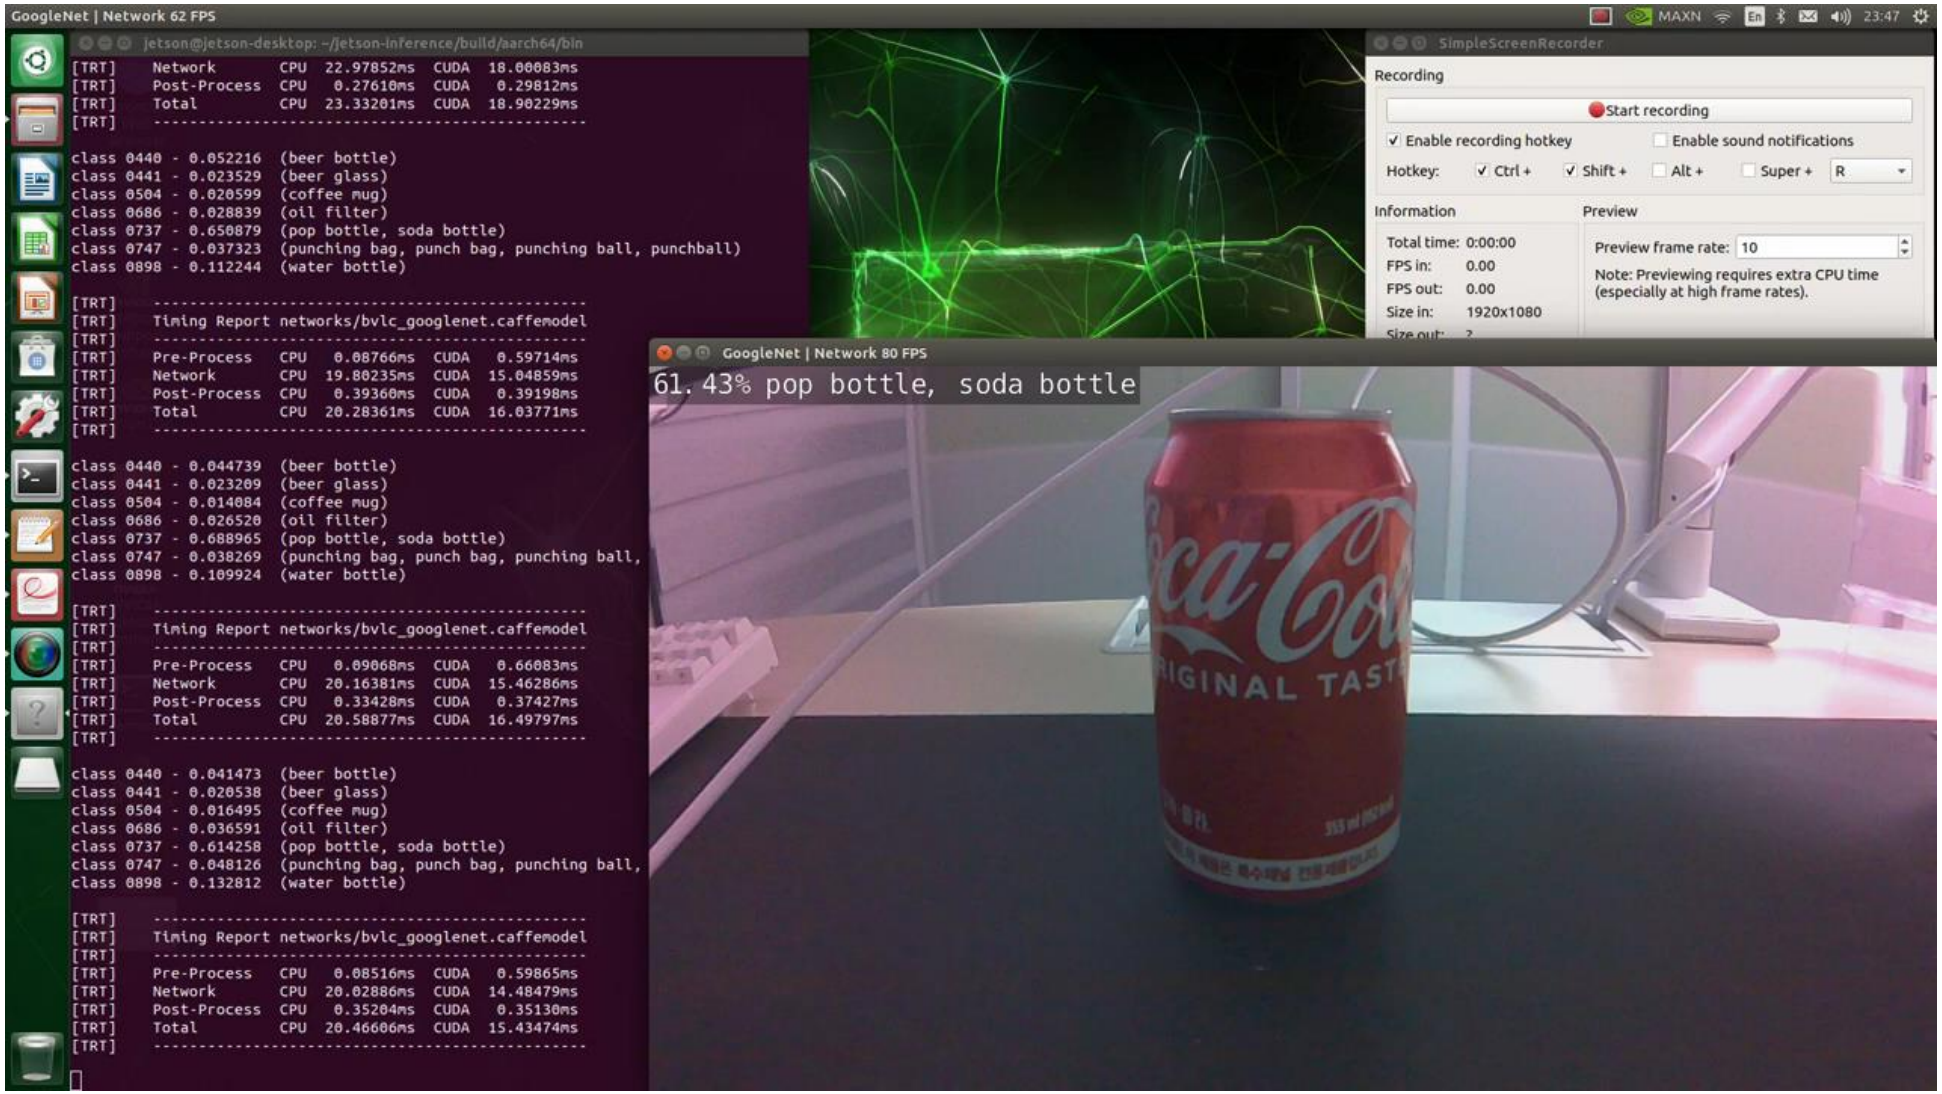

https://youtu.be/uMiXN-Z-hg4

Computer Vision Tasks

*Figure by Kim, et al., "Video Panoptic Segmentation" (CVPR 2020)

Model Complexity ⇧ Output dimension ⇧

Image

Video

Image Classification	Object Detection	Semantic Segmentation	Instance Segmentation	Panoptic Segmentation
	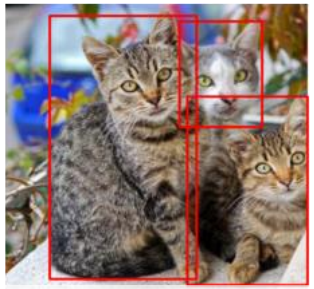	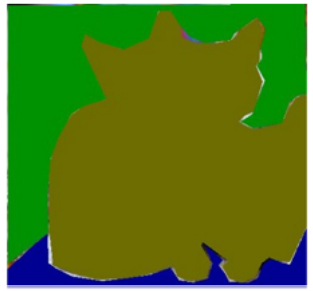	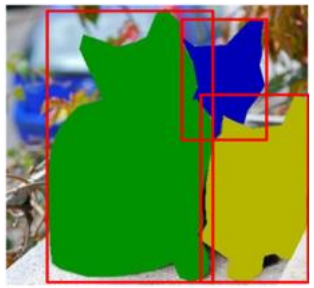	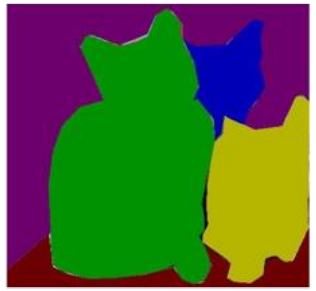
"Image-level"	"Box-level"	"Pixel-level"	"Pixel-level" + "Instance-level"	"Unknown pixel"
	Fast RCNN, CVPR 2015 **Faster RCNN** NeurIPS 2015 **YOLO,** CVPR 2016 **SSD,** ECCV 2016 **Cascade RCNN,** CVPR 2017 **CornerNet,** ECCV 2018 **HTC,** CVPR 2019	Dilated FCN, ICLR 2015 **DeepLab,** ICLR 2015 **PSPNet,** CVPR 2016 **Deform FCN,** ICCV 2017 **EncNet,** CVPR 2018 **DeepLabv2,** ECCV 2018 **DANet,** CVPR 2019 **Deformable V2,** ICCV 2019	DeepMask, NeurIPS 2015 MNC, CVPR 2016 FCIS, CVPR 2017 Cascade Mask-RCNN, PAMI 2019 **HTC,** CVPR2019 ⋮	
Karpathy et.al., CVPR 2014	**Kai et.al.,** CVPR 2016	**Xizhou et.al.,** CVPR 2017	**Linjie et.al.,** ICCV 2019	**Kim et.al.,** CVPR 2020
Video Classification	Video Object Detection	Video Semantic Segmentation	Video Instance Segmentation	Video Panoptic Segmentation

KENTECH
Korea Institute of Energy Technology

Object Detection with Jetson

Limitation of Image Classification: Dog or Cat?

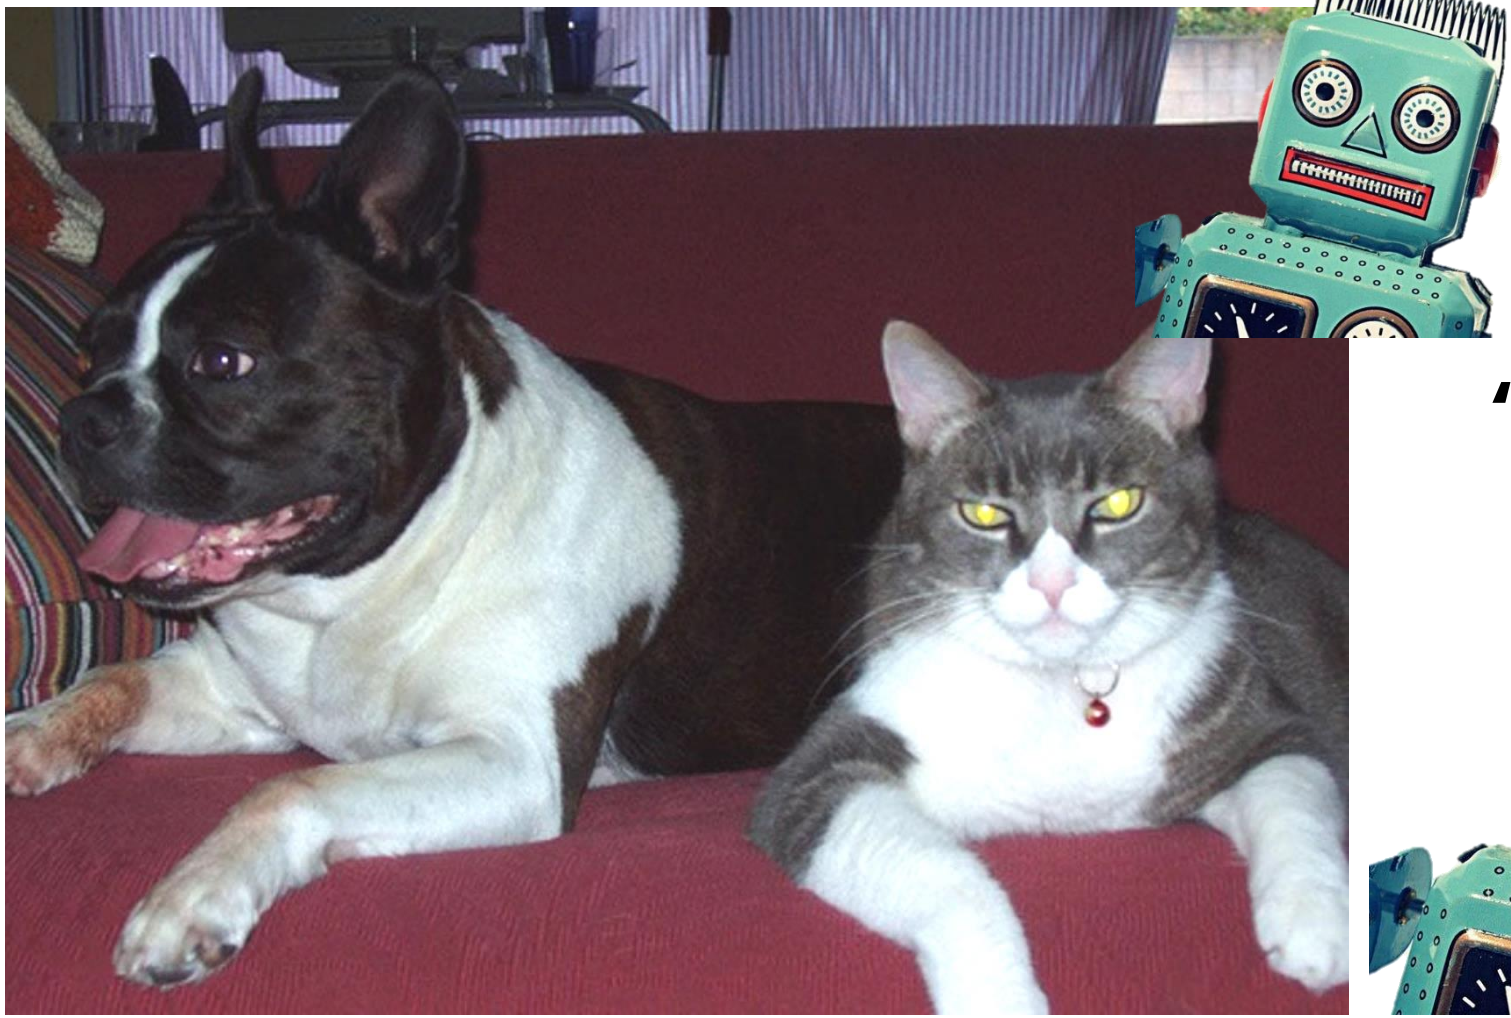

"Sofa?"

"Dog?"

"Cat?"

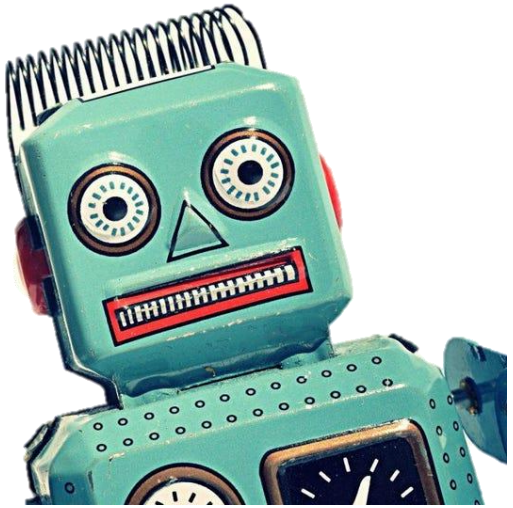

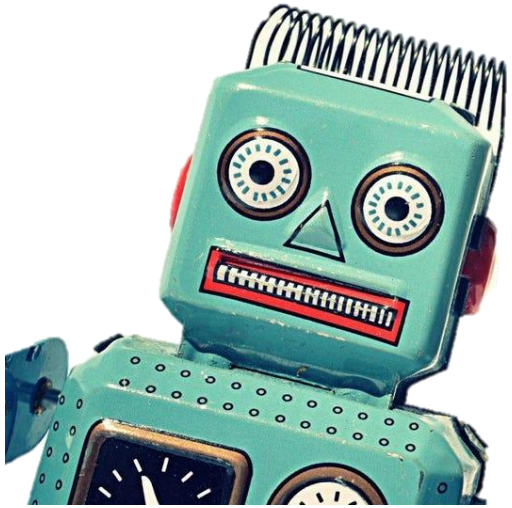

[1] S. Ren, et al., "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks." *NIPS* (2015).

Object Detection: What and Where?

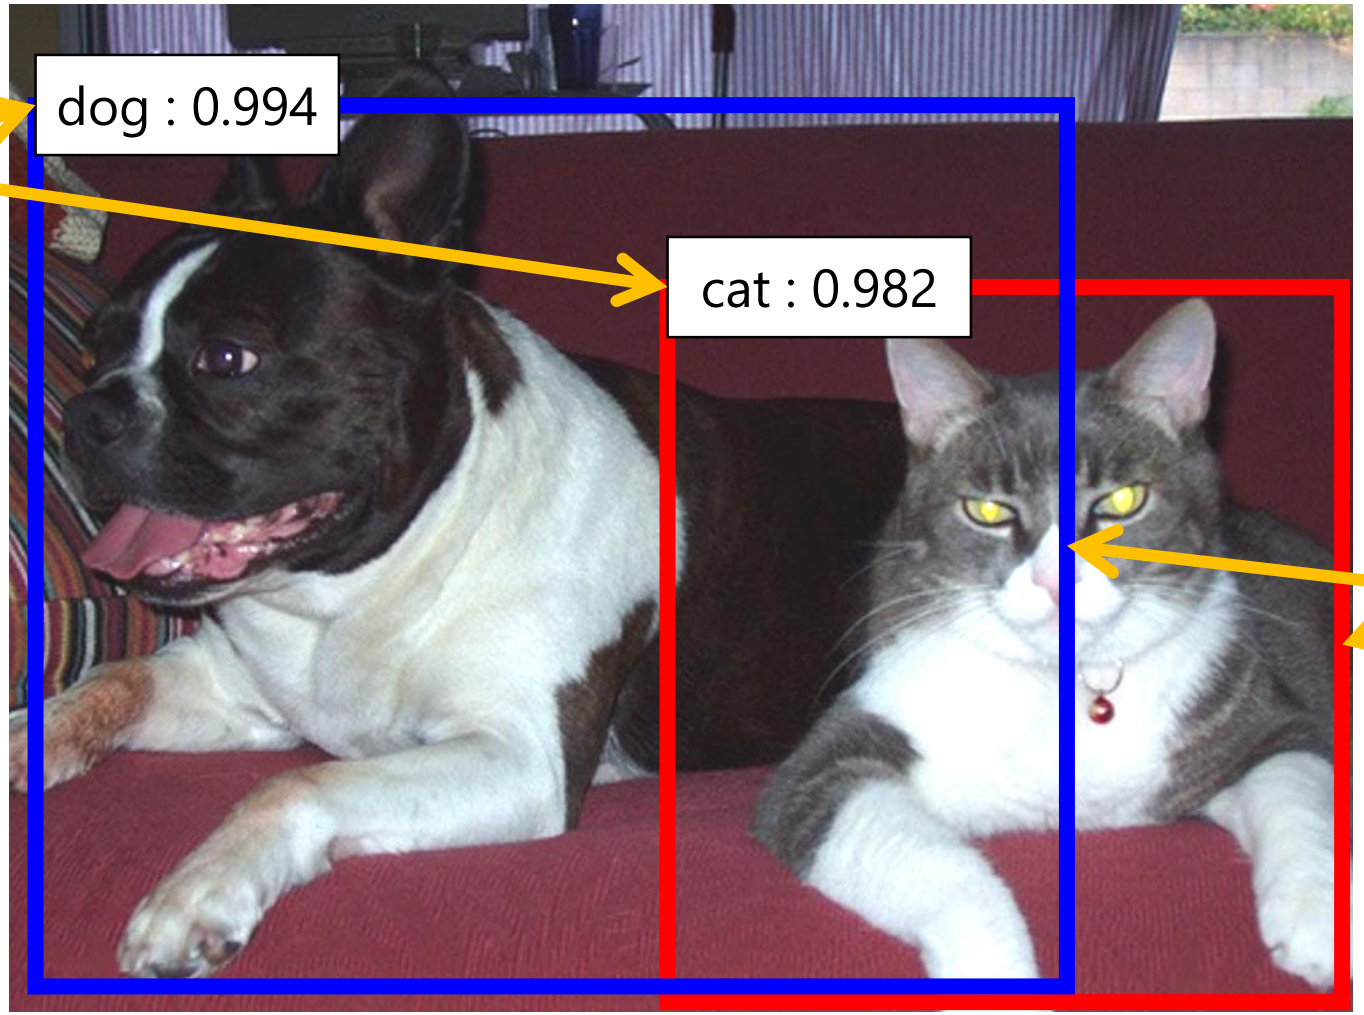

[1] S. Ren, et al., "Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks." *NIPS* (2015).

Object Detection: Input & Output

: Task of assigning **labels** & **bounding boxes** to all objects in the image.

Classify **which class** the object belongs to.

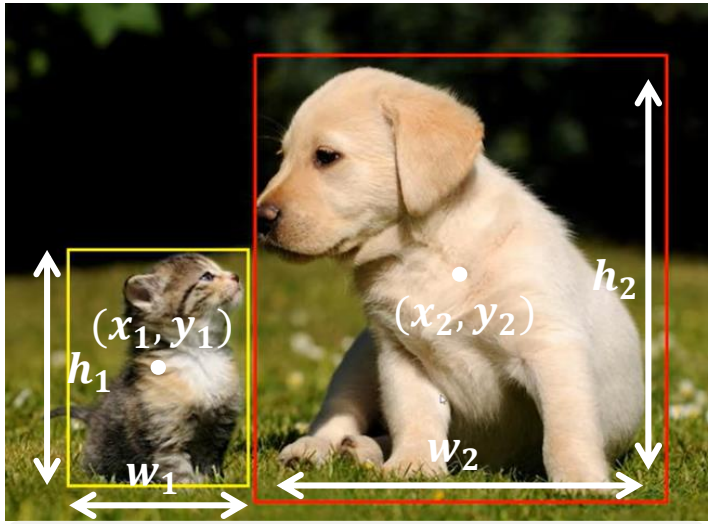

Find the coordinates of the bounding box **where** the object is located, and classify **which class** it belongs to.

Images	Class (=label)
I_1	cat
I_2	cat
I_3	dog

Classification

Images	Class (=label)	x	y	w	h
I_1	cat	60	210	100	180
I_1	dog	200	50	340	360
I_2	car	46	250	100	80

Classification + Regression

Object Detection: How to Design Models?

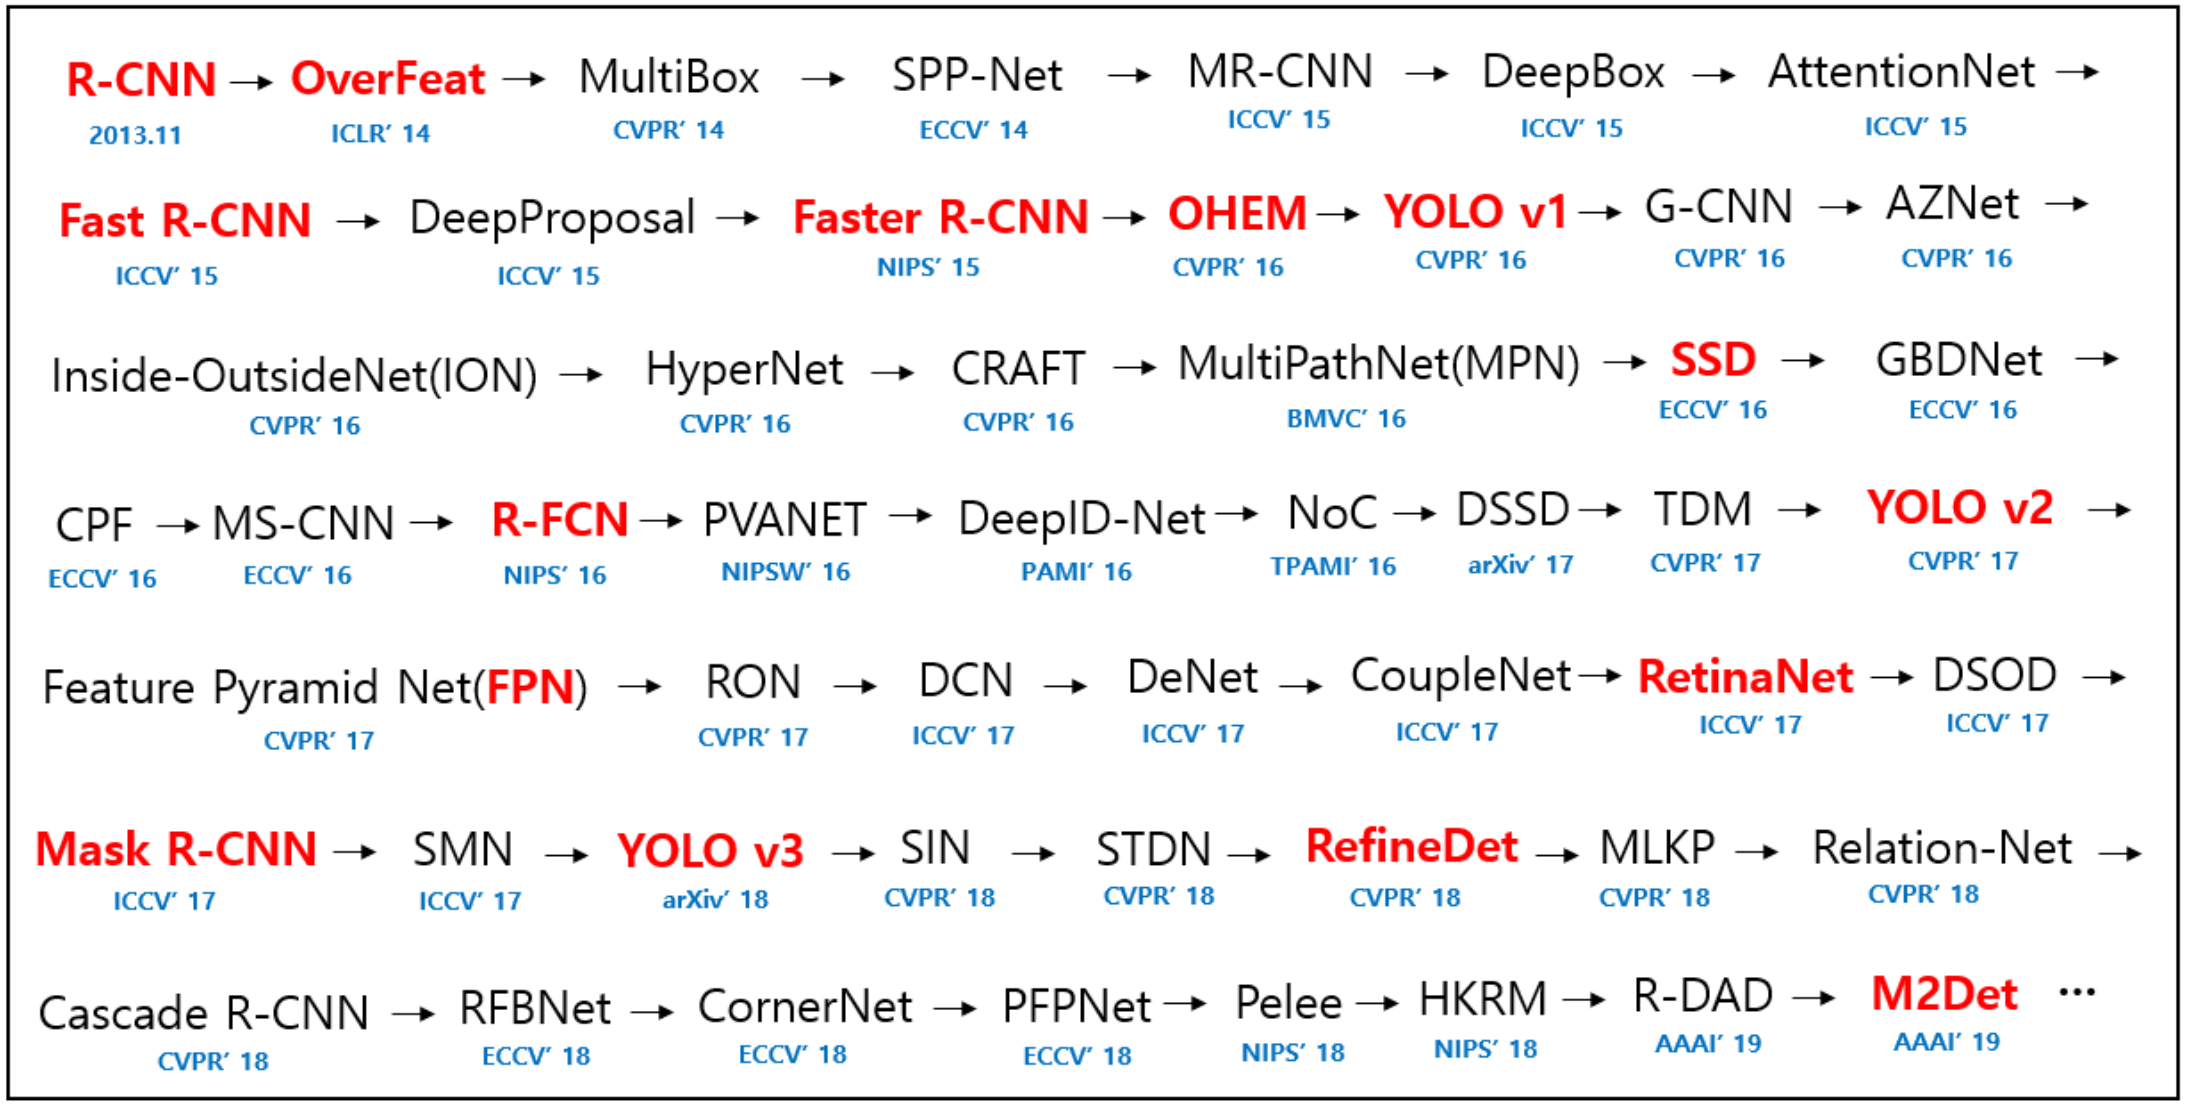

https://github.com/hoya012/deep_learning_object_detection

Live Demo for Object Detection

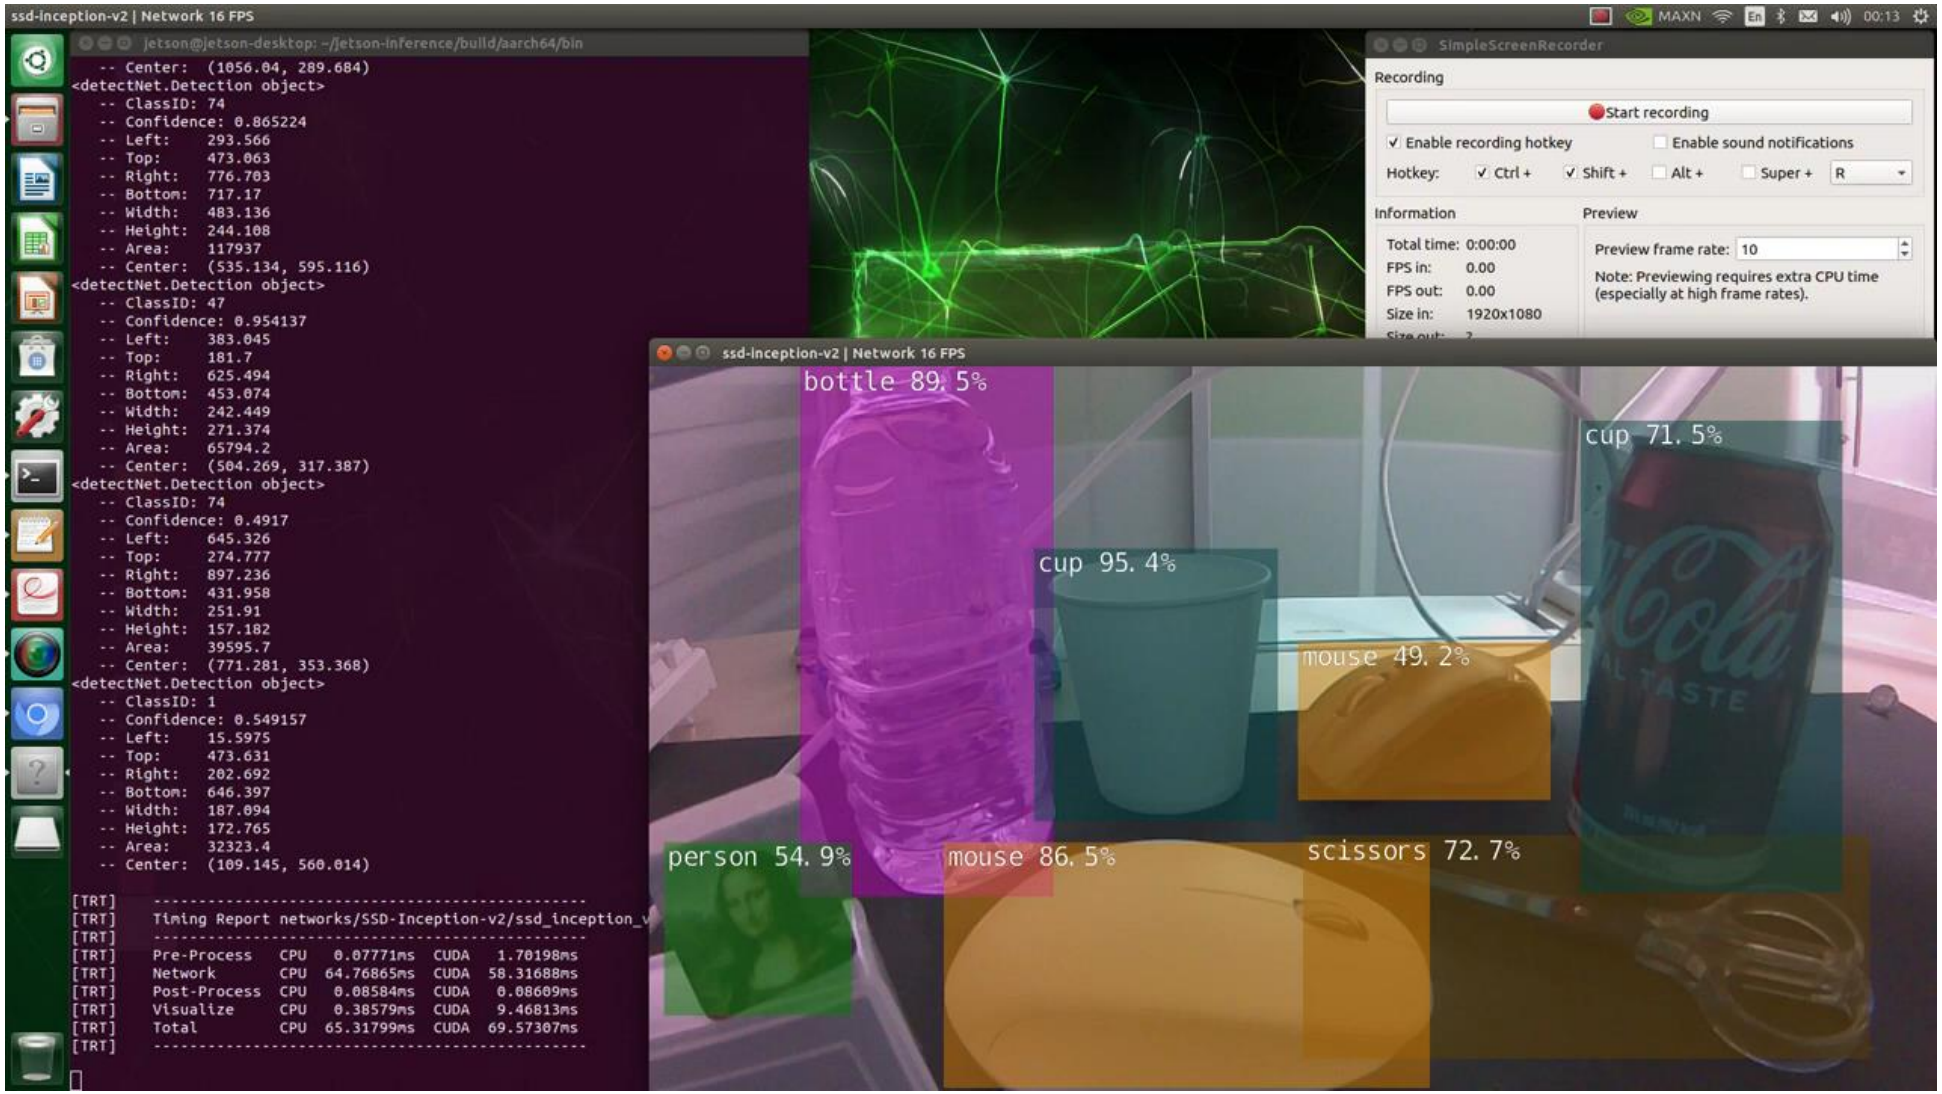

https://youtu.be/YU37WJ5z4Lg

Experiments – Object Detection

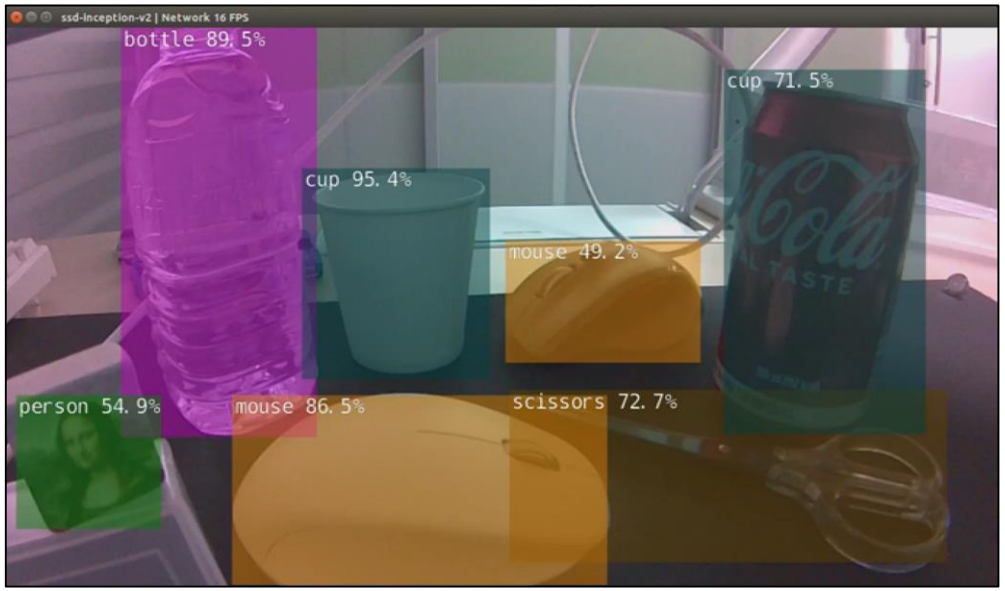

Your mission :

"Try to detect objects as much as you can!"

*Rule: the predicted class must be correct.

→ Submit the captured detection image.

*Due: ~11/7 (Mon)

Useful Commands

*Your basic workspace is here: "cd ~/jetson-inference/build/aarch64/bin"

Run "python detectnet.py --flip-method=rotate-180". Which model are you running?

Tip 1: Try Different Models for better results

You can download other models by running "./download_models.sh" at "~/jetson_inference/tools".

Run "python detectnet.py --network=ssd-inception-v2 --flip-method=rotate-180".

Tip 2: Try Different Thresholds for better results

Control threshold (default threshold =0.5 e.g., 0.3 & 0.7) by runing "python detectnet.py --threshold=0.3 --flip-method=rotate-180". What does threshold mean?

JetRacer Mission: Stop-and-Go

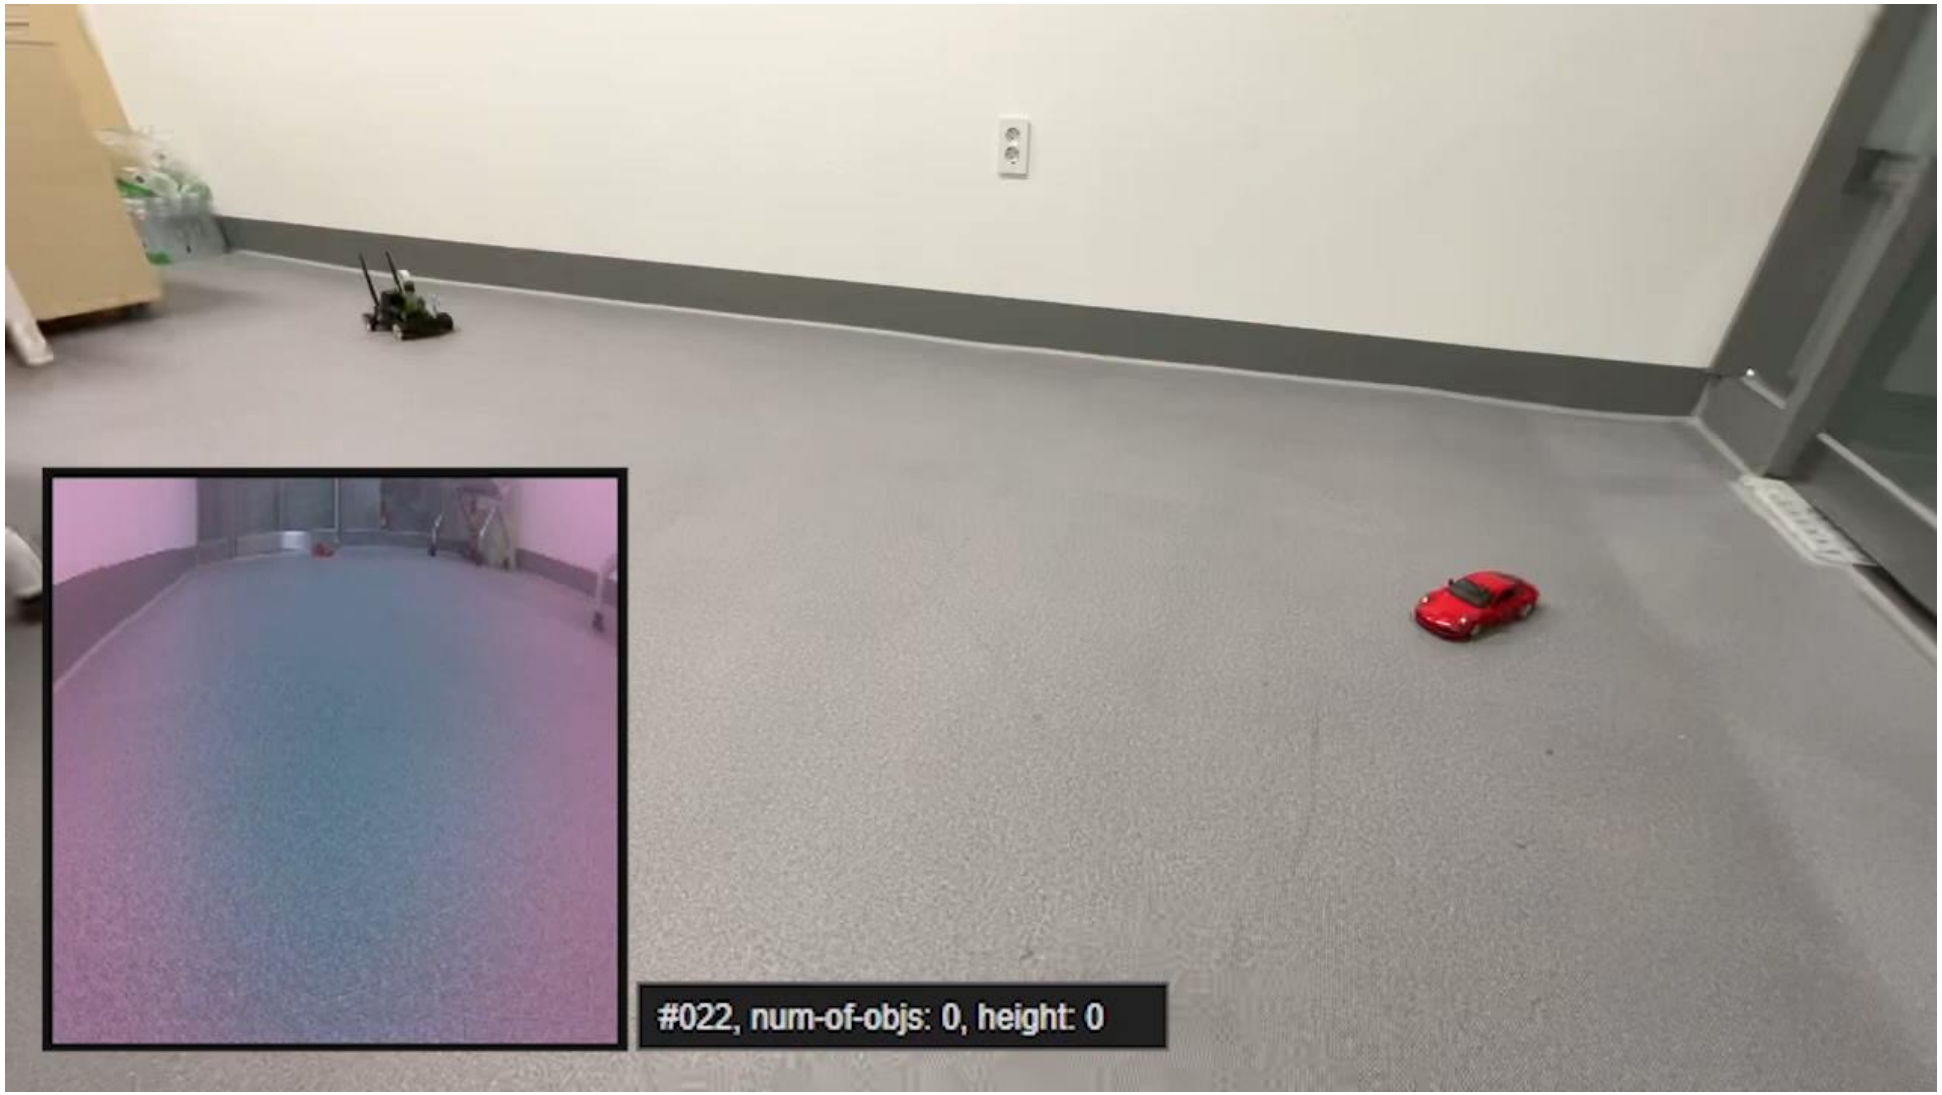

https://youtu.be/XCbgLJEEtDg

Experiments

Some Useful Tips while Debugging

*Sometimes, the python process does not respond. In this case, please terminate the process with `ctrl+c`. If it still does not respond at all, forcibly stop the process with `ctrl+z`, and check the running process name with the `ps -a` command, and then type `sudo pkill -9 [name-of-process]` command to kill the process. If you don't shut it down, it will remain as a 🧟zombie🧟 and keep occupying the processor (CPU or GPU) in the background.

*Sometimes, the best solution for resolving an issue is just rebooting the system.


Q&A
KENTECH
Korea Institute of Energy Technology